

AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Source: Artificial Analysis (artificialanalysis.ai) — April 2026


Audience: AI Engineering Leads & Product Managers evaluating LLM |

11

Frontier Models Evaluated

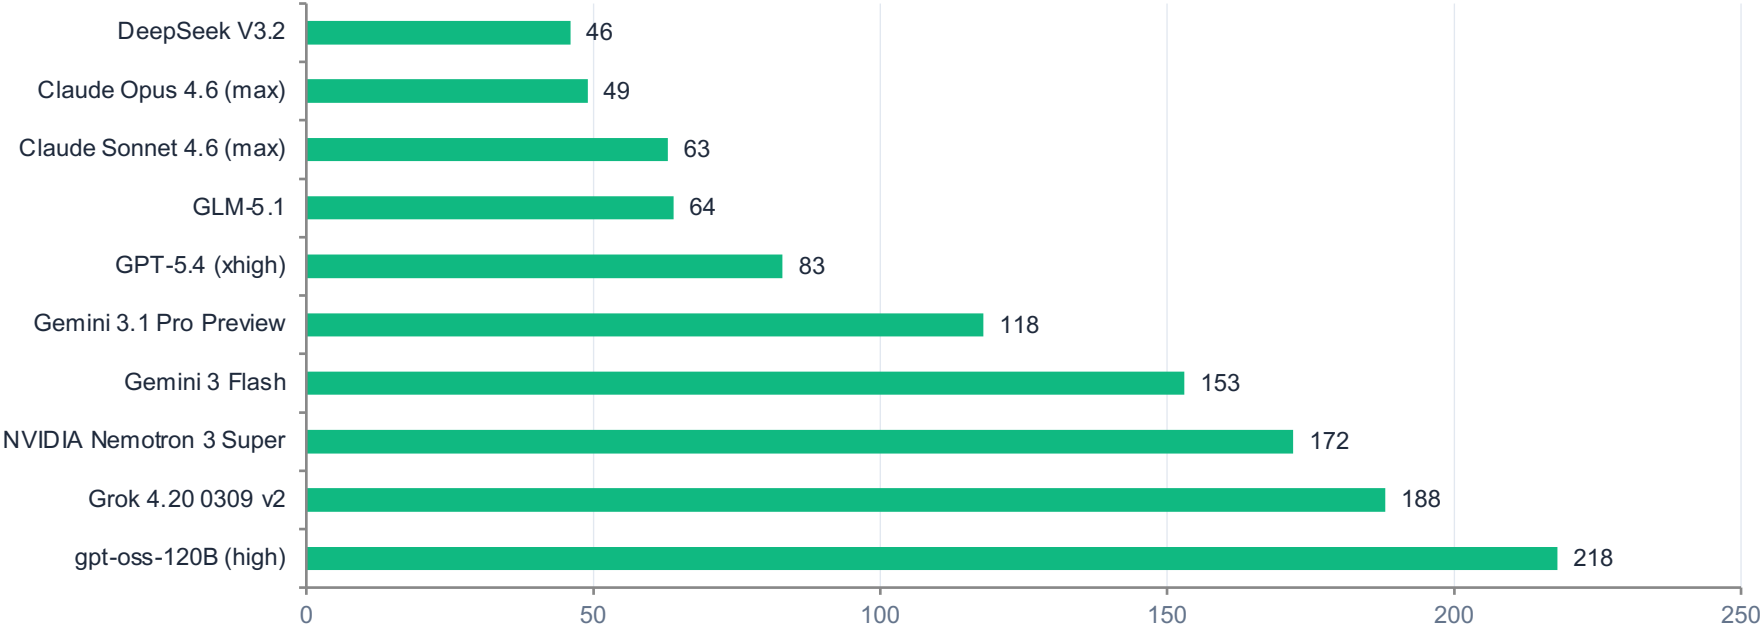
Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index — composite score across reasoning, knowledge, and instruction-following benchmarks



gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Output tokens per second — critical metric for latency-sensitive applications



Key insight: Speed and intelligence are inversely correlated — the fastest model (218 tok/s) scores only 33 on intelligence.

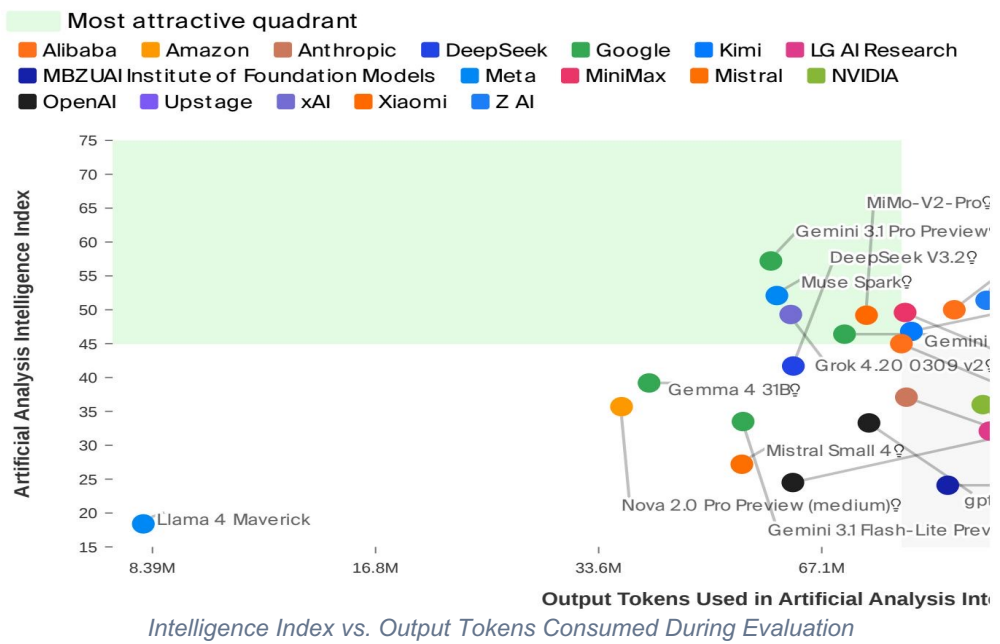
From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

USD per 1 million output tokens — reflects differences in compute, margin strategies, and competitive positioning



33x cost gap: Budget-sensitive workloads can access frontier-adjacent intelligence at \$0.30/M tokens vs. \$10.00 for top proprietary.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Proprietary vs. Open-Weight Leaders

Metric	Proprietary	Open-Weight
Intelligence	Gemini 3.1 Pro Preview: 57	GLM-5.1: 51
Speed	Gemini 3 Flash: 153 tok/s	gpt-oss-120B: 218 tok/s
Price	Gemini 3 Flash: \$1.10/M	gpt-oss-120B: \$0.30/M

Key Takeaways:

- Open-weight leads on speed & price
- Proprietary leads on intelligence (by slim 6-pt margin)
- No single model dominates all three axes

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 both score 57 on Intelligence Index

Speed ≠ Intelligence: Fastest model (218 tok/s) scores only 33; smartest operate at 49–118 tok/s

33x pricing spread: \$0.30/M (gpt-oss-120B, DeepSeek V3.2) to \$10.00/M (Claude Opus 4.6)

Open-weight closing gap: GLM-5.1 at 51 intelligence — within 10% of proprietary leaders, at \$2.10/M

Actionable Guidance for Model Selection

Max intelligence: Gemini 3.1 Pro Preview or GPT-5.4 (score 57)

Latency-critical: gpt-oss-120B (218 tok/s) or Grok 4.20 (188 tok/s)

Budget-optimized: gpt-oss-120B / DeepSeek V3.2 (\$0.30/M) — 33x cheaper than top proprietary

Balanced value: GLM-5.1 (51 intelligence, \$2.10/M) — best open-weight intelligence-per-dollar