

ARTIFICIAL ANALYSIS • APRIL 2026

AI Model Landscape

April 2026

Benchmarking Intelligence, Speed, and Cost
Across 11 Frontier Models



11

MODELS EVALUATED
Source: artificialanalysis.ai

3

KEY DIMENSIONS


33×

PRICING SPREAD

57

TOP INTEL. SCORE

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57




Full Rankings

#	Model	Provider	Score
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

Co-leaders Open-weight

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4× Faster Than Claude Opus 4.6



218 tokens/s

gpt-oss-120B sets the pace — 4× faster than Claude Opus 4.6 (49 t/s)



Inverse relationship

The two smartest models (Gemini 3.1 Pro, GPT-5.4) hit only 118 and 83 t/s



Open-weight advantage


4 of the top 5 fastest are open-weight or smaller-parameter models



Speed ≠ intelligence

Latency-sensitive workloads can pick fast/cheap models without sacrificing throughput

From \$0.30 to \$10 per Million Tokens: A 33× Pricing Gap Across Models



\$0.30

/1M tokens

gpt-oss-120B & DeepSeek V3.2
lowest available price

33×

price spread


\$0.30 cheapest → \$10.00
Claude Opus 4.6 most expensive

\$1.10

/1M tokens

Gemini 3 Flash: best mid-range
value (153 t/s, score 46)

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Intelligence Index vs. Output Tokens Consumed (evaluation efficiency frontier)

Proprietary vs. Open-Weight

Metric	Proprietary	Score	Open-Weight	Score
Intelligence	Gemini 3.1 Pro	57	GLM-5.1	51
Speed (t/s)	Gemini 3 Flash	153	gpt-oss-120B	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B	\$0.30

- Frontier intelligence requires proprietary models — for now.
- Open-weight leads on speed AND price across all dimensions.
- No single model dominates all three axes simultaneously.
- Trade-offs require use-case-driven selection.

Open Weights Close the Gap:

GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

01 Two-way tie at the top

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 — making Google and OpenAI co-leaders as of April 2026.

02 Speed and intelligence are inversely correlated

Fastest model (gpt-oss-120B, 218 t/s) scores only 33 on intelligence; smartest models run at 49–118 t/s.

03 Massive pricing spread

33× cost gap: \$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2) vs. \$10.00/1M tokens (Claude Opus 4.6).

04 Open-weight models are closing the gap

GLM-5.1 at 51 intelligence is within 10% of proprietary leaders — at just \$2.10/1M tokens.

Model Selection Guidance

INTELLIGENCE

Gemini 3.1 Pro
or **GPT-5.4**

Score 57 — best for complex reasoning

SPEED

gpt-oss-120B

218 t/s — latency-critical apps

COST

gpt-oss-120B or
DeepSeek V3.2

\$0.30/1M — high-volume / budget

BALANCE

Gemini 3 Flash

\$1.10/1M • 153 t/s • Score 46