

11

frontier models benchmarked

57

Top Intelligence Index (tie)

218

Max Output Tokens/s


33x

Price Gap (cheapest vs
priciest)

AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57



Key Insight


Two models tie at 57, but the spread runs from 57 to 33 — a 73% gap from top to bottom.

Open-weight GLM-5.1 (51) is within 10% of the leaders.

Provider Spread

8 organizations across the US, China, and Europe — a genuinely global competition.

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6



Speed vs Intelligence


The fastest model (218 tokens/s) scores only 33 on intelligence. The smartest models run at 49-118 tokens/s.

Open-weight models dominate throughput.

Sweet Spot

Gemini 3 Flash pairs 153 tokens/s with Intelligence 46 — a compelling balance for latency-sensitive apps.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



33x

price gap from cheapest (\$0.30) to most expensive (\$10.00)

Value Picks

Budget:

gpt-oss-120B & DeepSeek V3.2 at \$0.30/1M tokens


Balanced:

Gemini 3 Flash: \$1.10 + 153 tok/s + Intel 46

Open-weight:

GLM-5.1: Intel 51 at just \$2.10/1M tokens

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Intelligence Index vs. output tokens consumed during evaluation (not throughput)

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tok)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Proprietary vs. Open-Weight Leaders

Every model requires trade-offs

High intelligence = lower speed, higher cost
 High speed = lower intelligence scores
 Low cost = sacrifices on capability

Selection must be use-case driven.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Intelligence is contested

Gemini 3.1 Pro Preview and GPT-5.4 tie at 57. The 57-to-33 spread shows not all "frontier" models are equal — verify benchmarks for your domain.

Speed and intelligence are inversely correlated

The fastest model (218 tok/s) scores only 33 on intelligence. For latency-sensitive pipelines, Gemini 3 Flash balances 153 tok/s with Intelligence 46.

Open-weight models are closing the gap

GLM-5.1 (51 intelligence, \$2.10/M tokens) is within 10% of proprietary leaders. gpt-oss-120B leads on speed (218 tok/s) and ties on price (\$0.30/M tokens).

Recommendation: Define your workload priorities first, then select models that optimize the axis that matters most — intelligence, speed, or cost.