

AI Model Landscape

April 2026

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models



Source: Artificial Analysis | April 2026

For: AI Engineering Leads & Product Managers

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

#	Model	Provider	Score
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33


gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6

#	Model	Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

#	Model	\$/1M Tokens
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Intelligence Index vs. Output Tokens Consumed

Source: Artificial Analysis, April 2026

Proprietary vs. Open-Weight Leaders

Metric	Proprietary	Score	Open-Weight	Score
Intelligence	Gemini 3.1 Pro	57	GLM-5.1	51
Speed	Gemini 3 Flash	153	gpt-oss-120B	218
Price	Gemini 3 Flash	\$1.10	gpt-oss-120B	\$0.30

Key Insight

The two intelligence leaders (Gemini 3.1 Pro Preview, GPT-5.4) operate at 83-118 tokens/s. The fastest model (gpt-oss-120B at 218 tokens/s) scores only 33 on intelligence. Every model requires trade-offs.

Open Weights Close the Gap

GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

- 1 Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 both score 57 on the Intelligence Index.
- 2 Speed and intelligence are inversely correlated: gpt-oss-120B at 218 tokens/s scores only 33 intelligence.
- 3 33x pricing gap: From \$0.30 (gpt-oss-120B, DeepSeek V3.2) to \$10.00 (Claude Opus 4.6) per 1M tokens.
- 4 Open-weight models are closing the gap: GLM-5.1 at 51 intelligence is within 10% of proprietary leaders at \$2.10/M tokens.
- 5 No single model wins on all three axes — model selection must be driven by your specific workload priorities.