

April 2026

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Data source: Artificial Analysis (captured April 10, 2026).

Intelligence leader (tie)

57

Fastest output speed

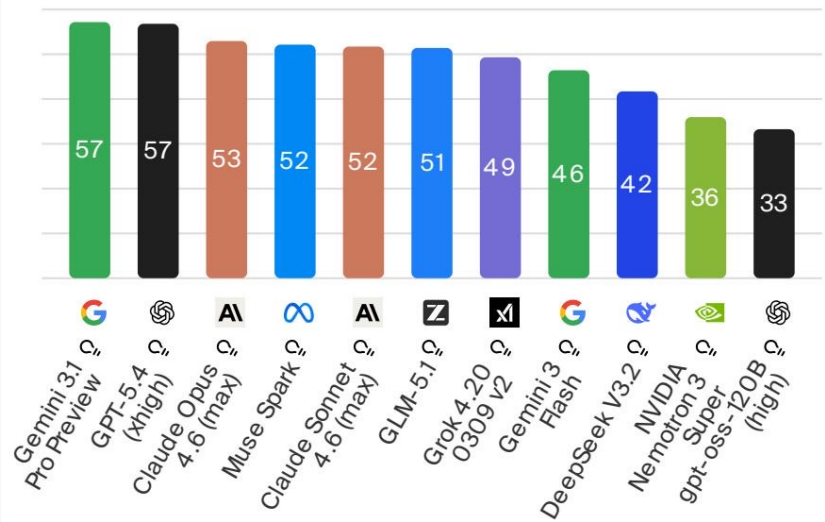
218 tok/s

Cheapest published price

\$0.30 / 1M

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

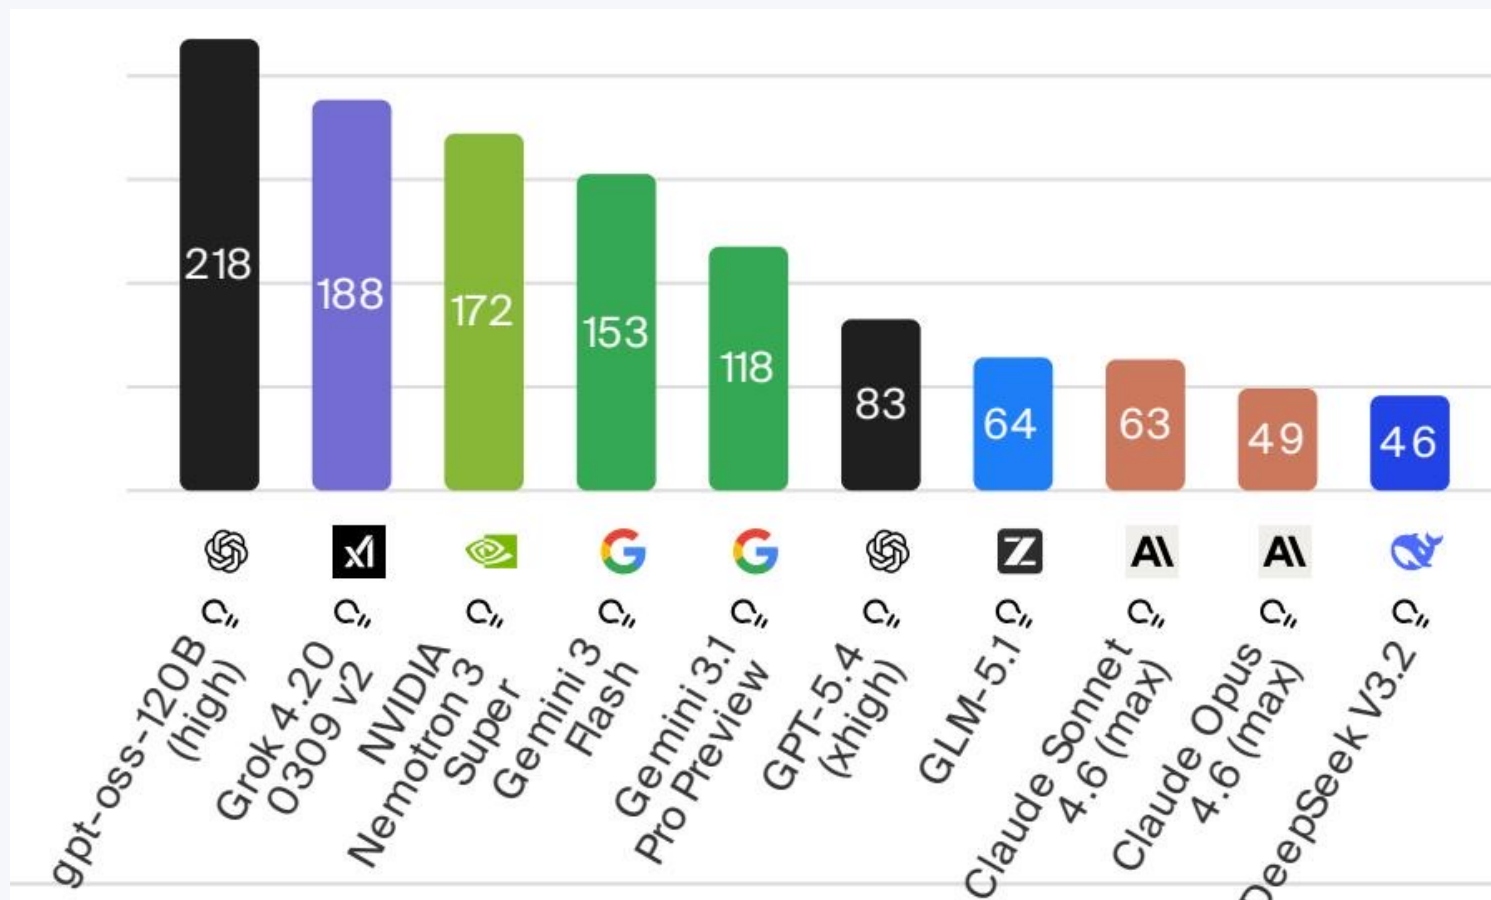
Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33



Index spread: 57 → 33 (top-to-bottom).

Open-weight momentum: GLM-5.1 reaches 51 (6 points behind the leaders).

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6



Fastest

218 tok/s

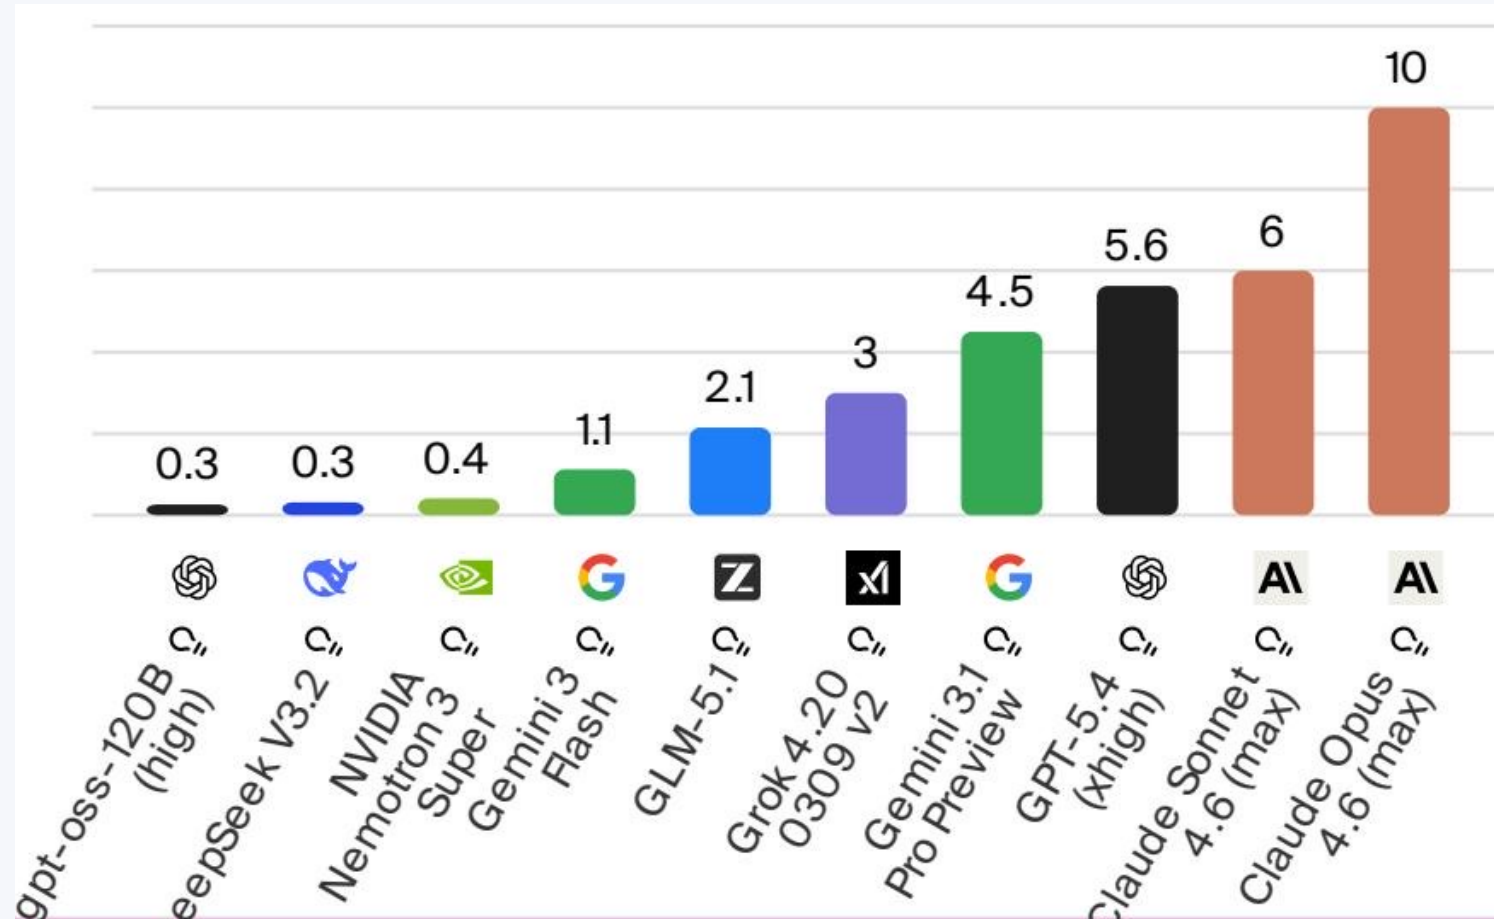
Slowest (in speed table)

46 tok/s

Interpretation

Speed tends to trade off with top-end intelligence; the co-leaders sit mid-pack at 118 and 83 tok/s.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



Cheapest

\$0.30 / 1M

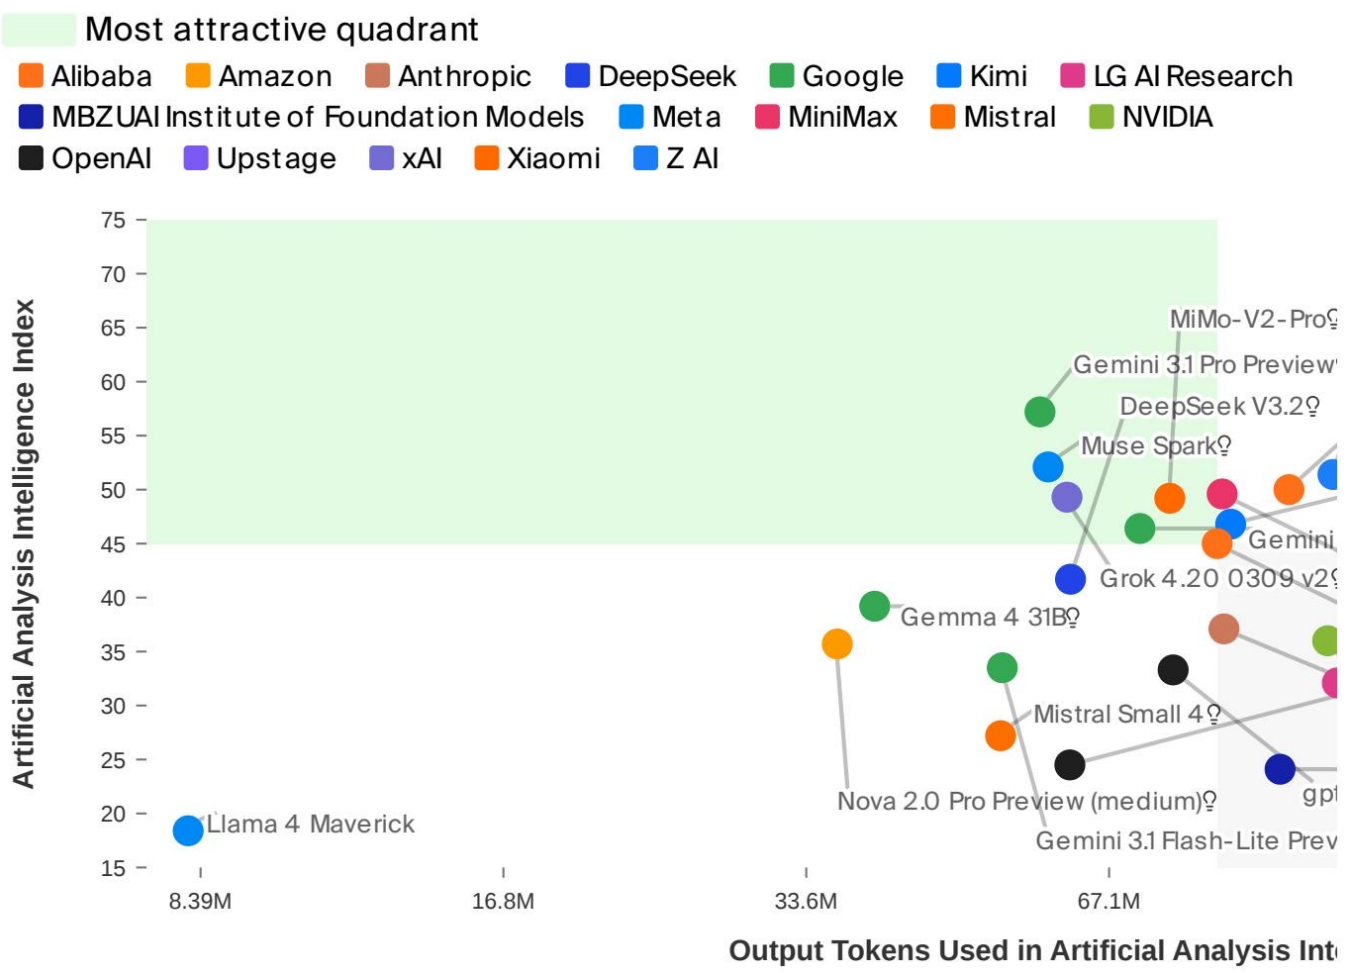
Most expensive

\$10.00 / 1M

Value signal

Gemini 3 Flash pairs \$1.10 / 1M with 153 tok/s and an Intelligence Index of 46.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Metric	Proprietary Leader	Score	Open-Weight Leader
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)

What this implies

Proprietary models still lead on raw intelligence — but by a thin margin.

Open-weight models dominate throughput and cost on these benchmarks.

Selection is workload-driven: decide which constraint is non-negotiable.

Scatter plot: Intelligence Index vs output tokens consumed during evaluation (token efficiency, not throughput).

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

If intelligence is the priority

Start with the co-leaders: Gemini 3.1 Pro Preview or GPT-5.4 (xhigh) at 57.
Validate application-specific reliability with your own evals.

If latency/throughput is the priority

Prefer open-weight speed leaders: gpt-oss-120B at 218 tok/s.
Use mid-tier options when you need more capability: Grok 4.20 at 188 tok/s.

If cost is the priority

For budget-sensitive deployments, \$0.30 / 1M options are decisive (gpt-oss-120B, DeepSeek V3.2).
Consider balanced value points like Gemini 3 Flash (\$1.10 / 1M).

Bottom line: no single model dominates intelligence, speed, and price — align model choice to your workload constraints and re-evaluate as open weights keep improving.

Source: Artificial Analysis (captured April 10, 2026)