Frontier LLM Landscape: April 2026

Competitive landscape of frontier AI models for engineering and product decisions

- Two-way tie at the frontier: Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on intelligence.
- No single model dominates intelligence, speed, and price simultaneously.
- Open-weight models now lead on throughput and cost, while closing the intelligence gap.
- Selection should be workload-specific, not leaderboard-only.

57 Top intelligence score

218 Fastest output tokens/s

33x Price spread (min to max)

Axis 1 — Intelligence: Who Leads on Capability?

Artificial Analysis Intelligence Index (full 11-model ranking)

Rank	Model	Provider	Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

What this implies

- Google and OpenAI are tied at the frontier (57).
- Top open-weight model (GLM-5.1) reaches 51.
- Gap to frontier is now 6 points, not double-digits.

Decision lens

For high-stakes reasoning tasks, shortlist from the top 5 and optimize speed/cost in secondary routing.

Axis 2 — Speed: Latency and Throughput Reality

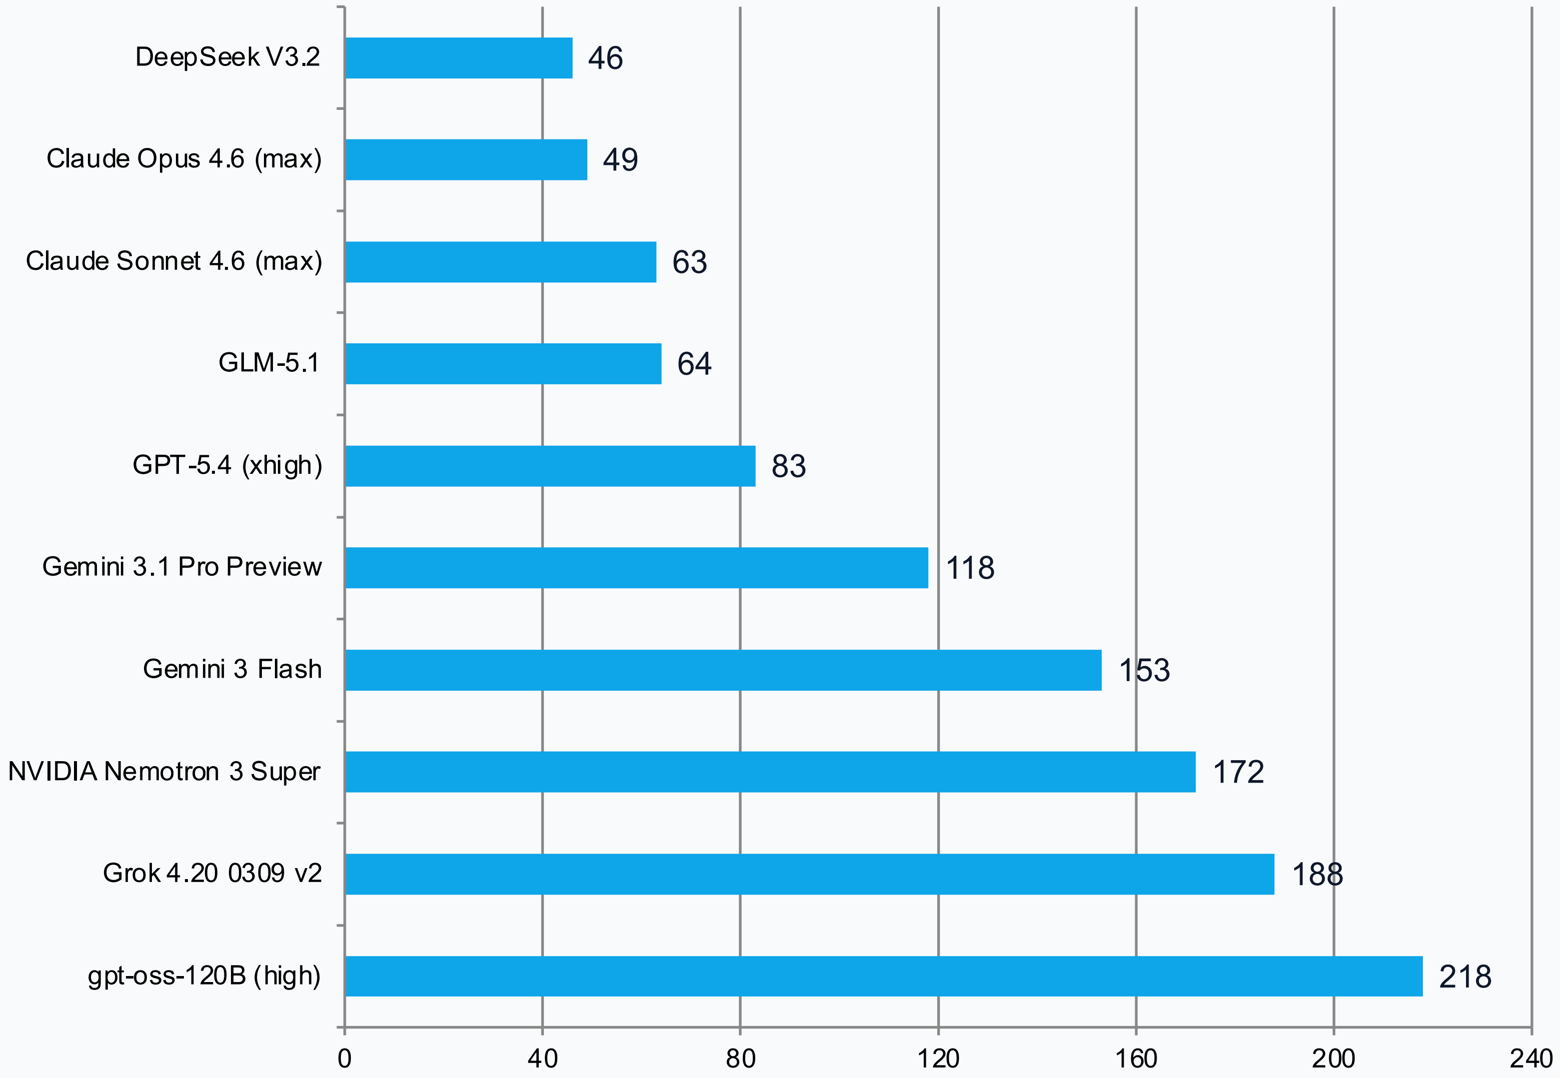

Throughput highlights

- gpt-oss-120B leads at 218 tokens/s
- 4x Claude Opus 4.6 speed
- Frontier intelligence leaders are mid-table on speed

PM / Eng takeaway

Use fast open-weight models for draft generation and low-latency UX. Escalate only complex prompts to frontier models.

Axis 3 — Pricing: Unit Economics Across Providers

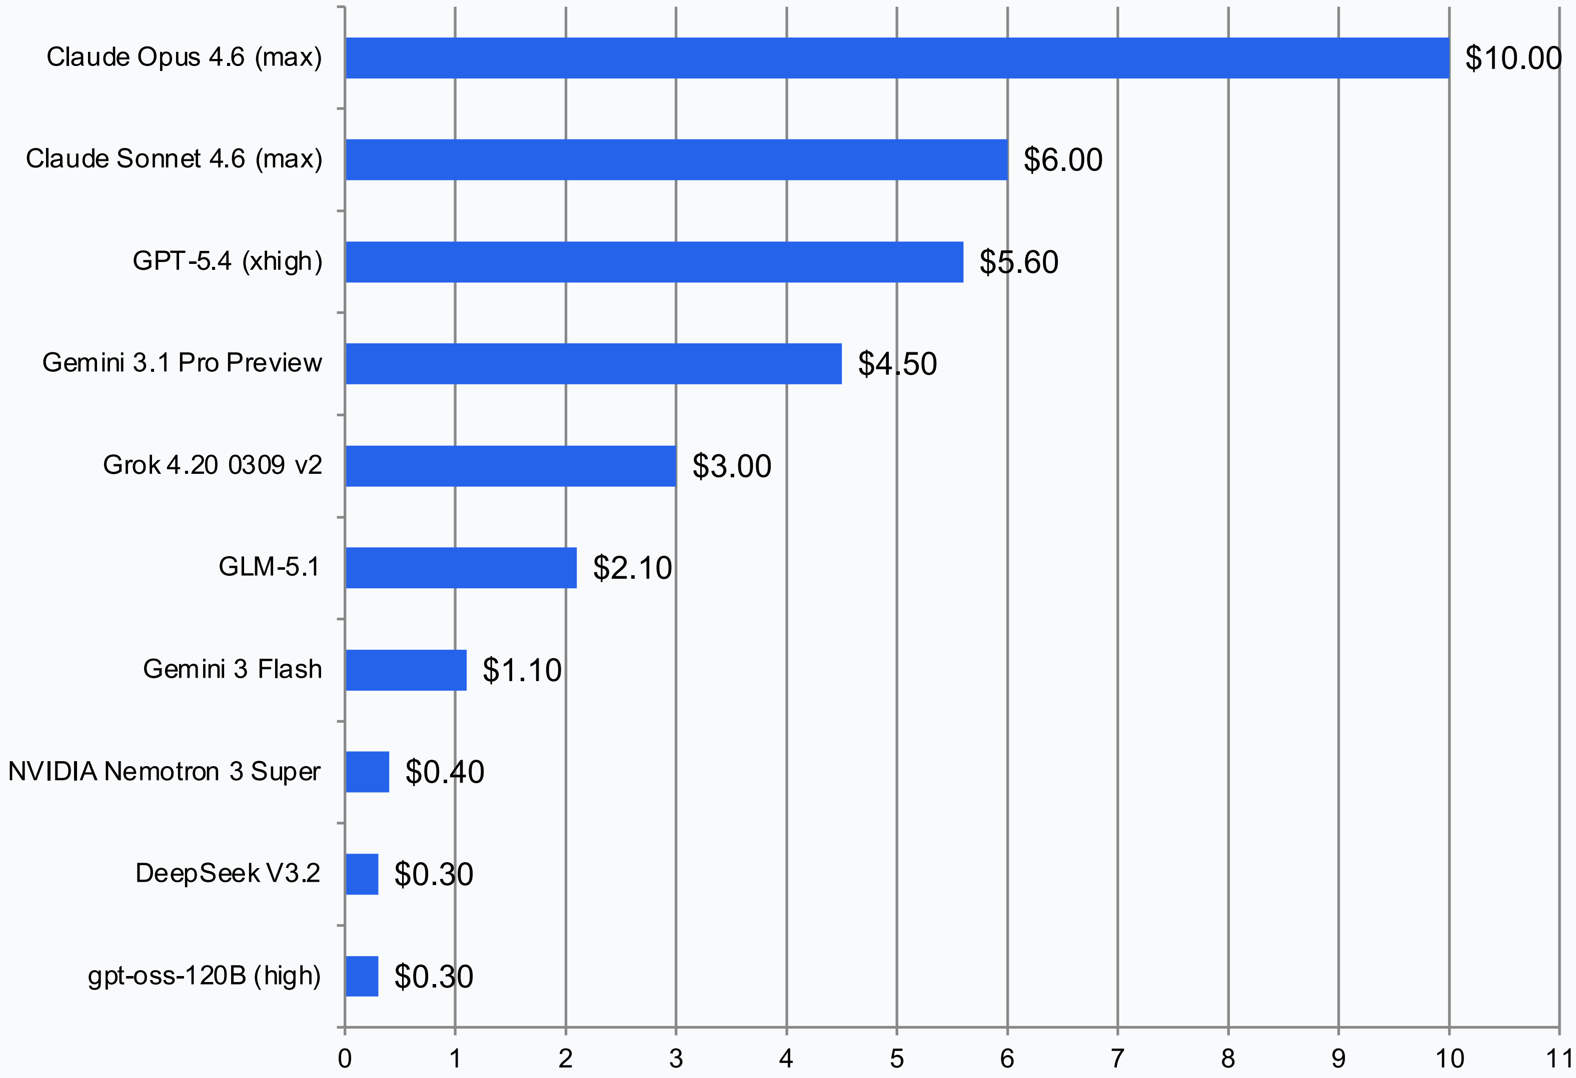

33x price gap

$0.30 → $10.00 per 1M tokens

Budget guidance

- Set spend tiers by workload criticality
- Reserve premium models for highest-value interactions
- Use low-cost models for bulk and background jobs

Trade-offs: Intelligence vs Efficiency Frontier

Intelligence Index vs output tokens consumed during evaluation

Proprietary vs Open-Weight

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price ($/1M tokens)	Gemini 3 Flash	$1.10	gpt-oss-120B (high)	$0.30

Key insight: no model leads every axis; optimize by workload objective.

Implications: Choosing the Right Model for the Job

Actionable model-selection strategy for AI product teams

1) Prioritize Intelligence

Use frontier leaders for complex reasoning, planning, and high-risk outputs where answer quality dominates latency and cost.

2) Prioritize Speed

Use high-throughput open-weight models for interactive drafting, autocomplete, and real-time UX.

3) Prioritize Cost

Route bulk generation and background jobs to lowest-cost models; enforce budget-aware model policies.

4) Hybrid Routing

Adopt tiered routing: default fast/cheap model, escalate to premium only when confidence or task complexity demands.

Bottom line: open-weight models are now strategic defaults for many workloads; proprietary frontier models remain critical for top-end capability.