


AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro Preview and GPT-5.4 tie at 57 on the Intelligence Index while a 33x pricing gap separates the cheapest and most expensive frontier models.




01 Intelligence

02 Speed

03 Pricing

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57



Full 11-model Artificial Analysis Intelligence Index ranking

Rank	Model	Provider	Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

Index score bars

Gemini 3.1 Pro Preview

GPT-5.4 (xhigh)

Claude Opus 4.6 (max)

Muse Spark

Claude Sonnet 4.6 (max)

GLM-5.1

Grok 4.20 0309 v2


Gemini 3 Flash

DeepSeek V3.2

NVIDIA Nemotron 3 Super

gpt-oss-120B (high)


57-to-33 spread



gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6




Output speed leaderboard (Muse Spark excluded; 10 models with speed data)





Interpretation

Open-weight and smaller models generate tokens far faster than frontier proprietary models; the two intelligence leaders sit in the middle of the speed table at 118 and 83 tokens/s.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



API price per million tokens (Muse Spark excluded; 10 models with pricing data)




Cost signal

For budget-sensitive workloads, the source brief identifies the price differential as decisive; Gemini 3 Flash is a mid-range value at \$1.10 with 153 tokens/s and 46 intelligence.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Scatter plot: Intelligence Index vs. output tokens consumed during evaluation



Proprietary vs. open-weight leaders

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

What the frontier shows

Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster in the upper region. Grok 4.20 0309 v2 and several Google models balance intelligence and token efficiency. MiMo-V2-Pro appears in the upper-right corner with very high intelligence but substantial token usage. No single model dominates intelligence, token efficiency, and cost.

Note: the scatter plot measures output tokens consumed during evaluation, not throughput in tokens/s.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Model selection should start with workload priorities rather than brand preference: maximize intelligence for hard reasoning, speed for latency-sensitive UX, and price for high-volume deployments.



If quality matters most

Start with Gemini 3.1 Pro Preview or GPT-5.4; both co-lead at 57 on the Intelligence Index.



If latency matters most

Evaluate gpt-oss-120B at 218 tokens/s; it leads the speed ranking.



If cost matters most

Benchmark gpt-oss-120B and DeepSeek V3.2 at \$0.30 per 1M tokens.



If openness matters

Test GLM-5.1: 51 intelligence, within 10% of proprietary leaders, at \$2.10/M tokens.

Bottom line: No model wins on all three axes; treat provider choice as a workload-specific portfolio decision.