



AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while
a 33x pricing gap separates the cheapest and most expensive
frontier models.

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57




gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6



From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Proprietary vs. Open-Weight Snapshot

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Two-way tie at the top:

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index.

Speed and intelligence are inversely correlated:

The fastest model (gpt-oss-120B at 218 tokens/s) scores only 33 on intelligence, while the smartest models operate at 49-118 tokens/s.

Massive pricing spread:

A 33x cost gap separates the cheapest (\$0.30/1M) from the most expensive (\$10.00/1M).

Open-weight models are closing the gap:

GLM-5.1 at 51 intelligence is within 10% of proprietary leaders, but cheaper (\$2.10/1M).

No single model wins on all three axes:

Choosing requires explicit trade-offs among intelligence, speed, and cost based on use case.