

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33× pricing gap separates the cheapest and most expensive frontier models.

Data Source: Artificial Analysis • April 2026

Audience: AI Engineering Leads & Product Managers Evaluating LLM Providers


Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57



Google & OpenAI share the lead at 57. Open-weight GLM-5.1 reaches 51 — within 10% of the frontier. Spread: 57 to 33.

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4× Faster Than Claude Opus 4.6


Rank	Model	Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46



Key Insight: Speed and intelligence are inversely correlated — the fastest model (218 tok/s) scores only 33 on intelligence, while the smartest models operate at 49–118 tok/s.


From \$0.30 to \$10 per Million Tokens: A 33× Pricing Gap Across Models

Rank	Model	\$/1M Tokens
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00



Budget-sensitive workloads: gpt-oss-120B and DeepSeek V3.2 at \$0.30/1M tokens are 33× cheaper than Claude Opus 4.6 at \$10.00. Gemini 3 Flash offers compelling mid-range value at \$1.10 with 153 tok/s and intelligence score of 46.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Intelligence Index vs. Output Tokens Consumed During Evaluation (not throughput)

Proprietary vs. Open-Weight Leaders

Metric	Proprietary	Score	Open-Weight	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120B	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B	\$0.30

- Gemini 3.1 Pro Preview & DeepSeek V3.2 cluster in the upper region — strong intelligence scores
- Grok 4.20 and Google models sit in the middle-right — balancing intelligence & token efficiency
- Open-weight models dominate speed (+43%) and price (73% cheaper) vs. proprietary leaders
- Every model requires explicit trade-offs — selection must be workload-driven

The scatter plot measures output tokens consumed during evaluation (efficiency), not throughput in tokens/s. Models in the upper-left quadrant achieve high intelligence with fewer tokens — the "most attractive" operating point.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Two-way tie at the top

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index, making them co-leaders as of April 2026.

Speed ↔ intelligence trade-off

The fastest model (gpt-oss-120B at 218 tok/s) scores only 33 on intelligence, while the smartest models operate at 49–118 tok/s.

33× pricing spread

From \$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2) to \$10.00/1M tokens (Claude Opus 4.6) — cost is a decisive factor for budget-sensitive workloads.

Open-weight models closing the gap

GLM-5.1 at 51 intelligence is within 10% of proprietary leaders while costing \$2.10/1M tokens. Open-weight models lead on speed and price.

Actionable Guidance for Model Selection

Max Intelligence

Gemini 3.1 Pro Preview or GPT-5.4 (xhigh) — score 57, best for complex reasoning tasks

Max Throughput

gpt-oss-120B — 218 tok/s, ideal for high-volume, latency-sensitive pipelines

Min Cost

gpt-oss-120B or DeepSeek V3.2 — \$0.30/1M tokens for budget-constrained workloads

Best Value

Gemini 3 Flash — \$1.10/1M tokens, 153 tok/s, intelligence 46; strong all-around option

Open-Weight Pick

GLM-5.1 — intelligence 51 at \$2.10/1M tokens; best open-weight intelligence

No single model dominates all three dimensions. Match your model to your workload: prioritize intelligence, speed, or cost based on your specific requirements.