INDEPENDENT LLM BENCHMARK REPORT

AI Model Landscape April 2026

Benchmarking intelligence, speed, and cost across 11 frontier models—so your team can make data-driven provider decisions.

57
Co-leading Intelligence Score (Gemini 3.1 Pro & GPT-5.4)

218
Fastest: tokens/s (gpt-oss-120B)

33×
Pricing gap ($0.30 → $10/1M tokens)

MODELS EVALUATED	PROVIDERS	EVALUATION AXES	SOURCE
11	8 Organizations	Intelligence · Speed · Price	artificialanalysis.ai

INTELLIGENCE INDEX

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead at 57

#	MODEL	PROVIDER	SCORE
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

Two-Way Tie at the Frontier

Google Gemini 3.1 Pro Preview and OpenAI GPT-5.4 (xhigh) both score **57**, the highest on record. A genuine dead heat between the two leading labs.

Open-Weight Models in the Mix

GLM-5.1 (Z AI, open-weight) scores **51**—just 6 points off the co-leaders. Grok 4.20 reaches 49. The proprietary moat is thinning.

57 → 33 Spread Across 11 Models

The gap between the top (57) and bottom (gpt-oss-120B at **33**) represents a 24-point range—significant for complex reasoning tasks.

8 Orgs, 3 Continents

US (OpenAI, Anthropic, xAI, NVIDIA, Meta), China (DeepSeek, Z AI), Europe—a genuinely global frontier.

OUTPUT SPEED

gpt-oss-120B Tops at 218 Tokens/s — 4× Faster Than Claude Opus 4.6

Open-weight models dominate throughput

The top 3 fastest models—gpt-oss-120B (**218** t/s), Grok 4.20 (**188** t/s), NVIDIA Nemotron (**172** t/s)—are all open-weight or efficiency-optimised.

Intelligence leaders sit in the middle

The top-intelligence models, Gemini 3.1 Pro and GPT-5.4, run at only **118** and **83** t/s respectively—fast enough for most workloads but not real-time streaming.

Inverse relationship with intelligence

Speed and intelligence trade off sharply. gpt-oss-120B (218 t/s) scores only 33 on intelligence. Teams optimising for latency must accept capability constraints.

Gemini 3 Flash: the pragmatic midpoint

153 t/s with an intelligence score of 46—the fastest model to crack the 40+ tier, and only $1.10/1M tokens.

API PRICING

From $0.30 to $10 per Million Tokens: A 33× Pricing Gap

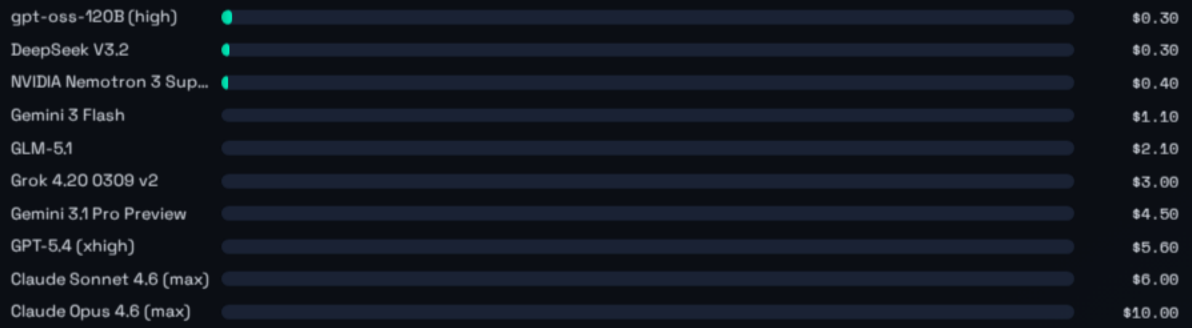

33×

Cost gap from cheapest (gpt-oss-120B, DeepSeek V3.2 at $0.30) to most expensive (Claude Opus 4.6 at $10.00) per 1M tokens

Open-weight price leaders

gpt-oss-120B and DeepSeek V3.2 both price at $0.30/1M tokens—the lowest on record—while NVIDIA Nemotron 3 Super follows at $0.40.

Best mid-range value: Gemini 3 Flash

At $1.10/1M tokens, Gemini 3 Flash pairs solid speed (153 t/s) with an intelligence score of 46—a compelling balance for most production workloads.

Premium tier: Anthropic & OpenAI

Claude Opus 4.6 ($10.00), Claude Sonnet 4.6 ($6.00), and GPT-5.4 ($5.60) price at a significant premium—justified only for the highest-stakes reasoning tasks.

TRADE-OFF ANALYSIS

No Model Wins on All Three Axes: Navigating the Frontier

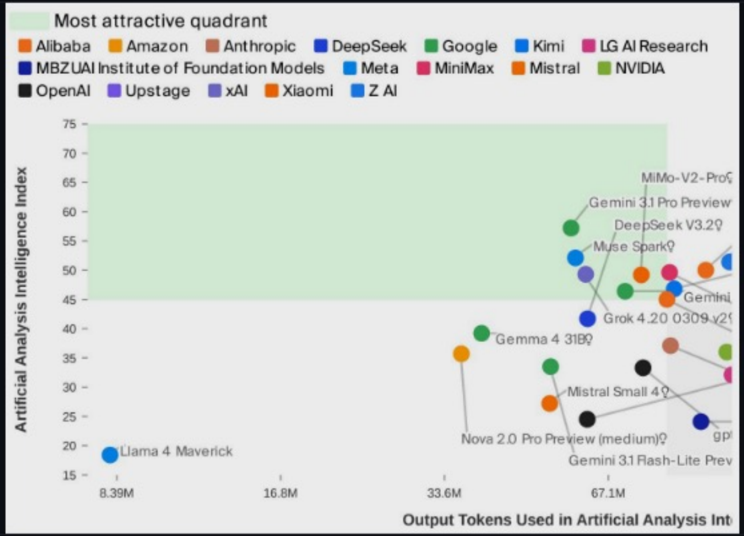

Intelligence vs. output tokens consumed during evaluation. Upper-left = high capability, efficient token use. Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster near the top; MiMo-V2-Pro achieves high intelligence at high token cost.

METRIC	PROPRIETARY LEADER	SCORE	OPEN-WEIGHT LEADER	SCORE
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (t/s)	Gemini 3 Flash	153	gpt-oss-120B	218
Price ($/1M)	Gemini 3 Flash	$1.10	gpt-oss-120B	$0.30

No single provider dominates all three dimensions. The intelligence leaders (Google, OpenAI) sacrifice speed and price. The speed and cost leaders (open-weight models) score 10–24 points below the intelligence frontier. Selection must be workload-specific.

Open weights are closing the gap fast. GLM-5.1 sits at 51 intelligence—within 10% of the co-leaders—while costing **$2.10/1M tokens** vs. $4.50–$10.00 for comparable proprietary models.

Speed ↔ Intelligence trade-off is real. Fastest (218 t/s) scores 33 on intelligence. Smartest (57) runs at 83–118 t/s. No model achieves both extremes simultaneously.

KEY FINDINGS & RECOMMENDATIONS

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at $2.10/M Tokens

⚡ Two-Way Tie at the Frontier

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score **57** on the Artificial Analysis Intelligence Index—a genuine dead heat as of April 2026.

↔ Speed & Intelligence Are Inversely Correlated

The fastest model (gpt-oss-120B, **218 t/s**) scores only 33 on intelligence. The smartest models operate at 49–118 t/s. No model dominates both.

$ Massive 33× Pricing Spread

$0.30/1M (gpt-oss-120B, DeepSeek V3.2) to **$10.00/1M** (Claude Opus 4.6). Budget-sensitive teams have strong cost-efficient options without sacrificing mid-tier capability.

↑ Open-Weight Models Are Closing the Gap

GLM-5.1 at **51 intelligence** is within 10% of the proprietary leaders at $2.10/1M tokens. gpt-oss-120B leads on speed and price. The open-weight tier is production-viable.