

# AI Model Landscape

## April 2026

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Artificial Analysis

Independent Benchmarks

11 Models · 3 Axes

## Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57



Two-way tie at the top · 24-point spread from leader to last place · Open-weight models hold mid-tier positions


# **gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6**



Speed and intelligence are inversely correlated · Top intelligence models sit mid-table on throughput · Muse Spark excluded (no speed data)


PRICING

# From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



33x cost gap from cheapest to most expensive · Gemini 3 Flash offers compelling mid-range value at \$1.10 · Muse Spark excluded (no pricing data)

# No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



## Proprietary vs. Open-Weight Leaders

| METRIC        | PROPRIETARY             | OPEN-WEIGHT           |
|---------------|-------------------------|-----------------------|
| Intelligence  | Gemini 3.1 Pro · 57     | GLM-5.1 · 51          |
| Speed (tok/s) | Gemini 3 Flash · 153    | gpt-oss-120B · 218    |
| Price (\$/1M) | Gemini 3 Flash · \$1.10 | gpt-oss-120B · \$0.30 |

Open-weight models lead on speed and price; proprietary models still edge intelligence by 6 points. Every model demands a trade-off.

# Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

- **Two-way tie at the top:** Gemini 3.1 Pro Preview and GPT-5.4 both score 57 on the Intelligence Index, co-leading as of April 2026.
- **Speed and intelligence are inversely correlated:** The fastest model (gpt-oss-120B at 218 tokens/s) scores only 33 on intelligence, while the smartest models run at 49-118 tokens/s.
- **Massive pricing spread:** A 33x cost gap separates the cheapest (\$0.30/1M tokens) from the most expensive (\$10.00/1M tokens).
- **No single model wins on all three axes:** Choosing a model requires explicit trade-offs among intelligence, speed, and cost depending on the use case.

## Recommendation

Map your workload to the right axis: latency-sensitive pipelines should prioritize speed (gpt-oss-120B, Grok 4.20), reasoning-heavy tasks should optimize for intelligence (Gemini 3.1 Pro, GPT-5.4), and budget-constrained deployments should leverage open-weight models (GLM-5.1, DeepSeek V3.2) that now deliver within 10% of proprietary intelligence at a fraction of the cost.