AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33x pricing gap separates the cheapest and most expensive frontier models.

SCOPE

Cross-provider comparison across three evaluation axes: intelligence, output speed, and pricing.

GOAL

Equip decision-makers with a clear, data-driven comparison so teams can select the right model for their workload.

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Intelligence Index (11 models)

COMPOSITE CAPABILITY SCORE • HIGHER IS BETTER

TOP = 57 | SPREAD: 57 → 33 | TIE AT #1

Rank	Model	Provider	Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

VISUAL RANKING

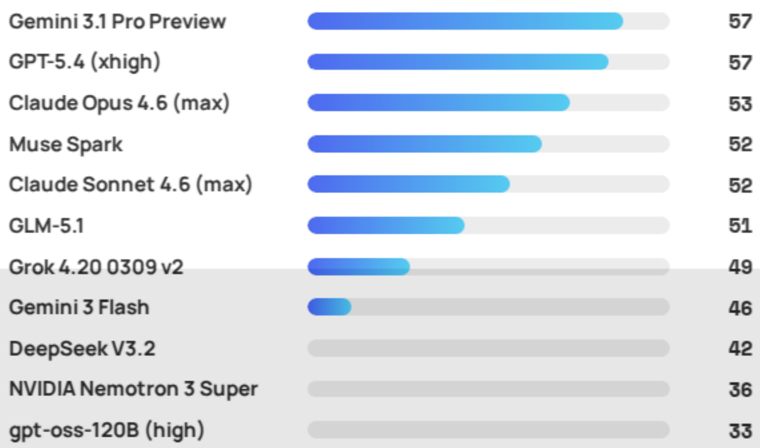

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Output Tokens per Second (10 models)

INFERENCE THROUGHPUT • MUSE SPARK EXCLUDED

FASTEST: 218 TOK/S

SLOWEST: 46 TOK/S

Rank	Model	Tok/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46

NOTES

PATTERN

Open-weight and smaller models dominate throughput.

TRADE-OFF

Speed and intelligence are inversely correlated in these April 2026 results.

CONTEXT

Intelligence co-leaders (57) sit mid-table at 118 and 83 tok/s.

From $0.30 to $10 per Million Tokens: A 33x Pricing Gap Across Models

Scatter plot metric: output tokens consumed during evaluation (not tokens/s).

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at $2.10/M Tokens

April 2026 results show a tight frontier at the top—and a widening field of strong, economical options.

Key findings (from benchmark)

WHAT THE NUMBERS IMPLY

- FRONTIER Two-way intelligence tie at 57 (Gemini 3.1 Pro Preview, GPT-5.4).
- SPEED gpt-oss-120B leads throughput at 218 tok/s while scoring 33 on intelligence.
- COST Cheapest models at $0.30/$1M are 33x below the $10.00/$1M ceiling.
- OPEN GLM-5.1 reaches 51 intelligence at $2.10/$1M—within 10% of the leaders.

Selection playbook

USE-CASE DRIVEN

- MAX IQ Prioritize the 57-tier when correctness and complex reasoning dominate.
- LATENCY For interactive UX, optimize for tokens/s (e.g., 150–218 tok/s models).
- BUDGET For high-volume tasks, start with $0.30–$1.10/$1M options, then tier up selectively.
- PORTFOLIO No single winner: standardize routing by workload class (reasoning vs latency vs cost).

Data source: Artificial Analysis (captured Apr 10, 2026).