

FRONTIER MODEL BENCHMARK • 11 MODELS

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Artificial Analysis benchmarks show a split frontier: Gemini 3.1 Pro Preview and GPT-5.4 tie on intelligence, while open-weight and smaller models pressure incumbents on speed and price.

INTELLIGENCE LEADERS

Gemini 3.1 Pro Preview + GPT-5.4

57

PRICING SPREAD

Cheapest to most expensive

33x

AXIS 1 • INTELLIGENCE INDEX

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

The full 11-model ranking spans 57 to 33. Proprietary frontier models still lead, but GLM-5.1 reaches 51 and open-weight challengers occupy the mid-tier.

57

Co-leader score

24

Point spread to rank 11

8+

Organizations represented



AXIS 2 • OUTPUT SPEED

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Speed is measured in output tokens per second. Muse Spark is excluded because speed data is not published in the brief.

218

gpt-oss-120B tokens/s

118

Gemini 3.1 Pro Preview

83

GPT-5.4 (xhigh)




From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

Per-million-token pricing varies enormously. Budget-sensitive workloads can see decisive impact from the \$0.30 options, while Claude Opus 4.6 sits at \$10.00.



RANK	MODEL	USD / 1M TOKENS
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00



Scatter measures output tokens consumed during evaluation, not throughput in tokens/s.

TRADE-OFF FRONTIER

No Model Wins on All Three Axes: Navigating the Intelligence- Efficiency-Cost Frontier

High intelligence, token efficiency, and low cost do not converge on a single winner. Teams need workload-priority routing, not one default model.

METRIC	PROPRIETARY LEADER	SCORE	OPEN-WEIGHT LEADER	SCORE
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

SELECTION GUIDANCE

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Open-weight models are close enough to change architecture decisions: use explicit trade-offs by workload rather than assuming the most intelligent model is always best.

01

Maximize reasoning quality

Start with Gemini 3.1 Pro Preview or GPT-5.4 when intelligence is the gating constraint.

02

Prioritize latency

Evaluate gpt-oss-120B, Grok 4.20, NVIDIA Nemotron, and Gemini 3 Flash for speed-sensitive paths.

03

Control unit economics

Route budget-sensitive volume to \$0.30–\$1.10 per million token options where quality permits.

04

Adopt portfolio routing

No single model wins intelligence, speed, and cost; match models to task priority and fallback needs.