

ARTIFICIAL ANALYSIS · APRIL 2026

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

An independent evaluation of frontier LLMs across intelligence, throughput, and pricing — equipping engineering leads and product managers with data-driven selection criteria.

Source: [artificialanalysis.ai](#) · Captured April 10, 2026



Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Full 11-model ranking by Artificial Analysis Intelligence Index (composite of reasoning, knowledge, and instruction-following benchmarks).



gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6

Output tokens per second leaderboard. Open-weight and smaller models generate tokens far faster than frontier proprietary models.

RANK	MODEL	TOKENS/S
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153


The two intelligence leaders (Gemini 3.1 Pro at 118, GPT-5.4 at 83) sit in the middle of the speed table — highlighting the inverse correlation between raw intelligence and throughput.



From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Scatter plot of Intelligence Index vs. output tokens consumed during evaluation, plus proprietary vs. open-weight leader comparison.



Models achieving high scores while consuming fewer tokens represent better engineering. Gemini 3.1 Pro and DeepSeek V3.2 cluster in the upper region; MiMo-V2-Pro shows very high intelligence with substantial token usage.

METRIC	PROPRIETARY		OPEN-WEIGHT	
Intelligence	Gemini 3.1 Pro	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B	\$0.30

Proprietary models still lead on raw intelligence, but the margin is thin. GLM-5.1 is within 10% of the co-leaders while open-weight models dominate speed and price.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

- **Two-way tie at the top** — Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index, making them co-leaders as of April 2026.
- **Speed and intelligence are inversely correlated** — The fastest model (gpt-oss-120B at 218 tokens/s) scores only 33 on intelligence, while the smartest models operate at 49–118 tokens/s.
- **Massive pricing spread** — A 33x cost gap separates the cheapest (\$0.30/1M for gpt-oss-120B and DeepSeek V3.2) from the most expensive (\$10.00/1M for Claude Opus 4.6).

Recommendation: Map your workload priorities first — then select the model that dominates your critical axis while meeting minimum thresholds on the others.