


AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33x pricing gap separates the cheapest and most expensive frontier models.

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57


Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33



Google and OpenAI are neck-and-neck at the frontier. Open-weight models like GLM-5.1 (51) are narrowing the gap to within 10% of the leaders.

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6


Rank	Model	Output Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46



Speed and intelligence are inversely correlated: the fastest model (218 tok/s) scores only 33 on intelligence, while the two smartest models operate at 83–118 tok/s.


From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

Rank	Model	USD / 1M Tokens
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00



A 33x cost gap separates the cheapest (\$0.30) from the most expensive (\$10.00). Gemini 3 Flash at \$1.10 offers compelling mid-range value with solid speed (153 tok/s) and intelligence (46).

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Proprietary vs. Open-Weight Leaders

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tok)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Key Observations

- Open-weight models lead on speed and price
- Proprietary models retain an intelligence edge, but the gap is narrowing (57 vs. 51)
- Models in the upper-left of the scatter plot achieve high intelligence with fewer output tokens consumed
- Every model requires trade-offs among intelligence, token efficiency, and cost

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key Findings

- Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 both score 57 on the Intelligence Index
- Speed and intelligence are inversely correlated: the fastest model (218 tok/s) scores only 33 on intelligence
- 33x pricing spread: from \$0.30 to \$10.00 per million tokens across 10 models
- Open-weight models closing the gap: GLM-5.1 at 51 intelligence is within 10% of proprietary leaders, at just \$2.10/M tokens
- No single model wins all three axes: teams must choose based on workload priorities

Actionable Guidance

Evaluate models on the specific axis most critical to your workload. For latency-sensitive applications, prioritize speed. For cost-sensitive deployments, consider open-weight alternatives. For maximum capability, budget for proprietary frontier models.