


AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Independent Analysis — Artificial Analysis, April 2026


Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33x pricing gap separates cheapest and most expensive.

For: AI Engineering Leads & Product Managers

April 2026



Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57



57

Co-leaders:

Gemini 3.1 Pro Preview & GPT-5.4


6 pts

behind: Claude Opus 4.6 (max) at 53

Open-weight GLM-5.1 reaches 51

33–57 range across 11 models

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6




218 tok/s
gpt-oss-120B (high) — fastest model

4x+
faster than Claude Opus 4.6 (49 tok/s)

Inverse trend
Smarter models run slower

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



33x

pricing gap

\$0.30/1M


gpt-oss-120B & DeepSeek V3.2

\$10.00/1M

Claude Opus 4.6 (max)

Gemini 3 Flash: \$1.10 — best mid-range value

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Intelligence Index vs. Output Tokens Consumed (evaluation efficiency)

Proprietary vs. Open-Weight

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price	Gemini 3 Flash	\$1.00	gpt-oss-120B (high)	\$0.30

No single model leads across all three dimensions. Open-weight models like GLM-5.1 (51 intelligence) and gpt-oss-120B (\$0.30/1M) are closing the gap.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

1

Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index — co-leaders as of April 2026.

2

Speed and intelligence are inversely correlated: The fastest model (gpt-oss-120B at 218 tok/s) scores only 33; the smartest run at 49–118 tok/s.

3

Massive pricing spread: A 33x cost gap separates the cheapest (\$0.30/1M for gpt-oss-120B and DeepSeek V3.2) from the most expensive (\$10.00/1M for Claude Opus 4.6).

4

Open-weight models closing the gap: GLM-5.1 at 51 intelligence is within 10% of proprietary leaders, at just \$2.10/1M tokens.

5

No single model wins on all three axes: Model selection requires explicit trade-offs among intelligence, speed, and cost for each use case.

Model Selection Guide

Latency-Critical Apps

gpt-oss-120B or Grok 4.20 — max throughput for chatbots & autocomplete

Budget Bulk Processing

gpt-oss-120B or DeepSeek V3.2 — \$0.30/1M tokens

Highest Reasoning Quality

Gemini 3.1 Pro Preview or GPT-5.4 — score 57 intelligence

Balanced Value

Gemini 3 Flash — \$1.10/1M, 153 tok/s, 46 intelligence score

Source: Artificial Analysis, April 2026