




AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost
Across 11 Frontier Models

Source: Artificial Analysis | Data captured April 10, 2026



Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57




Key Insights

57 co-leaders

- Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 both score 57
- 57-to-33 spread across 11 models shows wide capability range
- Open-weight GLM-5.1 scores 51, within 10% of the leaders
- 8 distinct organizations across US, China, and Europe reflect global competition

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6




Key Insights

218

tokens/s peak speed

- gpt-oss-120B is 4.5x faster than Claude Opus 4.6 (49 tokens/s)
- Intelligence leaders sit mid-table: Gemini 3.1 Pro at 118, GPT-5.4 at 83 tokens/s
- Speed and intelligence are inversely correlated: fastest model scores only 33 on intelligence
- Muse Spark excluded (no speed data published)

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models




Source: Artificial Analysis API Pricing Data, April 2026

Key Insights

33X
pricing gap

- Cheapest: gpt-oss-120B and DeepSeek V3.2 at \$0.30/1M tokens
- Most expensive: Claude Opus 4.6 at \$10.00/1M tokens
- Gemini 3 Flash offers compelling mid-range value at \$1.10 with 153 tokens/s speed and intelligence of 46
- Muse Spark excluded (no pricing data published)

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Intelligence Index vs. Output Tokens Consumed During Evaluation

Proprietary vs. Open-Weight Leaders

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tok)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Every model requires explicit trade-offs among intelligence, speed, and cost. Choose based on your workload priorities.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

57

Co-leaders on Intelligence Index
Gemini 3.1 Pro & GPT-5.4

33x

Pricing gap from \$0.30
to \$10.00 per 1M tokens

10%

GLM-5.1 within 10% of leaders
at just \$2.10/1M tokens

Actionable Recommendations for Model Selection

Maximize intelligence

Gemini 3.1 Pro Preview or GPT-5.4 (score 57). Accept higher cost and moderate speed.

Optimize for latency

gpt-oss-120B (218 tokens/s) or Grok 4.20 (188 tokens/s). Trade off raw intelligence.

Control cost at scale

gpt-oss-120B or DeepSeek V3.2 at \$0.30/1M tokens. Best for high-volume, tolerance-tolerant workloads.

Best balanced value

GLM-5.1: intelligence 51, \$2.10/1M tokens, 64 tokens/s. Open-weight with near-frontier capability.