

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Data-forward comparison of frontier models across three evaluation axes — intelligence, speed, and pricing — to support provider selection decisions.

Competitive landscape (April 2026):

Frontier proprietary labs co-lead on intelligence, while open-weight models lead on throughput and cost — implying deliberate trade-offs rather than a single dominant choice.

Source: Artificial Analysis (captured Apr 10, 2026)

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index (composite benchmark score)

Full 11-model ranking

Rank	Model	Provider	Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

Artificial Analysis (captured Apr 10, 2026)

What to notice

- Two-way tie at 57 (Google + OpenAI)
- 57 → 33 spread across the leaderboard
- Open-weight models sit mid-tier (gap narrowing)

57
57
53
52
52
51
49
46
42
36
33





Figure: Intelligence Index ranking (Artificial Analysis)

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Inference speed measured as output tokens per second (tokens/s)



Leaderboard (10 models)

Rank	Model	Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46

Pattern: the highest-intelligence leaders sit mid-table on throughput — reflecting an intelligence vs. speed trade-off

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

API pricing per 1M tokens (published pricing; Muse Spark excluded)

Pricing table (10 models)

Rank	Model	\$/1M
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60

9 Cheapest: \$0.30 (gpt-oss-120B, DeepSeek V3.2) vs. Most expensive: \$10.00 (Claude Opus 4.6) → 33x gap
10 Claude Opus 4.6 (max) \$10.00





Figure: Price per 1M tokens (Artificial Analysis)

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Scatter plot shows Intelligence Index vs. output tokens consumed during evaluation (token efficiency, not throughput).



Artificial Analysis (captured Apr 10, 2026)

Proprietary vs. open-weight leaders

Metric	Proprietary	Open-weight
Intelligence	Gemini 3.1 Pro Preview (57)	GLM-5.1 (51)
Speed (tokens/s)	Gemini 3 Flash (153)	gpt-oss-120B (high) (218)
Price (\$/1M tokens)	Gemini 3 Flash (\$1.10)	gpt-oss-120B (high) (\$0.30)

takeaway: selection is workload driven — optimize for the axis that dominates your latency, quality, and budget constraints.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Actionable guidance for selecting models by workload priorities

If quality dominates

- Start with the intelligence co-leaders (57): Gemini 3.1 Pro Preview or GPT-5.4.
- Expect mid-range throughput (118 / 83 tokens/s) and higher cost (\$4.50 / \$5.60 per 1M).

If latency dominates

- Prioritize high-throughput open-weight models (e.g., gpt-oss-120B at 218 tokens/s).
- Validate task fit: highest speed may come with lower intelligence (33).

If budget dominates

- Cheapest options: \$0.30 per 1M tokens (gpt-oss-120B, DeepSeek V3.2).
- Use mid-tier open weights to balance: GLM-5.1 scores 51 at \$2.10 per 1M.

Decision rule-of-thumb

No single model dominates intelligence, speed, and price. Use a two-stage approach: (1) pick the axis your workload cannot compromise (quality / latency / budget), then (2) shortlist 2–3 candidates and validate with your own evals and cost model.