

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models


Artificial Analysis snapshot (captured April 10, 2026)

Audience: AI engineering leads & product managers

Key premise: no single model dominates intelligence, throughput, and cost simultaneously.

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Output speed leaderboard (Muse Spark excluded: no speed data)




Model	Output tokens/s
gpt-oss-120B (high)	218
Claude Opus 4.6 (max)	49
GPT-5.4 (xhigh)	83
Gemini 3.1 Pro Preview	118

Throughput skews toward open-weight / smaller models; intelligence leaders sit mid-pack on speed.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

Published API pricing (10 models; Muse Spark excluded)




Cheapest: gpt-oss-120B and DeepSeek V3.2 at \$0.30/M
Most expensive: Claude Opus 4.6 at \$10.00/M
Price spread: 33x

Mid-range value callout:
Gemini 3 Flash = \$1.10/M, 153 tokens/s, intelligence 46.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence vs output tokens consumed during evaluation (not throughput in tokens/s)



Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Interpretation: frontier performance forms a trade-off curve; choose models by workload priority, not headline rank alone.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Actionable guidance for AI engineering and product teams

- No universal winner: prioritize intelligence, speed, or cost by use case.
- For highest capability tasks, Gemini 3.1 Pro Preview and GPT-5.4 co-lead at 57.
- For latency-sensitive products, favor high-throughput options (e.g., gpt-oss-120B at 218 tokens/s).
- For budget-sensitive scale, exploit the 33x pricing spread and route workloads accordingly.
- Open-weight progress is material: GLM-5.1 (51) narrows the gap while remaining cost-competitive.
- Recommended operating model: multi-model portfolio with task-based routing and continuous benchmark refresh.