

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33x pricing gap separates the cheapest and most expensive frontier models.

Evaluation axes

- Intelligence
- Speed
- Pricing

INTELLIGENCE LEADERS

57

Gemini 3.1 Pro Preview + GPT-5.4

PRICE SPREAD

33x

\$0.30 to \$10.00 per 1M tokens

FASTEST OUTPUT

218 t/s

gpt-oss-120B (high)

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57



Source: Artificial Analysis, captured April 10, 2026. Higher Intelligence Index scores indicate stronger overall capability.

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Output tokens per second leaderboard



Fastest model

218 t/s

gpt-oss-120B (high)

Claude Opus 4.6

49 t/s

Fastest model is more than 4x faster

Intelligence leaders

118 / 83

Gemini 3.1 Pro Preview / GPT-5.4

Muse Spark excluded: no output-speed data in source brief.

Pattern: open-weight and smaller models generate tokens far faster than frontier proprietary models, while the two intelligence leaders sit mid-table on throughput.

Metric: inference speed measured in output tokens per second.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

Cheapest published price

\$0.30

gpt-oss-120B + DeepSeek V3.2

Most expensive

\$10.00

Claude Opus 4.6 per 1M tokens

Mid-range value

\$1.10


Gemini 3 Flash: 153 t/s + Index 46

USD per 1M tokens



Budget implication: the cheapest options are over 33x less expensive than Claude Opus 4.6, making price a decisive factor for high-volume deployments.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Proprietary vs. open-weight leaders

Metric	Proprietary leader	Score	Open-weight leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Decision signal

Raw intelligence, output efficiency, throughput, and cost produce different winners. Use-case priorities should drive provider selection.

Scatter plot measures Intelligence Index vs. output tokens consumed during evaluation — not throughput in tokens/s.

Open-weight models are closing the gap: GLM-5.1 trails the co-leaders by 6 intelligence points, while gpt-oss-120B leads speed and price.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens



Max intelligence

Choose Gemini 3.1 Pro Preview or GPT-5.4 when top composite capability is the primary requirement.



Latency-sensitive apps

Consider gpt-oss-120B, Grok 4.20, NVIDIA Nemotron 3 Super, or Gemini 3 Flash for high output tokens/s.



Budget-sensitive scale

Start with gpt-oss-120B or DeepSeek V3.2 at \$0.30 per 1M tokens; compare quality thresholds before routing.



Balanced open-weight path

Evaluate GLM-5.1: Intelligence Index 51 at \$2.10 per 1M tokens shows the gap is narrowing.

Operating model recommendation

Do not select a single default provider by benchmark rank alone. Build a routing matrix by workload: intelligence-critical reasoning, low-latency UX, and high-volume cost control.

Key finding: no single model wins on all three axes; explicit trade-offs among intelligence, speed, and cost are required.