


AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Equip decision-makers with a clear, data-driven comparison of top LLMs across intelligence, throughput, and cost so they can select the right model for their use case.


Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33


gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6



From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 both score 57.

Speed and intelligence are inversely correlated: The fastest model scores only 33 on intelligence, while the smartest models operate at 49-118 tokens/s.

Massive pricing spread: A 33x cost gap separates the cheapest (\$0.30/1M) from the most expensive (\$10.00/1M).

Open-weight models are closing the gap: GLM-5.1 at 51 intelligence is within 10% of proprietary leaders, while being significantly cheaper.

No single model wins on all three axes: Choosing a model requires explicit trade-offs depending on the use case.