AI Model Landscape April 2
Benchmarking Intelligence
and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a
separates the cheapest and most expensive frontier models.

11

Frontier Models Benchmarked

33x

Pricing Gap (Cheapest vs. Most Expensive)

Source: Artificial Analysis | April 2026

Audience: AI Engineering Leads & Product Managers

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index — composite score across reasoning, knowledge, and instruction-following

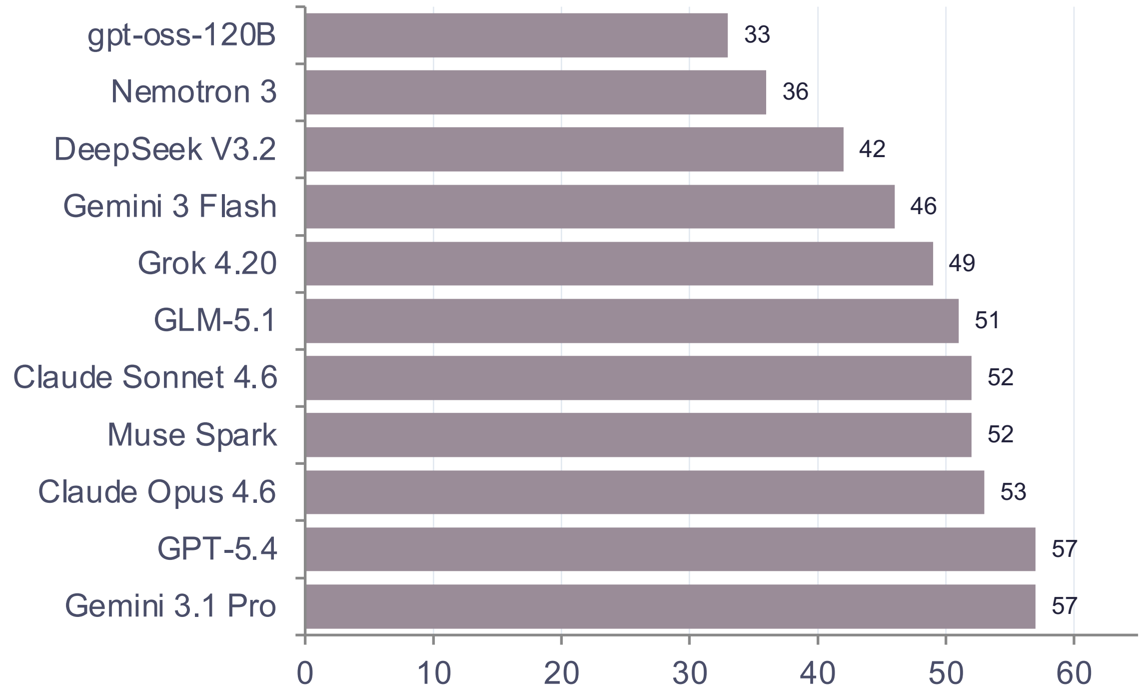

Key Insights

- **Two-way tie at the top**

Gemini 3.1 Pro Preview and GPT-5.4 both score 57, making them co-leaders.

- **Open-weight gap narrows**

GLM-5.1 at 51 is within 10% of the leaders. Muse Spark enters strongly at 52.

- **57-to-33 spread**

A 24-point range separates the co-leaders from gpt-oss-120B.

Source: Artificial Analysis, April 2026

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Inference throughput measured in output tokens per second — critical for latency-sensitive applications

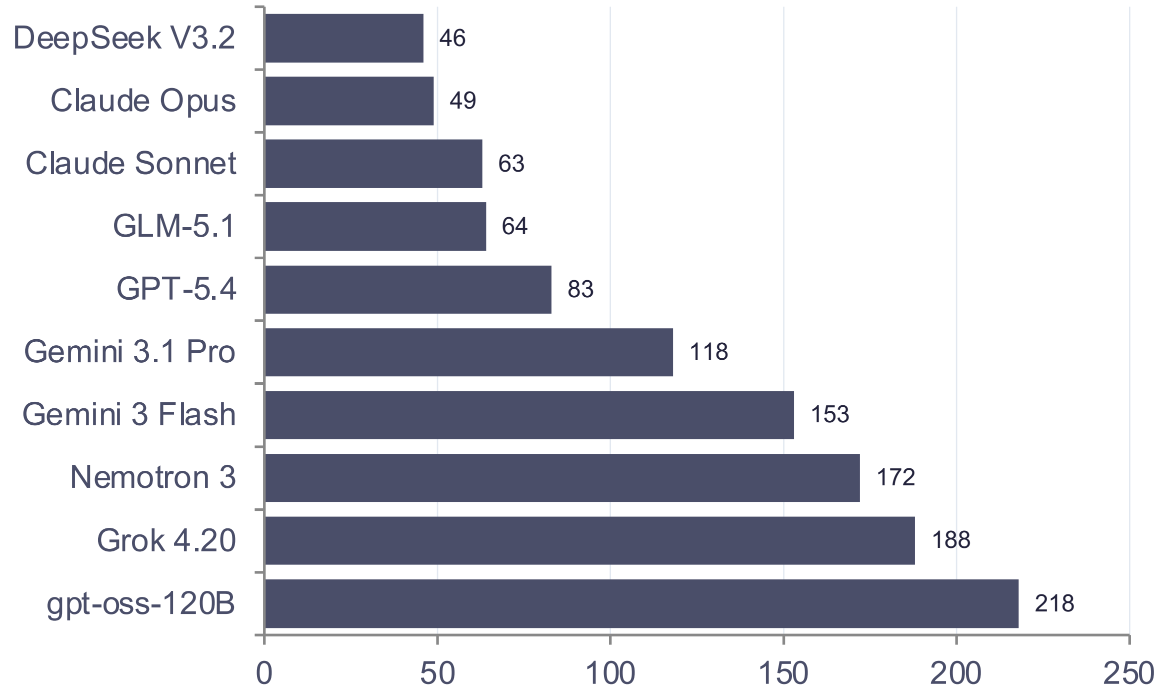

Key Insights

- **Open-weight dominance**

gpt-oss-120B outputs 218 tokens/s — 4x+ Claude Opus and nearly 5x DeepSeek V3.2.

- **Inverse intelligence correlation**

The two intelligence leaders (Gemini 3.1 Pro at 118, GPT-5.4 at 83) sit mid-table.

- **Latency trade-off**

For real-time apps, raw speed may outweigh marginal intelligence gains.

Source: Artificial Analysis, April 2026

From $0.30 to $10 per Million Tokens: A 33x Pricing Gap Across Models

API pricing per million tokens — reflecting compute requirements, margin strategies, and competitive positioning

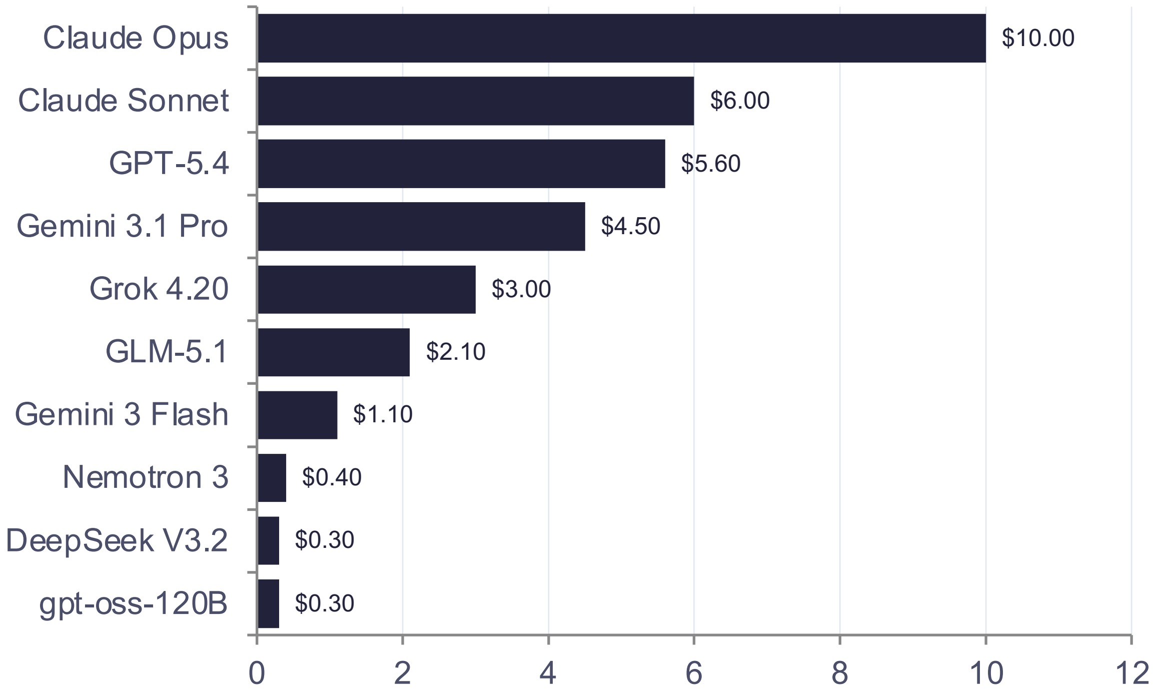

Key Insights

- **33x cost spread**

Cheapest ($0.30) vs. most expensive ($10.00) per million tokens.

- **Budget sweet spot**

Gemini 3 Flash at $1.10 pairs solid speed (153 t/s) with reasonable intelligence (46).

- **Open-weight price advantage**

gpt-oss-120B and DeepSeek V3.2 both at $0.30 make high-volume inference extremely affordable.

Source: Artificial Analysis, April 2026

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Scatter plot measures output tokens consumed during evaluation (not throughput). Lower = more token-efficient.

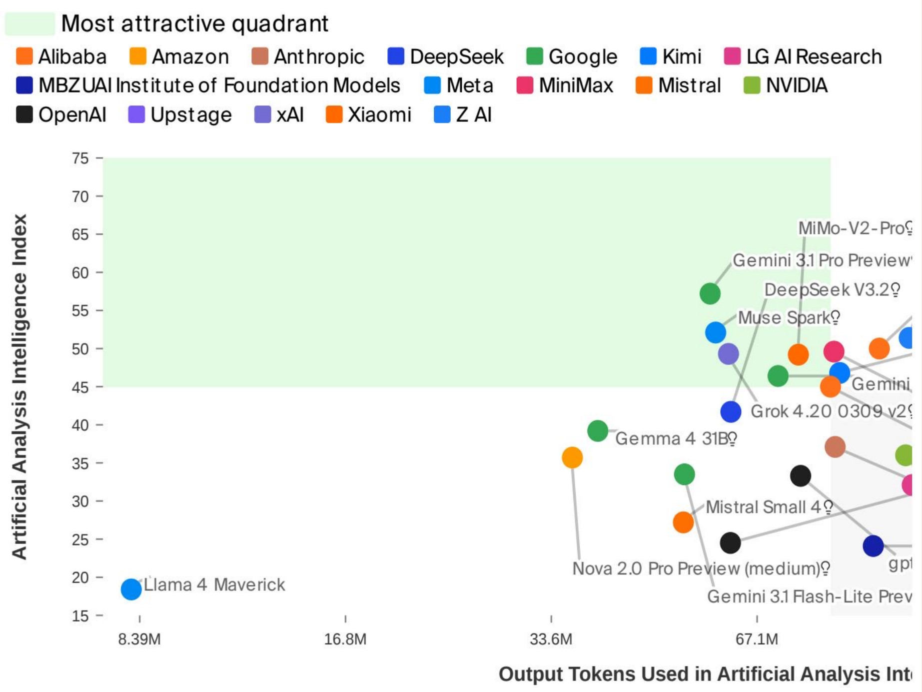

Source: Artificial Analysis, April 2026

Proprietary vs. Open-Weight Leaders

Metric	Proprietary Leader	Score	Open-Weight
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B
Price ($/1M tokens)	Gemini 3 Flash	$1.10	gpt-oss-120B

Scatter Observations

- Gemini 3.1 Pro Preview & DeepSeek V3.2 cluster upper region — strong intelligence.
- Grok 4.20 & Google models balance intelligence with token efficiency.
- MiMo-V2-Pro shows very high intelligence but substantial token usage.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at $2.10/M Tokens

Actionable guidance for model selection based on workload priorities

Key Findings

- **Two-way tie at the top**

Gemini 3.1 Pro Preview and GPT-5.4 both score 57 on the Intelligence Index.

- **Speed and intelligence are inversely correlated**

Fastest model (gpt-oss-120B at 218 t/s) scores 33; smartest models run 49–118 t/s.

- **Massive pricing spread**

33x gap from $0.30 to $10.00 per million tokens.

- **Open-weight models are closing the gap**

GLM-5.1 at 51 intelligence is within 10% of proprietary leaders at $2.10/M tokens.

- **No single model wins on all three axes**

Selection requires explicit trade-offs among intelligence, speed, and cost.

Recommendations

- **Maximize intelligence**

Choose Gemini 3.1 Pro Preview or GPT-5.4 for complex reasoning tasks.

- **Minimize latency**

Deploy gpt-oss-120B or Grok 4.20 for real-time, high-throughput workloads.

- **Optimize cost**

Use gpt-oss-120B or DeepSeek V3.2 at $0.30/M tokens for budget-sensitive scale.

- **Best balance**

Gemini 3 Flash offers compelling mid-range value: speed + intelligence + $1.10/M.

- **Future-proof with open weights**

GLM-5.1 proves open-weight models are approaching parity — evaluate for self-hosted deployments.

Source: Artificial Analysis, April 2026