AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33× pricing gap separates the cheapest and most expensive frontier models.

Data Source: Artificial Analysis • Captured April 10, 2026 • For AI engineering leads & product managers

11 Frontier Models Benchmarked

3 Evaluation Axes: Intelligence · Speed · Cost

8+ Organizations US · China · Europe

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index — composite score across reasoning, knowledge, and instruction-following benchmarks

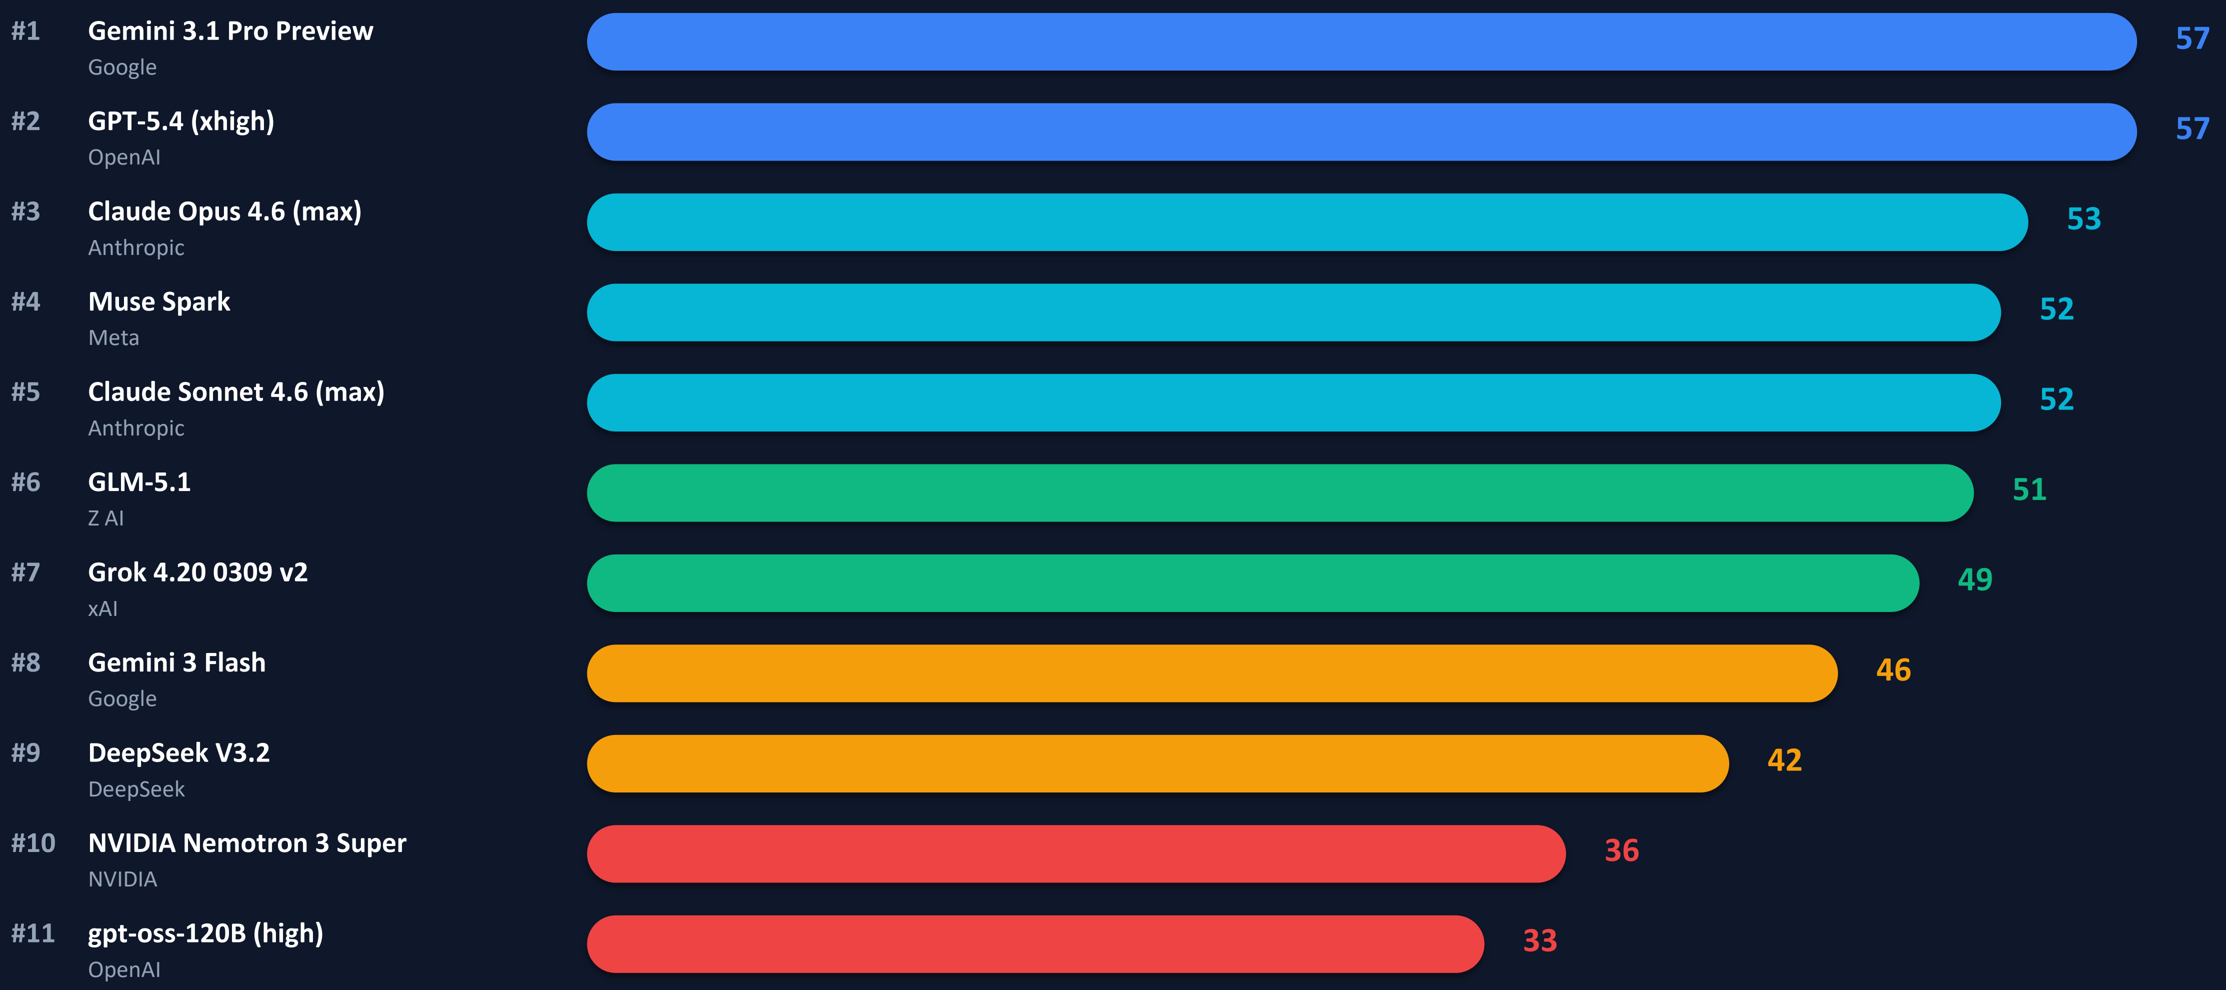

⚡ Two-way tie at the top • Open-weight GLM-5.1 (51) within 10% of leaders • 24-point spread across the full field (57 → 33)

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4× Faster Than Claude Opus 4.6

Output tokens per second — critical for latency-sensitive applications (Muse Spark excluded: no speed data published)

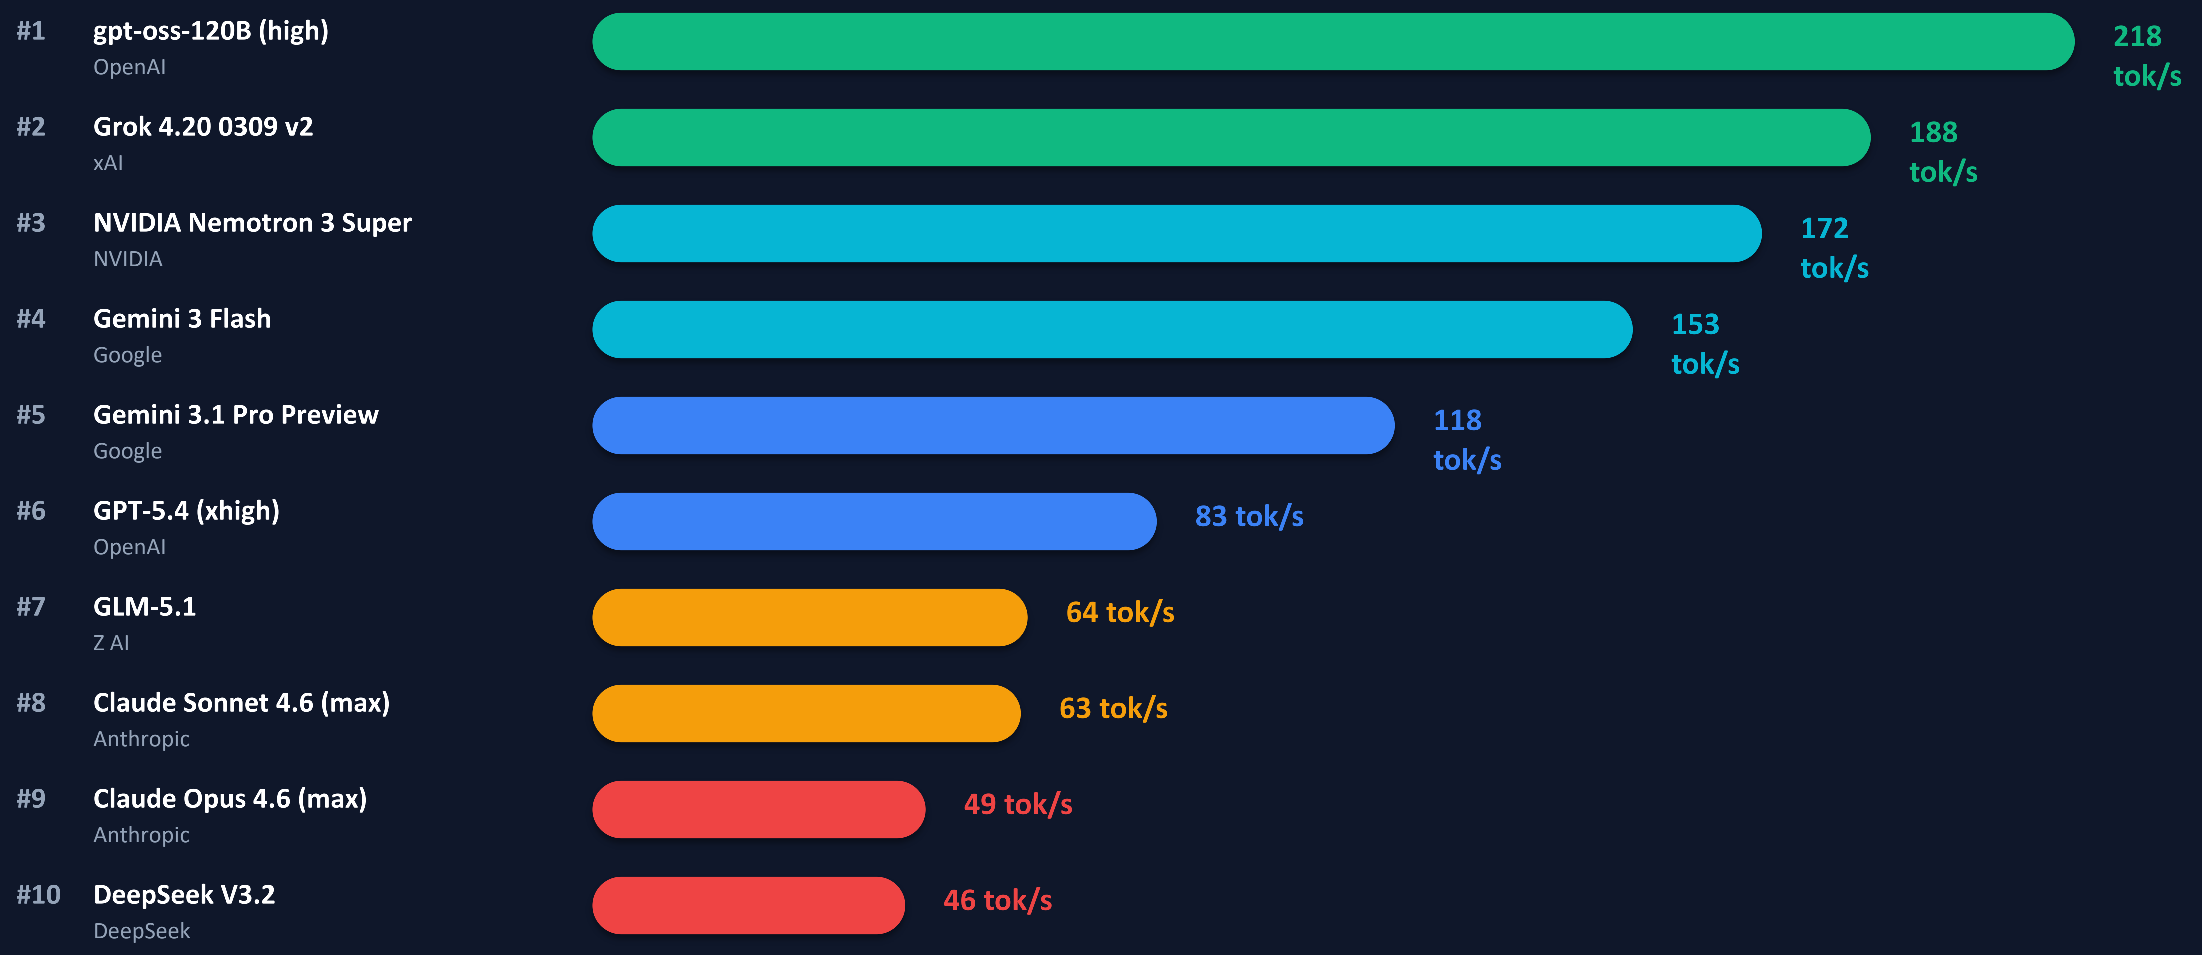

⚡ Speed ↔ Intelligence inverse correlation • Intelligence co-leaders sit mid-table: Gemini 3.1 Pro (118 tok/s), GPT-5.4 (83 tok/s) • Open-weight gpt-oss-120B is fastest AND cheapest

From $0.30 to $10 per Million Tokens: A 33× Pricing Gap Across Models

USD per 1M output tokens — Muse Spark excluded (no published pricing data)

⚡ 33× cost gap from cheapest to most expensive • Gemini 3 Flash at $1.10 offers compelling mid-range value (153 tok/s, intelligence 46) • Budget-sensitive? gpt-oss-120B & DeepSeek V3.2 at $0.30

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence vs. Output Tokens Consumed During Evaluation

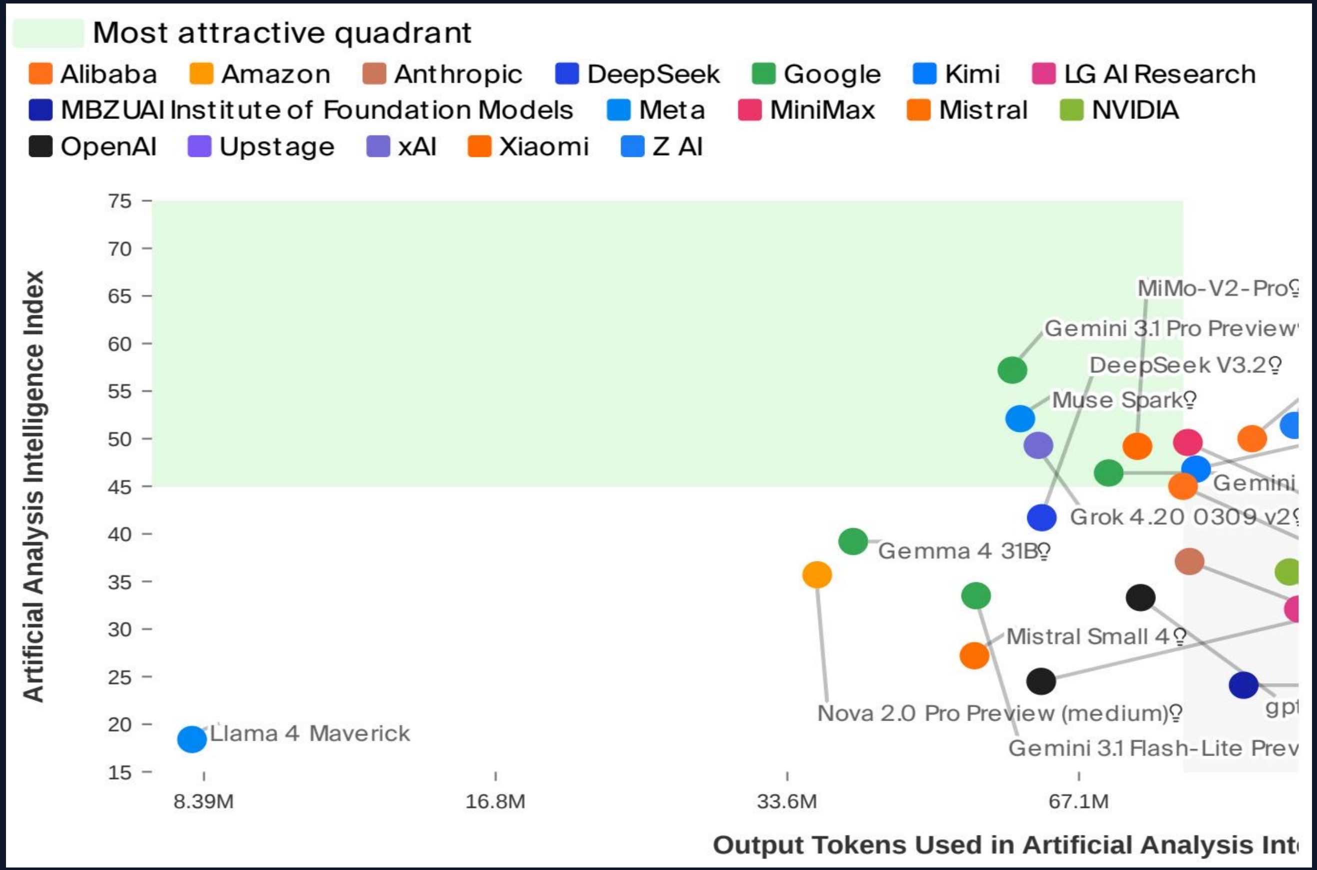

Upper-left quadrant = high intelligence, low token usage (most efficient)

Proprietary vs. Open-Weight Leaders

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price ($/1M tok)	Gemini 3 Flash	$1.10	gpt-oss-120B (high)	$0.30

Scatter-plot insights:

- Gemini 3.1 Pro Preview & DeepSeek V3.2 cluster in the high-intelligence region
- Grok 4.20 & Google models balance intelligence and token efficiency mid-right
- MiMo-V2-Pro: very high intelligence but substantial token usage (upper-right)
- Llama 4 Maverick: low intelligence, low tokens — neither efficient nor capable

The trade-off is structural:

- No single model dominates intelligence + speed + cost simultaneously
- Open-weight models already lead on speed & price; intelligence gap is <10%

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at $2.10/M Tokens

Two-way tie at the top

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index — the co-leaders as of April 2026.

Speed & intelligence are inversely correlated

The fastest model (gpt-oss-120B at 218 tok/s) scores only 33 on intelligence; the smartest models run at 49–118 tok/s.

33× pricing spread

From $0.30/M tokens (gpt-oss-120B, DeepSeek V3.2) to $10.00/M (Claude Opus 4.6) — cost is a decisive factor for budget-sensitive workloads.

Open-weight models are closing the gap

GLM-5.1 at 51 intelligence is within 10% of proprietary leaders, at $2.10/M tokens. Open-weight already leads on speed & price.

ACTIONABLE GUIDANCE

→ Max Intelligence

Gemini 3.1 Pro Preview or GPT-5.4 when accuracy is paramount. Budget ~$4.50–$5.60/M tokens.

→ Max Throughput

gpt-oss-120B (218 tok/s) or Grok 4.20 (188 tok/s) for real-time & streaming use cases.

→ Cost-Optimized

gpt-oss-120B or DeepSeek V3.2 at $0.30/M tokens. Gemini 3 Flash at $1.10 for mid-range value.

→ Best Open-Weight

GLM-5.1: 51 intelligence, $2.10/M — strong all-rounder. Self-host for data sovereignty.

Bottom line: No single model wins on all three axes. Match your model to your workload — intelligence for quality-critical tasks, speed for real-time apps, cost for scale. Open-weight options are now viable for most production workloads.