

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro Preview and GPT-5.4 tie at 57 on the Intelligence Index while a 33× pricing gap separates the cheapest and most expensive frontier models.

INTELLIGENCE

Composite reasoning
& knowledge score

SPEED

Output tokens
per second

COST

Price per
million tokens

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index · April 2026 · Higher = stronger reasoning, knowledge & instruction-following

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57 ★
2	GPT-5.4 (xhigh)	OpenAI	57 ★
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51 ●
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42 ●
10	NVIDIA Nemotron 3 Super	NVIDIA	36 ●
11	gpt-oss-120B (high)	OpenAI	33 ●

Source: Artificial Analysis, April 2026 | artificialanalysis.ai

57 / 57

Two-way tie: Gemini 3.1 Pro Preview & GPT-5.4 (xhigh)

51 pts

GLM-5.1 (open-weight) just 6 pts behind co-leaders

33 → 57

24-point spread across the full 11-model field

● = open-weight model ★ = co-leaders

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4× Faster Than Claude Opus 4.6

Output tokens per second · Higher = faster inference · Muse Spark excluded (no published speed data)

Rank	Model	Tokens / s	Intelligence
1	gpt-oss-120B (high) ●	218	33
2	Grok 4.20 0309 v2 ●	188	49
3	NVIDIA Nemotron 3 Super ●	172	36
4	Gemini 3 Flash	153	46
5	Gemini 3.1 Pro Preview	118	57
6	GPT-5.4 (xhigh)	83	57
7	GLM-5.1 ●	64	51
8	Claude Sonnet 4.6 (max)	63	52
9	Claude Opus 4.6 (max)	49	53
10	DeepSeek V3.2 ●	46	42

4.7× speed gap

gpt-oss-120B (218 tok/s)
vs Claude Opus 4.6 (49 tok/s)
— same vendor, opposite ends

Inverse pattern

The 2 intelligence leaders
(Gemini 3.1 Pro, GPT-5.4) rank
5th and 6th on speed.

● = open-weight model

From \$0.30 to \$10 per Million Tokens: A 33× Pricing Gap Across Models

USD per 1 M tokens (blended) · Muse Spark excluded (no published price)

Rank	Model	\$/1M Tokens	Intelligence
1	gpt-oss-120B (high) ●	\$0.30	33
2	DeepSeek V3.2 ●	\$0.30	42
3	NVIDIA Nemotron 3 Super ●	\$0.40	36
4	Gemini 3 Flash	\$1.10	46
5	GLM-5.1 ●	\$2.10	51
6	Grok 4.20 0309 v2 ●	\$3.00	49
7	Gemini 3.1 Pro Preview	\$4.50	57
8	GPT-5.4 (xhigh)	\$5.60	57
9	Claude Sonnet 4.6 (max)	\$6.00	52
10	Claude Opus 4.6 (max)	\$10.00	53

33× cost spread

\$0.30 (gpt-oss-120B / DeepSeek)
vs \$10.00 (Claude Opus 4.6)

Best mid-range value

Gemini 3 Flash: \$1.10/1M
153 tok/s · Intelligence 46
— solid on all three axes


Open-weight advantage

Top-4 cheapest models
are all open-weight

● = open-weight model

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence Index vs. output tokens consumed during evaluation · Most-attractive quadrant = upper-left



Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price \$/1M	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Open-weights lead on speed and cost — gpt-oss-120B beats every proprietary model on both axes.

Intelligence gap is narrowing — GLM-5.1 (open) scores 51 vs 57 for the proprietary co-leaders.

No single model dominates — every selection requires explicit trade-offs across the three axes.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key findings & actionable model-selection guidance

Two-way tie at the top

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) share rank #1 with a score of 57 on the Artificial Analysis Intelligence Index.

Speed and intelligence are inversely correlated

The fastest model (gpt-oss-120B, 218 tok/s) scores only 33 on intelligence; the smartest models run at 49–118 tok/s.

Massive 33× pricing spread

\$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2) vs. \$10.00/1M tokens (Claude Opus 4.6) — cost must be a first-class selection criterion.

Open-weight models are closing the gap

GLM-5.1 scores 51 — within 10% of the proprietary leaders — at \$2.10/1M tokens, demonstrating viable alternatives for most workloads.

No single model wins on all three axes

Every deployment decision requires explicit prioritization of intelligence, speed, and cost based on the target use case.

SELECTION GUIDANCE BY WORKLOAD

Max intelligence (reasoning, coding, research)

→ Gemini 3.1 Pro Preview or GPT-5.4 (xhigh)

High throughput / real-time UX

→ gpt-oss-120B or Grok 4.20 0309 v2 (open-weight)

Budget-sensitive / high-volume

→ gpt-oss-120B (\$0.30) or DeepSeek V3.2 (\$0.30)

Balanced performance at mid-cost

→ Gemini 3 Flash (\$1.10, Intel 46, 153 tok/s)

Open-weight near-frontier intelligence

→ GLM-5.1 (Intel 51, \$2.10/1M tokens)