

AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models


Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33x pricing gap separates the cheapest and most expensive frontier models.

Source: Artificial Analysis | Data captured April 10, 2026 | artificialanalysis.ai

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index — composite score across reasoning, knowledge, and instruction-following benchmarks

Score



Key Takeaways

Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 both score 57.

24-point spread from leader (57) to last place (33).

Anthropic's Claude Opus 4.6 trails by 4 points.


Meta's Muse Spark enters at 52 — strong debut.

Open-weight GLM-5.1 at 51 is within 10% of the proprietary leaders.

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6

Output tokens per second — critical for latency-sensitive applications (Muse Spark excluded, no speed data published)

Score



Key Takeaways

gpt-oss-120B outputs 218 tokens/s — 4.4x faster than Claude Opus 4.6.

Open-weight models dominate the top 3 speed positions.


Intelligence leaders sit mid-table: Gemini 3.1 Pro at 118 tok/s, GPT-5.4 at 83 tok/s.

Speed and intelligence are inversely correlated: the fastest model scores only 33 on intelligence.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

API pricing per million tokens (Muse Spark excluded, no pricing data published)

Score



Key Takeaways

33x gap: \$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2) vs. \$10.00 (Claude Opus 4.6).


Budget tier (under \$1): gpt-oss-120B, DeepSeek V3.2, NVIDIA Nemotron 3 Super, Gemini 3 Flash.

Gemini 3 Flash offers a compelling mid-range value at \$1.10 with 153 tok/s and intelligence of 46.

Claude Opus 4.6 is 33x more expensive than the cheapest option — justified only for tasks requiring maximum intelligence.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence vs. output tokens consumed during evaluation — models in the upper-left quadrant deliver the best intelligence-per-token efficiency



Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Scatter Plot Observations

- Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster in the upper region — strong intelligence with efficient token usage.
- Grok 4.20 and several Google models balance intelligence and token efficiency in the middle-right.
- Llama 4 Maverick occupies the low-intelligence, low-token region.
- MiMo-V2-Pro shows very high intelligence but with substantial token usage.

Every model requires explicit trade-offs among intelligence, speed, and cost. The right choice depends on your workload: latency-critical apps should prioritize speed, complex reasoning tasks need peak intelligence, and high-volume pipelines must optimize for cost.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key findings and actionable recommendations for model selection

- Two-way tie at the top**
Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index — co-leaders as of April 2026.
- Speed and intelligence are inversely correlated**
The fastest model (gpt-oss-120B at 218 tokens/s) scores only 33 on intelligence, while the smartest models operate at 49–118 tokens/s.
- Massive pricing spread**
A 33x cost gap separates the cheapest (\$0.30/1M tokens for gpt-oss-120B and DeepSeek V3.2) from the most expensive (\$10.00 for Claude Opus 4.6).
- Open-weight models are closing the gap**
GLM-5.1 at 51 intelligence is within 10% of the proprietary leaders, while also being significantly cheaper at \$2.10/1M tokens.
- No single model wins on all three axes**
Choosing a model requires explicit trade-offs among intelligence, speed, and cost depending on the use case.

Max Intelligence

Gemini 3.1 Pro Preview or GPT-5.4

Complex reasoning, research, high-stakes decision support

Latency-Critical

gpt-oss-120B or Grok 4.20

Real-time chat, interactive UI, streaming applications

Cost-Optimized

DeepSeek V3.2 or gpt-oss-120B

High-volume batch, classification, summarization pipelines

Balanced

Gemini 3 Flash

General-purpose workloads needing moderate intelligence at good speed and price