


AI Model Landscape


April 2026

Benchmarking Intelligence, Speed, and Cost
Across 11 Frontier Models




Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

11 models ranked by the Artificial Analysis Intelligence Index — higher scores indicate stronger overall capability




gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6

10 models ranked by output tokens per second — Muse Spark excluded (no published speed data)




From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

10 models with published pricing data — Muse Spark excluded



No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence Index vs. output tokens consumed during evaluation — efficient models appear toward the lower-right



Proprietary vs. Open-Weight Summary			Weight
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1
		51	
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120
		218	
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120
		\$0.30	

Source: Artificial Analysis Intelligence Index vs. Output Tokens Consumed, April 2026

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Actionable guidance for AI engineering leads and product managers evaluating model selection

01

Two-way tie at the top

Gemini 3.1 Pro Preview and GPT-5.4 both score 57 on the Intelligence Index — the highest scorers as of April 2026.

02

Speed and intelligence are inversely correlated

gpt-oss-120B delivers 218 tokens/s but scores only 33 in intelligence. Frontier models trade throughput for capability.

03

A 33x pricing gap separates cheapest and most expensive

gpt-oss-120B and DeepSeek V3.2 cost \$0.30/1M tokens vs. Claude Opus 4.6 at \$10.00/1M tokens.

04

Open-weight models are closing the gap

GLM-5.1 scores 51 intelligence (within 10% of leaders) at just \$2.10/1M tokens — a compelling value proposition.

05

No single model wins on all three axes

Model selection requires explicit trade-offs: prioritize intelligence for complex reasoning, speed for real-time apps, cost for high-volume workloads.