

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost

Independent benchmarks across 11 frontier models (captured April 10, 2026) spanning proprietary and open-weight providers.

Competitive landscape is global: models from at least 8 organizations across the US, China, and Europe.
Three evaluation axes used throughout this deck: intelligence, output speed (tokens/s), and API price (\$/1M tokens).
Core insight: the frontier is multi-dimensional — intelligence leaders are not always fastest or cheapest.

Evaluation Axes

- Intelligence Index (composite score)
- Output speed (tokens/s)
- Price (\$ per 1M tokens)

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index (higher is better)


Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

Co-leaders at 57; leaderboard spans 57 → 33 across the 11 models.

Source: Artificial Analysis Intelligence Index (captured April 10, 2026)

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than C


Output tokens per second (Muse Spark excluded: no speed data in source)



Fastest: gpt-oss-120B at 218 tokens/s.
218 is >4× Claude Opus 4.6 (49) and nearly
5× DeepSeek V3.2 (46).
Intelligence leaders sit mid-pack on speed:
Gemini 3.1 Pro (118), GPT-5.4 (83).

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

USD per 1M tokens (Muse Spark excluded: no published pricing in source)




Cheapest: \$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2).
Most expensive: \$10.00/1M tokens (Claude Opus 4.6) → 33× gap.
Mid-range value point in source: Gemini 3 Flash at \$1.10/1M tokens.

Source: Artificial Analysis pricing (captured April 10, 2026)

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence vs output tokens consumed during evaluation (token consumption ≠ throughput)



Metric	Proprietary Leader	Score	Open-Weight
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Actionable guidance for model selection (April 2026 snapshot)

Start with workload priorities: (1) quality/robustness, (2) latency/throughput, (3) unit economics — optimize for the constraint you cannot violate.

For maximum intelligence, Gemini 3.1 Pro Preview and GPT-5.4 co-lead at 57; expect mid-tier output speed (118 / 83 tokens/s) and higher cost (\$4.50 / \$5.60 per 1M tokens).

For latency-sensitive systems, open-weight options lead throughput (gpt-oss-120B at 218 tokens/s) while trading off intelligence (33). For budget-sensitive deployments, cheapest options are \$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2) — a 33× gap vs Claude Opus 4.6 at \$10.00.

Open weights are closing the gap: GLM-5.1 reaches 51 intelligence at \$2.10/1M tokens (within 10% of the co-leaders on Intelligence Index).