

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontiers

Audience: AI engineering leads and product managers | Source: Artificial Analysis (captured April 10, 2026)

Objective: Compare 11 frontier models on three decision-critical dimensions: intelligence, output speed, and price.
Headline: Gemini 3.1 Pro and GPT-5.4 tie at 57 intelligence, while pricing spans \$0.30 to \$10.00 per 1M tokens (33x gap).
Decision context: No single provider leads across all axes; selection must be workload-specific.

Intelligence

Capability ceiling

Speed

Latency/UX responsiveness

Pricing


Unit economics at scale

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57


Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

10-model speed leaderboard (Muse Spark excluded: no speed data)
Leader: gpt-oss-120B at 218 tokens/s.
gpt-oss-120B is >4x Claude Opus 4.6 (49 tokens/s).
Intelligence co-leaders sit mid-pack on speed: Gemini 3.1 Pro (118), GPT-5.4 (83).
Pattern: higher intelligence generally comes with lower throughput.




From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



10-model pricing comparison (Muse Spark excluded: no published price)
Cheapest: gpt-oss-120B and DeepSeek V3.2 at \$0.30 per 1M tokens.
Most expensive: Claude Opus 4.6 at \$10.00 per 1M tokens.
Total spread: 33x from lowest to highest.
Mid-range value example: Gemini 3 Flash at \$1.10 with 46 intelligence a

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Scatter plot x-axis is output tokens consumed during evaluation (efficiency)
Takeaway: frontier selection is a multi-objective optimization problem.
No model simultaneously maximizes intelligence, token-efficiency, and price

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Co-leaders on intelligence: Gemini 3.1 Pro Preview and GPT-5.4 at 57.

Throughput leader: gpt-oss-120B at 218 tokens/s.

Cost floor: \$0.30/1M tokens; ceiling: \$10.00/1M tokens (33x spread).

Open weights are increasingly competitive: GLM-5.1 reaches 51 intelligence

Implication: model choice should be workload-specific, not brand-specific

For highest-quality reasoning workflows: start with Gemini 3.1 Pro or GPT-5.4

For latency-critical UX: prioritize faster models and gate to high-IQ models

For budget-sensitive scale: evaluate gpt-oss-120B, DeepSeek V3.2, and Gemini 3.1 Pro

Use routing strategies by task tier (simple, standard, complex) to balance cost and quality

Re-benchmark quarterly; rankings and economics shift rapidly.