

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33x pricing gap separates the cheapest and most expensive frontier models.

CO-LEADERS

57

Gemini 3.1 Pro Preview + GPT-5.4

PRICE SPREAD

33x

\$0.30 to \$10 per 1M tokens

FASTEST MODEL

218

output tokens/s for gpt-oss-120B

Scope: 11 frontier models benchmarked across intelligence, speed, and pricing

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index: higher scores indicate stronger reasoning, knowledge, and instruction-following capability.

TOP SCORE

57

Google + OpenAI tie

OPEN-WEIGHT HIGH

51

GLM-5.1 from Z AI

RANK RANGE

57–33

Full 11-model leaderboard

Rank	Model	Provider	Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

What matters for evaluators

- Google and OpenAI are neck-and-neck at the frontier.
- Anthropic's Claude Opus 4.6 trails by 4 points.
- Meta's Muse Spark enters at a strong 52.
- Open-weight models such as GLM-5.1 and Grok 4.20 occupy the mid-tier.

Gemini 3.1 Pro Preview



GPT-5.4



Claude Opus 4.6



Muse Spark



Claude Sonnet 4.6




GLM-5.1



gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Output speed is measured in output tokens per second; Muse Spark is excluded because speed data was not provided.



FASTEST

218

gpt-oss-120B tokens/s

CO-LEADERS

118 / 83

Gemini Pro / GPT-5.4

Rank	Model	Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46

Pattern: open-weight and smaller models generate tokens far faster than frontier proprietary models; the intelligence leaders sit in the middle of the speed table.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

API pricing per million tokens varies by compute requirements, margin strategies, and competitive positioning.

CHEAPEST

\$0.30

gpt-oss-120B + DeepSeek V3.2

HIGHEST LISTED


\$10.00

Claude Opus 4.6

GAP

33x

Cheapest vs. most expensive




Rank	Model	USD / 1M
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00

Budget-sensitive workloads: the \$0.30 models are over 33x less expensive than Claude Opus 4.6. Gemini 3 Flash pairs solid speed (153 tokens/s) with reasonable intelligence (46) at \$1.10.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

The scatter plot uses output tokens consumed during evaluation, not throughput in tokens/s.



Proprietary vs. open-weight leaders

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Scatter-plot observations

- Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster in the upper region.
- Grok 4.20 0309 v2 and several Google models balance intelligence and token efficiency.
- Llama 4 Maverick occupies the low-intelligence, low-token region.
- MiMo-V2-Pro appears in the upper-right corner with substantial token usage.

Decision implication: choose explicitly among intelligence, token efficiency, speed, and cost rather than expecting a universal winner.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Actionable model-selection guidance for AI engineering leads and product managers.

If peak intelligence matters

Start with Gemini 3.1 Pro Preview or GPT-5.4; both score 57 on the Intelligence Index.

If latency matters

Evaluate gpt-oss-120B, Grok 4.20, NVIDIA Nemotron 3 Super, and Gemini 3 Flash; they lead throughput at 218, 188, 172, and 153 tokens/s.

If unit economics matter

Prioritize gpt-oss-120B or DeepSeek V3.2 at \$0.30 per 1M tokens, with NVIDIA Nemotron 3 Super at \$0.40.

If balancing capability and control

Put GLM-5.1 on the shortlist: 51 intelligence at \$2.10 per 1M tokens.

Bottom line

No single model wins on all three axes. Open-weight models are closing the gap: GLM-5.1 is within 10% of proprietary leaders, while gpt-oss-120B leads speed and price.

Recommended operating model: define workload priorities first, then benchmark a shortlist against intelligence, latency/token consumption, and cost constraints.