AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Target Audience: AI engineering leads and product managers evaluating LLM providers

Tone: Data-forward, professional

Source: Artificial Analysis, April 2026

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Rank	Model	Output Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46

Key takeaway: Open-weight and smaller models output tokens faster than frontier proprietary models, displaying an inverse cor

From $0.30 to $10 per Million Tokens: A 33x Pricing Gap Across Models

Rank	Model	USD per 1M Tokens
1	gpt-oss-120B (high)	$0.30
2	DeepSeek V3.2	$0.30
3	NVIDIA Nemotron 3 Super	$0.40
4	Gemini 3 Flash	$1.10
5	GLM-5.1	$2.10
6	Grok 4.20 0309 v2	$3.00
7	Gemini 3.1 Pro Preview	$4.50
8	GPT-5.4 (xhigh)	$5.60
9	Claude Sonnet 4.6 (max)	$6.00
10	Claude Opus 4.6 (max)	$10.00

The cheapest options are 33x less expensive than the most expensive, a decisive factor for budget-sensitive workloads. (Muse S

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

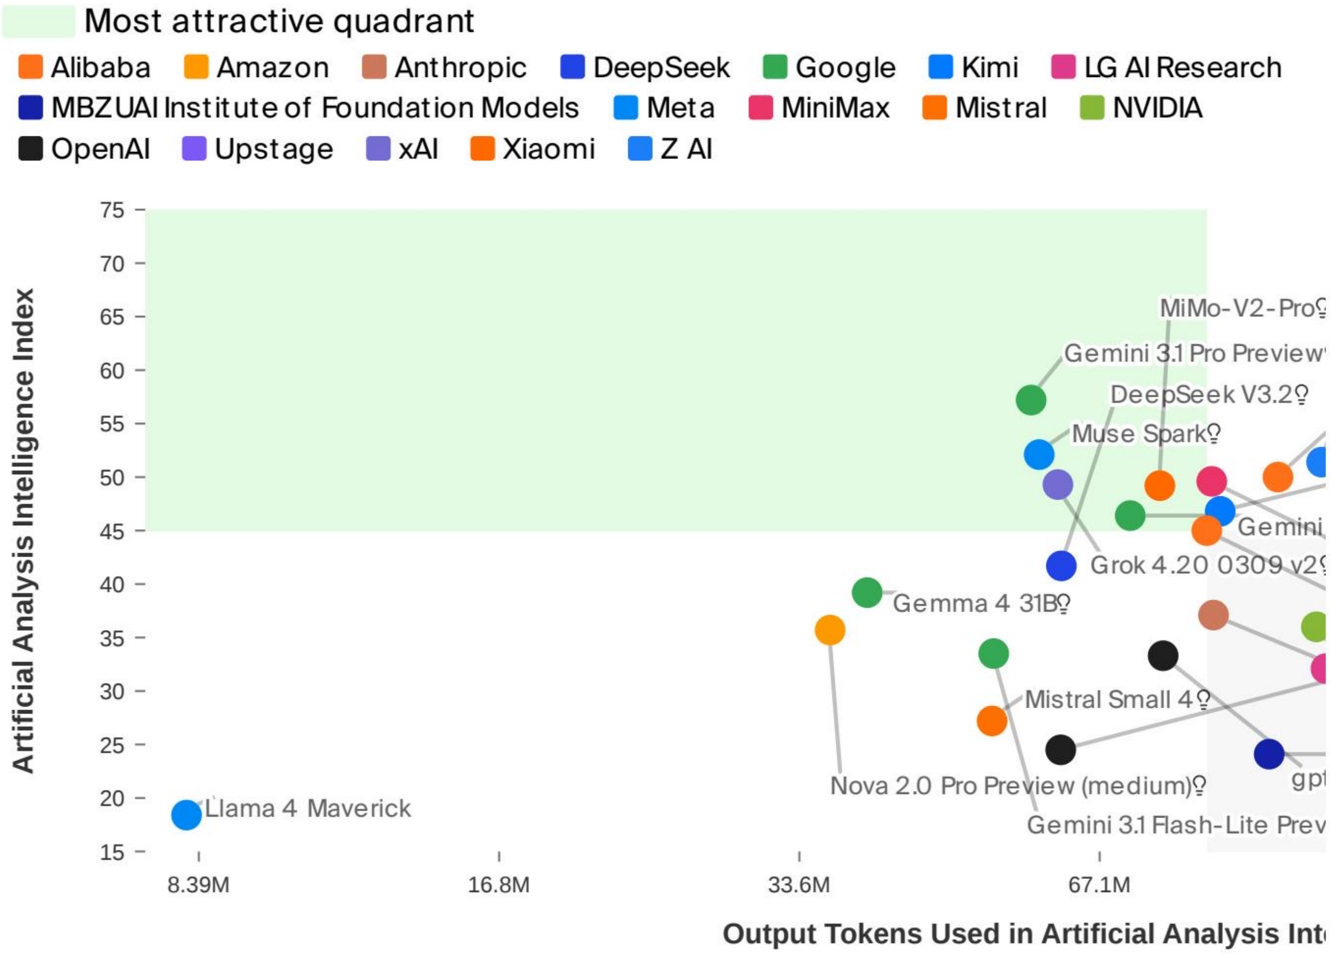

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price ($/1M tokens)	Gemini 3 Flash	$1.10	gpt-oss-120B (high)	$0.30

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at $2.10/M Tokens

- Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Artificial Analysis Intelligence Index.
- Speed and intelligence are inversely correlated: The fastest model scores 33 on intelligence, while smartest models operate at 49-118 tokens/s.
- Massive pricing spread: A 33x cost gap separates the cheapest ($0.30/1M tokens) from the most expensive ($10.00/1M tokens).
- Open-weight models are closing the gap: GLM-5.1 at 51 intelligence is within 10% of the proprietary leaders and significantly cheaper.
- No single model wins on all three axes: Explicit trade-offs among intelligence, speed, and cost must be made depending on the specific workload priority.