

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

A data-forward comparison of proprietary and open-weight LLMs | Artificial Analysis, April 2026

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at

The Artificial Analysis Intelligence Index aggregates reasoning, knowledge, and instruction-following benchmarks.

Google and OpenAI are neck-and-neck at the frontier. Anthropic's Claude Opus 4.6 trails by 4 points, while Meta's newly released Muse Spark enters at a strong 52.

Open-weight models such as GLM-5.1 and Grok 4.20 occupy the mid-tier, demonstrating that the gap between proprietary and

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Fast

Inference speed is critical for latency-sensitive applications. The leaderboard shows a striking pattern: open-weight and smaller models generate tokens far faster than frontier proprietary models.

The fastest model, gpt-oss-120B (high), outputs 218 tokens/s -- more than 4x the speed of Claude Opus 4.6 and nearly 5x that of DeepSeek V3.2.

This highlights the trade-off between raw intelligence and throughput: the two intelligence leaders sit in the middle of the speed

Rank	Model	Output Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

API pricing per million tokens varies enormously across models, reflecting differences in compute requirements, margin strategies, and competitive positioning.

The cheapest options -- gpt-oss-120B and DeepSeek V3.2 at \$0.30 per million tokens -- are over 33x less expensive than Claude Opus 4.6 at \$10.00.


For budget-sensitive workloads, this price differential is a decisive factor. Gemini 3 Flash offers a compelling mid-range value at

Rank	Model	USD per 1M Tokens
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency

A scatter plot of Intelligence Index versus output tokens consumed during evaluation reveals which models occupy the 'most attractive quadrant' -- high intelligence with efficient token usage.

Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster in the upper region. Grok 4.20 and several Google models sit in the middle-right, balancing intelligence and token efficiency.



Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.1

- Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index, making them the co-leaders as of April 2026.
- Speed and intelligence are inversely correlated: The fastest model (gpt-oss-120B at 218 tokens/s) scores only 33 on intelligence, while the smartest models operate at 49-118 tokens/s.
- Massive pricing spread: A 33x cost gap separates the cheapest (\$0.30/1M tokens) from the most expensive (\$10.00/1M tokens for Claude Opus 4.6).
- Open-weight models are closing the gap: GLM-5.1 at 51 intelligence is within 10% of the proprietary leaders, while also being significantly cheaper at \$2.10/1M tokens.
- No single model wins on all three axes: Choosing a model requires explicit trade-offs among intelligence, speed, and cost depending on the use case.

Recommendation: Prioritize intelligence for complex reasoning tasks, speed for real-time applications, and cost for high-volume workloads. Open-weight models now offer viable alternatives across all three dimensions.