

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

An independent benchmarking analysis of frontier AI models across intelligence, throughput, and cost

Source: Artificial Analysis | April 2026

11

Frontier Models
Benchmarked

57

Top Intelligence
Index Score


33x

Pricing Gap
Cheapest to Most Expensive

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index | Composite score across reasoning, knowledge, and instruction-following benchmarks

Intelligence Index



KEY INSIGHTS


- **Two-way tie at the top:** Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57.
- **Tight clustering:** Top 6 models span just 6 points (51–57), showing fierce competition.
- **Open-weight contender:** GLM-5.1 from Z AI scores 51 — within 10% of the co-leaders.
- **Wide spread:** Scores range from 57 down to 33, a 42% gap between top and bottom.
- **8 organizations** from the US, China, and Europe represented in the top 11.

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus

4.6

Output tokens per second | 10 models with published speed data (Muse Spark excluded)

Output Tokens/s



KEY INSIGHTS

- **Open-weight models dominate:** gpt-oss-120B, Grok, and Nemotron take the top 3 spots.
- **Inverse correlation:** The two intelligence leaders (57 pts each) sit mid-table at 118 and 83 tokens/s.
- **4.7x gap:** Fastest model (218 tok/s) vs. DeepSeek V3.2 (46 tok/s).
- **Latency-sensitive apps** should favor gpt-oss-120B or Grok for real-time workloads.
- **Intelligence cost:** Achieving top-tier reasoning comes with 2–4x throughput penalty.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

USD per 1M output tokens | 10 models with published pricing (Muse Spark excluded)


Rank	Model	Provider	\$/1M Tokens
1	gpt-oss-120B (high)	OpenAI	\$0.30
2	DeepSeek V3.2	DeepSeek	\$0.30
3	NVIDIA Nemotron 3 Super	NVIDIA	\$0.40
4	Gemini 3 Flash	Google	\$1.10
5	GLM-5.1	Z AI	\$2.10
6	Grok 4.20 0309 v2	xAI	\$3.00
7	Gemini 3.1 Pro Preview	Google	\$4.50
8	GPT-5.4 (xhigh)	OpenAI	\$5.60
9	Claude Sonnet 4.6 (max)	Anthropic	\$6.00
10	Claude Opus 4.6 (max)	Anthropic	\$10.00

KEY INSIGHTS

- **33x cost gap:** \$0.30 (gpt-oss-120B, DeepSeek V3.2) vs. \$10.00 (Claude Opus 4.6).
- **Budget tier (<\$1):** Three models under \$1/M tokens offer strong throughput for batch workloads.
- **Mid-range value:** Gemini 3 Flash at \$1.10 pairs 153 tok/s speed with 46 intelligence.
- **Premium tier:** Intelligence leaders Gemini 3.1 Pro and GPT-5.4 cost \$4.50–\$5.60/M tokens.
- **Cost per intelligence point** varies over 10x — model selection must match budget constraints.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence Index vs. output tokens consumed during evaluation | Models in the upper-left quadrant are most attractive



PROPRIETARY vs. OPEN-WEIGHT LEADERS

Metric	Proprietary	Score	Open-Weight	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120B	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B	\$0.30

KEY OBSERVATIONS

- **Upper-left quadrant** = high intelligence, low token usage = most efficient.
- **Gemini 3.1 Pro Preview** and DeepSeek V3.2 cluster in the high-intelligence region.
- **Open-weight speed lead:** gpt-oss-120B is 42% faster than the best proprietary model.
- **Open-weight price lead:** \$0.30 vs. \$1.10 — nearly 4x cheaper.
- **Intelligence gap narrowing:** GLM-5.1 (open) trails Gemini 3.1 Pro (proprietary) by just 6 points.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key findings and actionable recommendations for model selection

KEY FINDINGS

- **Two-way tie at the top:** Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index.
- **Speed–intelligence trade-off:** The fastest model (gpt-oss-120B, 218 tok/s) scores only 33 on intelligence. The smartest models run at 49–118 tok/s.
- **33x pricing spread:** From \$0.30/M tokens (gpt-oss-120B, DeepSeek V3.2) to \$10.00/M tokens (Claude Opus 4.6).
- **Open weights closing the gap:** GLM-5.1 scores 51 intelligence at \$2.10/M tokens — within 10% of proprietary leaders at a fraction of the cost.
- **No single model wins all three axes:** Every selection requires explicit trade-offs among intelligence, speed, and cost.

RECOMMENDATIONS BY WORKLOAD

INTELLIGENCE-FIRST

Complex reasoning, code generation, research

Gemini 3.1 Pro Preview or GPT-5.4 (xhigh)

57 intelligence | 83–118 tok/s | \$4.50–\$5.60/M

SPEED-FIRST

Real-time applications, chatbots, streaming UX

gpt-oss-120B (high) or Grok 4.20

188–218 tok/s | 33–49 intelligence | \$0.30–\$3.00/M

COST-FIRST

High-volume batch processing, classification, extraction

gpt-oss-120B (high) or DeepSeek V3.2

\$0.30/M tokens | 33–42 intelligence | 46–218 tok/s