

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro Preview and GPT-5.4 tie at 57 on the Intelligence Index while a 33× pricing gap separates the cheapest and most expensive frontier models.

Source: Artificial Analysis | April 2026

For AI Engineering Leads & Product Managers Evaluating LLM Providers

Intelligence

57

Speed

218 t/s


Cost

\$0.30/1M

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index — composite of reasoning, knowledge & instruction-following benchmarks | April 2026


Ran k	Model	Provider	Index Score
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33



gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4× Faster Than Claude Opus 4.6

Output speed in tokens per second | Muse Spark excluded (no speed data) | April 2026

Rank	Model	Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46




Key insight: Intelligence and speed are inversely correlated — the two smartest models (57 index) run at only 83–118 t/s

From \$0.30 to \$10 per Million Tokens: A 33× Pricing Gap Across Models


USD per 1M tokens (blended input+output) | Muse Spark excluded (no pricing data) | April 2026

Rank	Model	\$/1M Tokens
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00



33× cost gap: gpt-oss-120B & DeepSeek V3.2 (\$0.30) vs. Claude Opus 4.6 (\$10.00/1M tokens)

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Intelligence Index vs. Output Tokens Consumed During Evaluation
Upper-left quadrant = high intelligence + efficient token use (most attractive)

Proprietary vs. Open-Weight: Best-in-Class by Metric

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (t/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

- Key observations from the scatter plot:
- Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster in the upper region (high intelligence)
 - Grok 4.20 and several Google models balance intelligence and token efficiency
 - MiMo-V2-Pro: very high intelligence but high token usage (upper-right)
 - Llama 4 Maverick: low intelligence and low token region

No single model dominates all three axes — workload priorities must drive model selection

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Two-Way Tie at the Top

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Artificial Analysis Intelligence Index — the frontier is shared across providers.

Speed ↔ Intelligence Trade-off

The fastest model (gpt-oss-120B, 218 t/s) scores only 33 on intelligence. The two co-leaders run at 83–118 t/s.

Massive Pricing Spread

A 33× cost gap separates the cheapest (\$0.30/1M — gpt-oss-120B & DeepSeek V3.2) from the most expensive (\$10.00/1M — Claude Opus 4.6).

Open-Weight Models Closing the Gap

GLM-5.1 scores 51 — within 10% of the proprietary leaders — at \$2.10/1M tokens. Open-weight options now offer near-frontier intelligence at a fraction of cost.

No Single Model Wins

Selecting the right model requires explicit trade-offs. Use-case priorities — latency, intelligence, or cost — should drive model selection.

Actionable Guidance for Teams

Max intelligence workloads → Gemini 3.1 Pro Preview or GPT-5.4
Latency-sensitive / high-volume → gpt-oss-120B or Grok 4.20
Budget-constrained → gpt-oss-120B or DeepSeek V3.2 at \$0.30/1M