

# AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

---

Data source: Artificial Analysis — independent benchmarks across reasoning, knowledge, and instruction-following tasks


Coverage: 11 frontier models from 8 organizations (Google, OpenAI, Anthropic, Meta, Z AI, xAI, DeepSeek, NVIDIA)

Evaluation period: April 2026

**Key insight: No single model dominates across intelligence, speed, and pricing**

# Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index — composite score across reasoning, knowledge, and instruction-following benchmarks



## Key Takeaways

### Two-way tie at the top

Gemini 3.1 Pro Preview and GPT-5.4 both score 57, sharing the #1 position.

### 57-to-33 spread

A 24-point gap separates the leaders from gpt-oss-120B at the bottom.

### Anthropic's strong mid-tier


Claude Opus 4.6 (53) and Sonnet 4.6 (52) both rank in the top 5.

### Open-weight in the mix

GLM-5.1 scores 51 — within 10% of the proprietary leaders.

# gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6

Output tokens per second — critical metric for latency-sensitive applications (Muse Spark excluded: no speed data published)



## Key Takeaways

### 4.5x speed gap

gpt-oss-120B (218 tok/s) vs. DeepSeek V3.2 (46 tok/s) — a 4.5x throughput spread.

### Intelligence-speed trade-off

The two intelligence leaders sit mid-pack on speed: Gemini 3.1 Pro at 118 tok/s, GPT-5.4 at 83 tok/s.

### Open-weight speed advantage


3 of the top 4 fastest models are open-weight (gpt-oss-120B, Grok 4.20, Nemotron 3 Super).

### Gemini 3 Flash: sweet spot

153 tok/s with Intelligence 46 — strong throughput for cost-efficient deployments.

# From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

API pricing per million tokens — reflects compute requirements, margin strategies, and competitive positioning (Muse Spark excluded)



## Key Takeaways

### 33x cost gap

\$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2) vs. \$10.00/1M (Claude Opus 4.6).

### Frontier premium

The two intelligence co-leaders cost \$4.50–\$5.60/1M tokens — 15–19x the cheapest option.

### Best budget value


DeepSeek V3.2 at \$0.30/1M with Intelligence 42; Gemini 3 Flash at \$1.10/1M with Intelligence 46.

### Anthropic's pricing premium

Claude Opus 4.6 is the most expensive model at \$10.00/1M — 3.3x GPT-5.4 and 7.8x Gemini 3 Flash.

# No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Scatter plot: intelligence vs. output tokens consumed during evaluation (not throughput) — upper-left quadrant = most attractive



## Proprietary vs. Open-Weight Leaders

| Metric            | Proprietary Leader     | Score  | Open-Weight Leader  | Score  |
|-------------------|------------------------|--------|---------------------|--------|
| Intelligence      | Gemini 3.1 Pro Preview | 57     | GLM-5.1             | 51     |
| Speed (tok/s)     | Gemini 3 Flash         | 153    | gpt-oss-120B (high) | 218    |
| Price (\$/1M tok) | Gemini 3 Flash         | \$1.10 | gpt-oss-120B (high) | \$0.30 |

## Navigating the Frontier

### Every model demands trade-offs

High intelligence comes with lower speed and higher cost. Fast and cheap models sacrifice capability. No single model optimizes all three.

### Proprietary still leads intelligence

But the margin is thin: GLM-5.1 (open-weight) scores 51 — just 6 points behind the co-leaders at 57.

### Open-weight dominates speed & price

gpt-oss-120B is fastest (218 tok/s) and cheapest (\$0.30/1M) — but scores only 33 on intelligence.

### Gemini 3 Flash: balanced mid-range

\$1.10/1M tokens, 153 tok/s, Intelligence 46 — a viable option when cost matters more than peak capability.

## Recommendation

Map your workload priorities first, then select the model that optimizes the axis that matters most for your use case.

# Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key findings and actionable recommendations for model selection — Artificial Analysis, April 2026

## Key Findings

### 1. Two-way tie at the top

Gemini 3.1 Pro Preview and GPT-5.4 both score 57 on the Intelligence Index — no single model holds the crown.

### 2. Speed and intelligence are inversely correlated

The fastest model (gpt-oss-120B, 218 tok/s) scores only 33 on intelligence; the smartest models run at 49–118 tok/s.

### 3. 33x pricing spread

\$0.30 to \$10.00 per million tokens. Budget-sensitive workloads can achieve massive cost savings with minimal intelligence loss.

### 4. Open-weight models are closing the gap

GLM-5.1 at Intelligence 51 is within 10% of the proprietary leaders, at \$2.10/M tokens — making it a strong value contender.

### 5. No single model wins all three axes

Every model requires explicit trade-offs among intelligence, speed, and cost depending on the use case.

## Actionable Recommendations

### For maximum intelligence

Use Gemini 3.1 Pro Preview or GPT-5.4 — both score 57. Budget for \$4.50–\$5.60/M tokens and accept moderate speed (83–118 tok/s).

### For latency-sensitive applications

Use gpt-oss-120B (218 tok/s, \$0.30/M) or Grok 4.20 (188 tok/s, \$3.00/M). Accept lower intelligence (33–49).

### For cost-optimized deployments

Use DeepSeek V3.2 (\$0.30/M, Intelligence 42) or NVIDIA Nemotron 3 Super (\$0.40/M, Intelligence 36). Best for high-volume, low-stakes tasks.

### For balanced performance

Use Gemini 3 Flash (\$1.10/M, Intelligence 46, 153 tok/s) or GLM-5.1 (\$2.10/M, Intelligence 51, 64 tok/s).

### Bottom line

Map workload priorities first. Select models that optimize the axis that matters most. Re-evaluate quarterly — this landscape shifts fast.