AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro Preview and GPT-5.4 tie at 57 on the Intelligence Index
while a 33x pricing gap separates the cheapest and most expensive frontier models.

Source: Artificial Analysis, April 2026

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

The Intelligence Index aggregates perfo

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

Key Insight: 57-point spread between top (57) and bottom (33) — a 24-point gap

gpt-oss-120B Tops Output Speed at 218 Tokens/s More Than 4x Faster Than Claude Opus 4.6

Source: Artificial Analysis benchmarks, April 2

Rank	Model	Output Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46

Inverse relationship: The fastest model (218 tokens/s) scores only 33 on intelligence, while the smartest models operate at 49-118 tokens/s.

From $0.30 to $10 per Million Tokens: A 33x Pricing Gap Across Models

Rank	Model	USD per 1M Tokens
1	gpt-oss-120B (high)	$0.30
2	DeepSeek V3.2	$0.30
3	NVIDIA Nemotron 3 Super	$0.40
4	Gemini 3 Flash	$1.10
5	GLM-5.1	$2.10
6	Grok 4.20 0309 v2	$3.00
7	Gemini 3.1 Pro Preview	$4.50
8	GPT-5.4 (xhigh)	$5.60
9	Claude Sonnet 4.6 (max)	$6.00
10	Claude Opus 4.6 (max)	$10.00

Budget-sensitive note: At $0.30/M tokens, gpt-oss-120B and DeepSeek V3.2 are 33x cheaper than Claude Opus 4.6 ($10.00/M)

Source: Artificial Analysis pricing data, April

No Model Wins on All Three Axes:
Navigating the Intelligence-Efficiency-Cost Frontier

Source: Artificial Analysis Intelligence Index

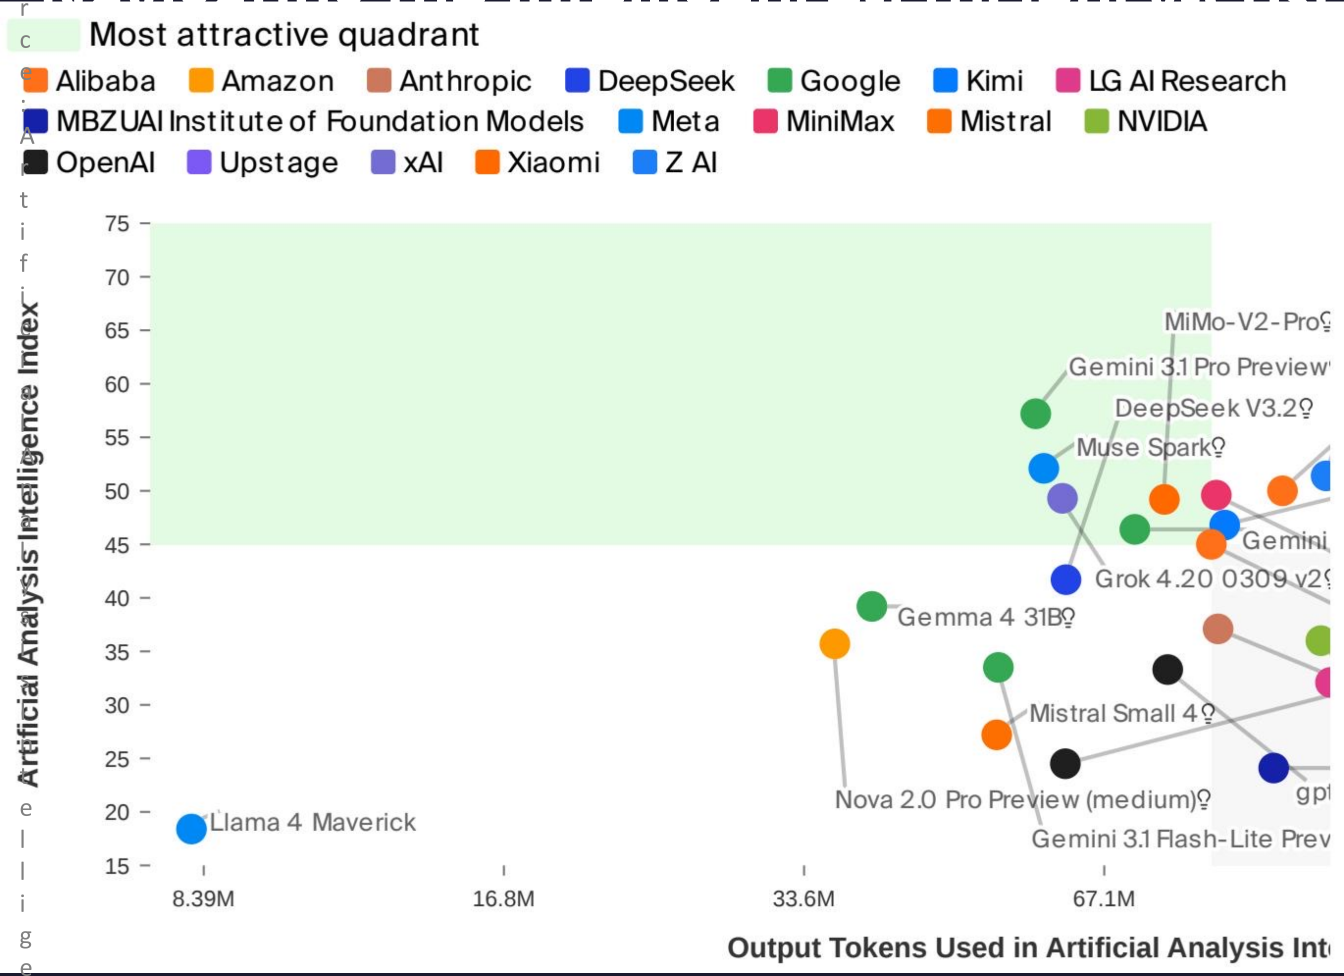

Metric	Proprietary Leader	Score	Open-
Intelligence	Gemini 3.1 Pro Preview	57	
Speed (tokens/s)	Gemini 3 Flash	153	gpt-o
Price ($/1M tokens)	Gemini 3 Flash	$1.10	gpt-o

- Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster in the upper region, indicating strong intelligence.
- Grok 4.20 and Google models balance intelligence and token efficiency.
- Open-weight models dominate speed but score lower on intelligence.
- Every model requires trade-offs among intelligence, speed, and cost.

04

[illegible] ($10.00/1M tokens).

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at $2.10/M Tokens

No single model dominates all three di

Model Selection Guidance

For raw intelligence: Choose Gemini 3.1 Pro Preview or GPT-5.4

For speed-critical workloads: Choose gpt-oss-120B or Grok 4.20

For budget-sensitive production: Choose DeepSeek V3.2 or gpt-oss-120B

For balanced use cases: Consider GLM-5.1 (51 intelligence at $2.10/M)