

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models


Independent benchmarks from Artificial Analysis (captured
April 10, 2026)

Scope: 11 frontier models across intelligence, output speed,
and pricing

Gemini 3.1 Pro Preview and GPT-5.4 Co- Lead the Intelligence Index at 57

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

Full 11-model ranking (Artificial Analysis Intelligence Index)




Intelligence Index = composite across reasoning/knowledge/instruction-following benchmarks

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Inference speed is measured as output tokens per second (Mistral Spark excluded: no speed data in source).

Rank	Model	Output Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46




Takeaway: open-weight models dominate throughput; frontier intelligence leaders sit mid-pack

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

Pricing shown as USD per 1M tokens (Mistral 3 excluded: no published pricing in source).

Rank	Model	USD per 1M Tokens
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00




Cheapest-to-most-expensive gap in source: \$0.30 → \$10.00 (33x)

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-

Scatter: Intelligence Index vs output tokens consumed during evaluation (token efficiency, not tokens/s)

Cost Function



Proprietary vs open-weight leaders (from source brief)

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Key observations from the frontier plot:

- Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster in the upper region (strong intelligence)
- Grok 4.20 0309 v2 and several Google models balance intelligence and token efficiency
- MiMo-V2-Pro appears high-intelligence but high token usage
- Llama 4 Maverick sits low-intelligence and low-token

Decision principle: no single model wins intelligence, throughput, and cost simultaneously

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

What the April 2026 benchmarks imply:

- Frontier intelligence is a two-way tie at 57 (Gemini 3.1 Pro Preview, GPT-5.4 (xhigh))
- Speed and intelligence are inversely correlated: 218 tokens/s (gpt-oss-120B) vs 49–118 tokens/s for top intelligence models
- Cost varies by 33x: \$0.30/1M (gpt-oss-120B, DeepSeek V3.2) vs \$10.00/1M (Claude Opus 4.6)
- Open weights close the gap: GLM-5.1 reaches 51 intelligence at \$2.10/1M tokens

How to select (use-case driven):

- Intelligence-first: choose Gemini 3.1 Pro Preview or GPT-5.4; validate latency budget
- Latency-sensitive / high-throughput: start with gpt-oss-120B, Grok 4.20, or Nemotron; accept lower peak intelligence
- Budget-sensitive: gpt-oss-120B or DeepSeek V3.2; use routing/fallback to a smarter model for hard prompts