

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost


Competitive landscape of frontier AI models

- As of April 2026, frontier model competition is global and tightly contested across at least 8 organizations.
- The intelligence leaders are tied: Gemini 3.1 Pro Preview and GPT-5.4 (xhigh), both scoring 57.
- Open-weight models increasingly pressure incumbents: GLM-5.1 reaches 51 intelligence while gpt-oss-120B leads on speed and price.
- For AI product teams, model selection is no longer a single-winner decision; it is a workload-specific optimization problem.

Framework for this deck: Intelligence, Speed, and Pricing — then integrated trade-offs and selection guidance.

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57


Full 11-model ranking from source brief



Rank	Model	Index
1	Gemini 3.1 Pro Preview	57
2	GPT-5.4 (xhigh)	57
3	Claude Opus 4.6 (max)	53
4	Muse Spark	52
5	Claude Sonnet 4.6 (max)	52
6	GLM-5.1	51
7	Grok 4.20 0309 v2	49
8	Gemini 3 Flash	46
9	DeepSeek V3.2	42
10	NVIDIA Nemotron 3 Super	36
11	gpt-oss-120B (high)	33

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than

Output tokens per second (higher is faster)




Key observations

- gpt-oss-120B leads at 218 tokens/s.
- Grok 4.20 is second at 188 tokens/s.
- Intelligence leaders are mid-pack on throughput: Gemini 3.1 Pro (118), GPT-5.4 (83).
- Claude Opus 4.6 runs at 49 tokens/s.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

USD per 1M tokens (lower is cheaper)




Pricing insight

- Cheapest: gpt-oss-120B and DeepSeek V3.2 at \$0.30.
- Most expensive: Claude Opus 4.6 at \$10.00.
- Cost spread is ~33x from lowest to highest.
- Gemini 3 Flash (\$1.10) is a mid-range value point.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost

Intelligence vs output tokens consumed + proprietary/open-weight comparison



Scatter plot measures output tokens consumed during evaluation (not throughput in tokens/s).

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Trade-off signal

No model leads intelligence, speed, and price simultaneously.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M

How AI teams should choose models by workload priorities

Decision framework

- Prioritize intelligence for complex reasoning and quality-critical workflows.
- Prioritize speed for interactive copilots and real-time UX constraints.
- Prioritize price for high-volume workloads and cost-sensitive deployments.
- Use portfolio routing: pair frontier proprietary models with fast, low-cost open-weight models.

Strategic takeaway

- Open-weight models are rapidly closing the intelligence gap while already winning on throughput and cost.
- For April 2026, leaders should treat model selection as a dynamic, workload-level optimization problem — not a one-time vendor choice.
- Re-benchmark regularly as the frontier shifts quickly across all three axes.