

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Competitive landscape of frontier AI models as of April 2026, compared across intelligence, output speed, and API pricing.

INTELLIGENCE LEADERS

57

Gemini 3.1 Pro Preview and GPT-5.4 tie on Intelligence Index

FASTEST OUTPUT

218 tokens/s

gpt-oss-120B leads throughput

PRICE SPREAD

33x

\$0.30 to \$10 per 1M tokens


Decision lens: no single model dominates all three dimensions, so teams need workload-specific model routing and explicit trade-off thresholds.

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Full 11-model ranking; higher Artificial Analysis Intelligence Index indicates stronger overall capability.

#	Model	Provider	Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

Index score by model




Frontier spread: co-leaders at 57 vs. gpt-oss-120B at 33; GLM-5.1 reaches 51 as an open-weight model.

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus

10-model speed leaderboard; Muse Spark excluded because speed data was not published in the brief.

Output tokens per second



Throughput pattern

218 tokens/s

gpt-oss-120B is fastest

>4x

faster than Claude Opus 4.6

118 / 83


tokens/s for the two intelligence leaders

Implication: latency-sensitive products may trade away peak intelligence for faster response generation.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

Per-million-token API pricing for 10 models with published pricing data; Muse Spark excluded.

USD per 1M tokens



Cost lens

\$0.30

gpt-oss-120B and DeepSeek V3.2 are cheapest

\$10.00


Claude Opus 4.6 is most expensive

Gemini 3 Flash: \$1.10 with 153 tokens/s and 46 intelligence.

Budget-sensitive deployments should treat this spread as a first-order architecture variable.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Scatter plot x-axis is output tokens consumed during evaluation, not throughput in tokens/s.



Frontier observations

- Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster in the upper region.
- Grok 4.20 and several Google models balance intelligence and token efficiency.
- MiMo-V2-Pro shows very high intelligence with substantial token usage.

Proprietary vs. open-weight leaders

Metric	Prop. leader	Score	Open leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Takeaway: proprietary models still lead raw intelligence, but open-weight models dominate speed and price.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Actionable guidance for AI engineering leads and product managers selecting LLM providers.

1

Maximize intelligence

Gemini 3.1 Pro Preview and GPT-5.4 co-lead at 57; use when benchmarked reasoning quality is the gating requirement.

2

Optimize latency

gpt-oss-120B leads at 218 tokens/s; open-weight and smaller models dominate throughput.

3

Control unit cost

gpt-oss-120B and DeepSeek V3.2 are \$0.30/M tokens, versus \$10.00/M for Claude Opus 4.6.

4

Use open weights strategically

GLM-5.1 scores 51 intelligence at \$2.10/M tokens, showing the gap is narrowing.

Recommended selection approach

Set workload priorities first, then benchmark shortlists by axis: frontier reasoning for quality-critical tasks, high-throughput open-weight models for latency-sensitive paths, and low-cost models for budget-sensitive volume.

Core insight: no single model wins on intelligence, speed, and price — portfolio routing beats one-size-fits-all standardization.