AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Independent Analysis of AI Models and API Providers

Source: Artificial Analysis, April 2026

Target Audience: AI Engineering Leads & Product Managers

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

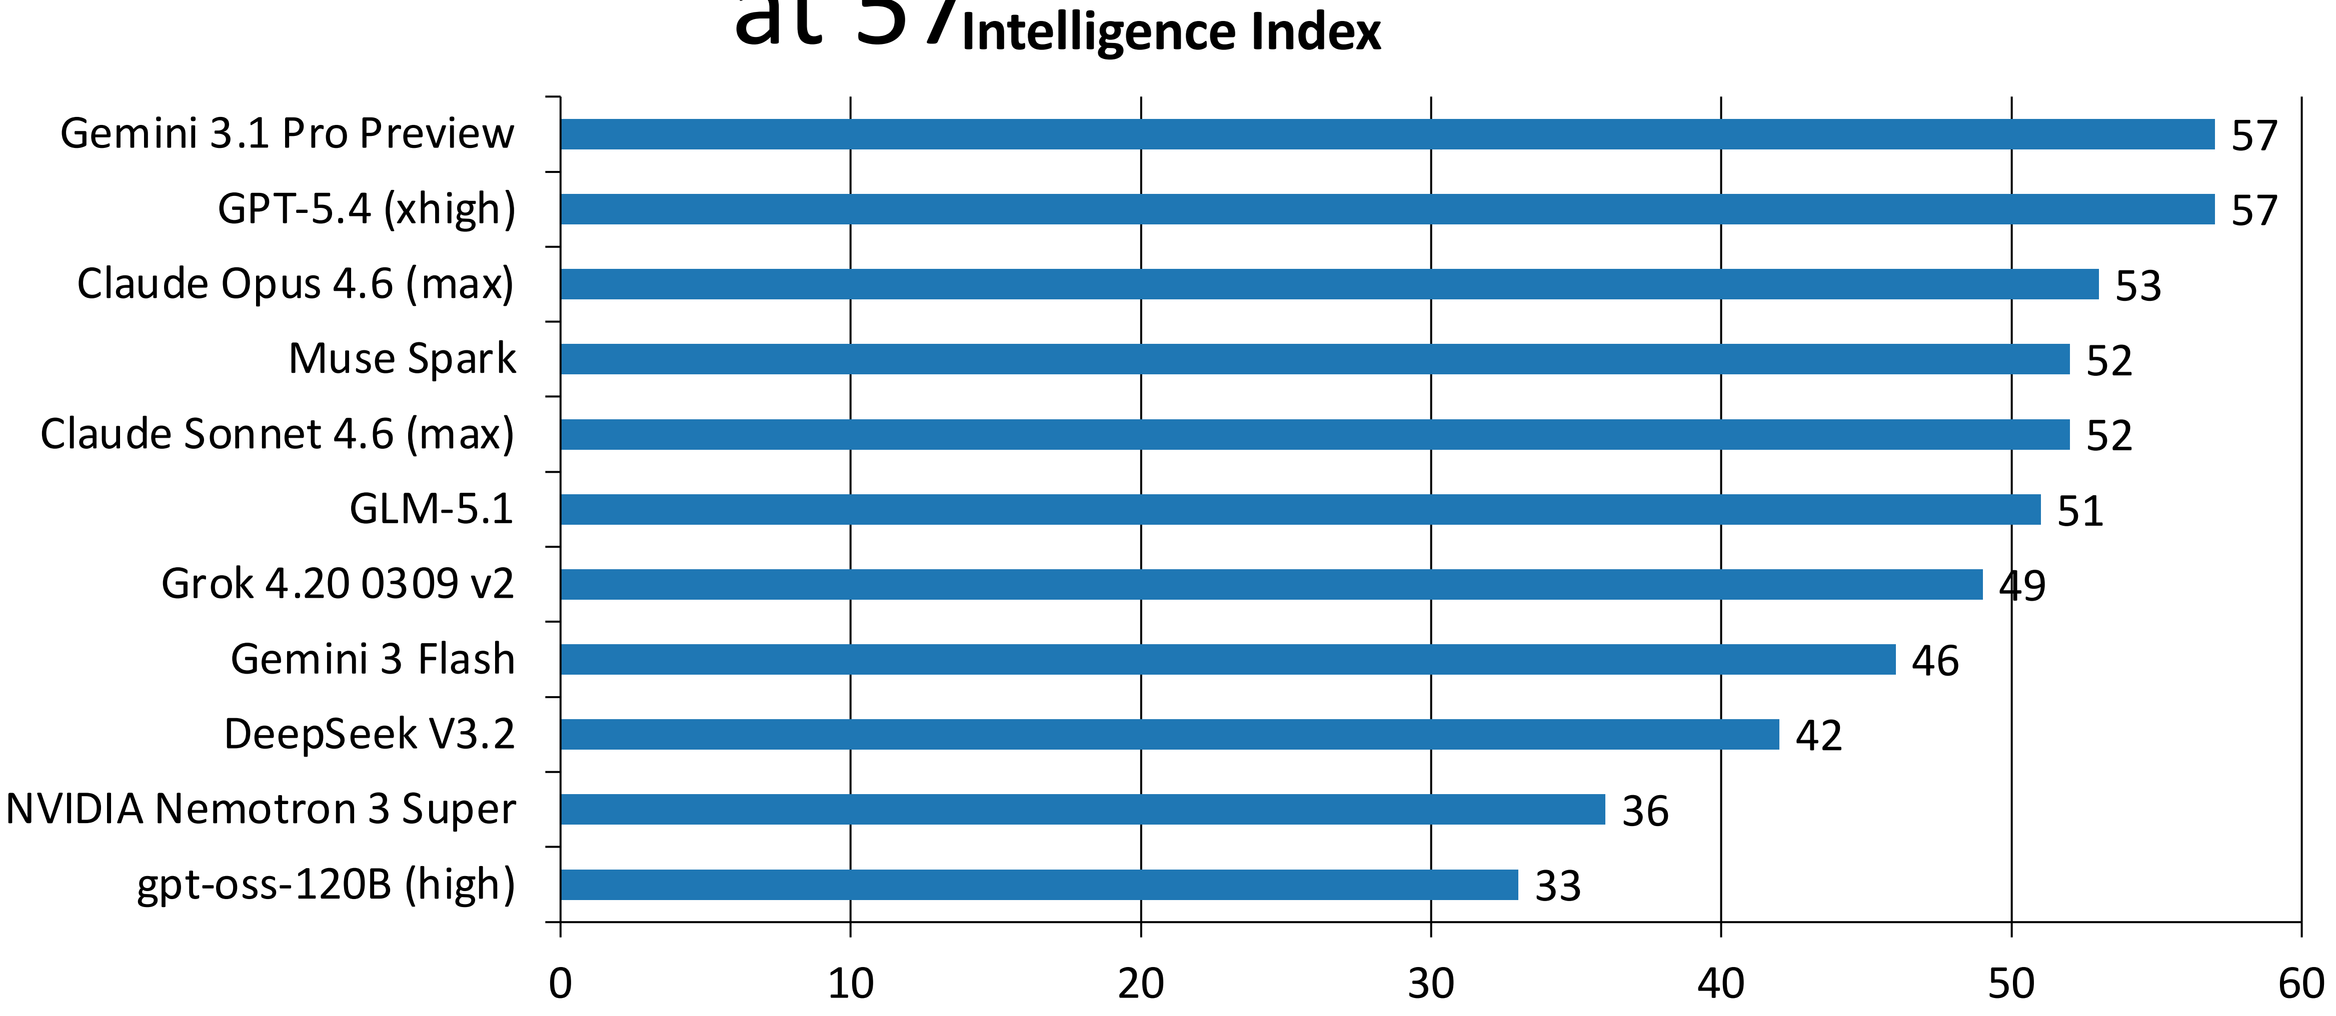

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

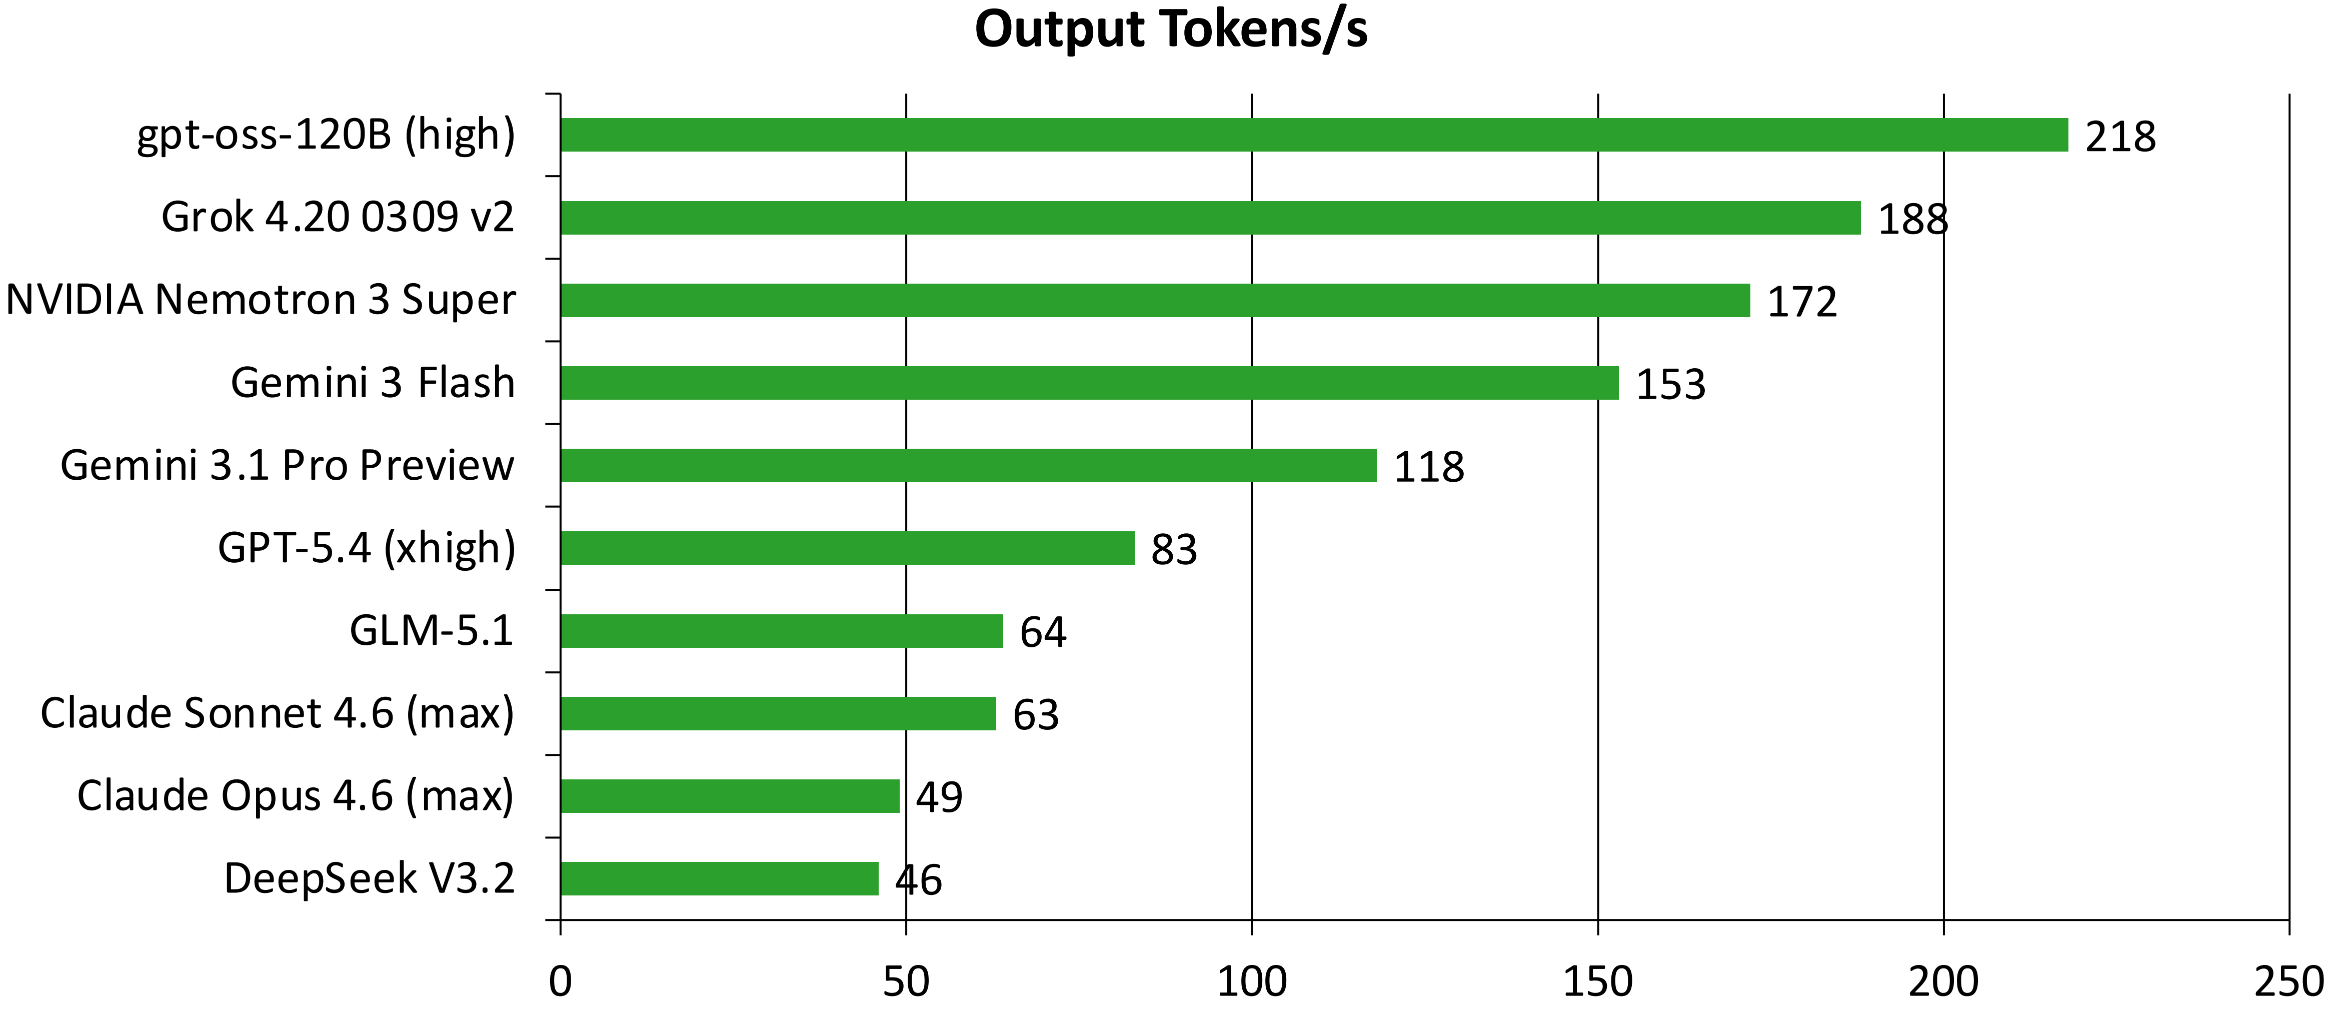

From $0.30 to $10 per Million Tokens: A 33x Pricing Gap Across Models

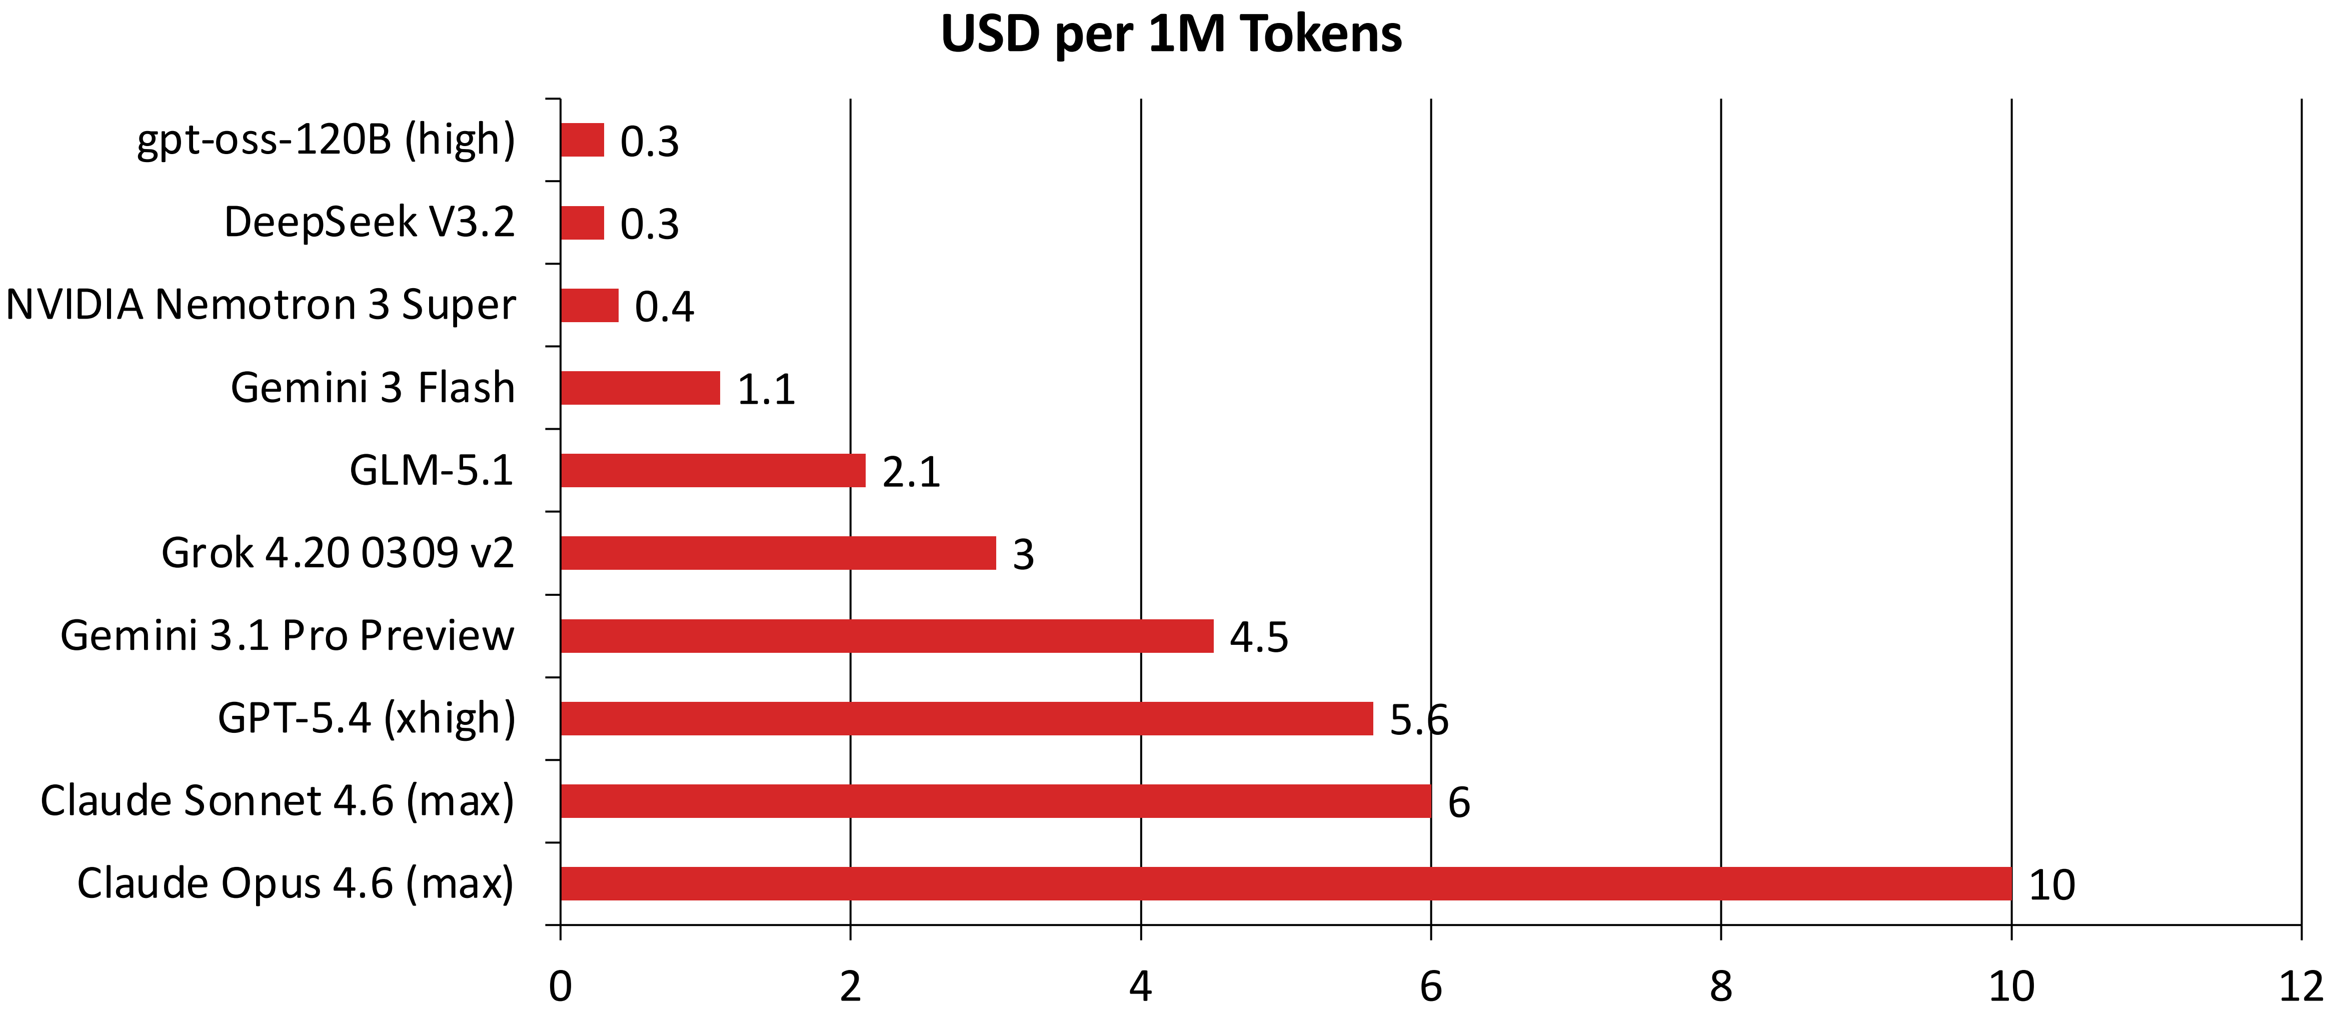

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

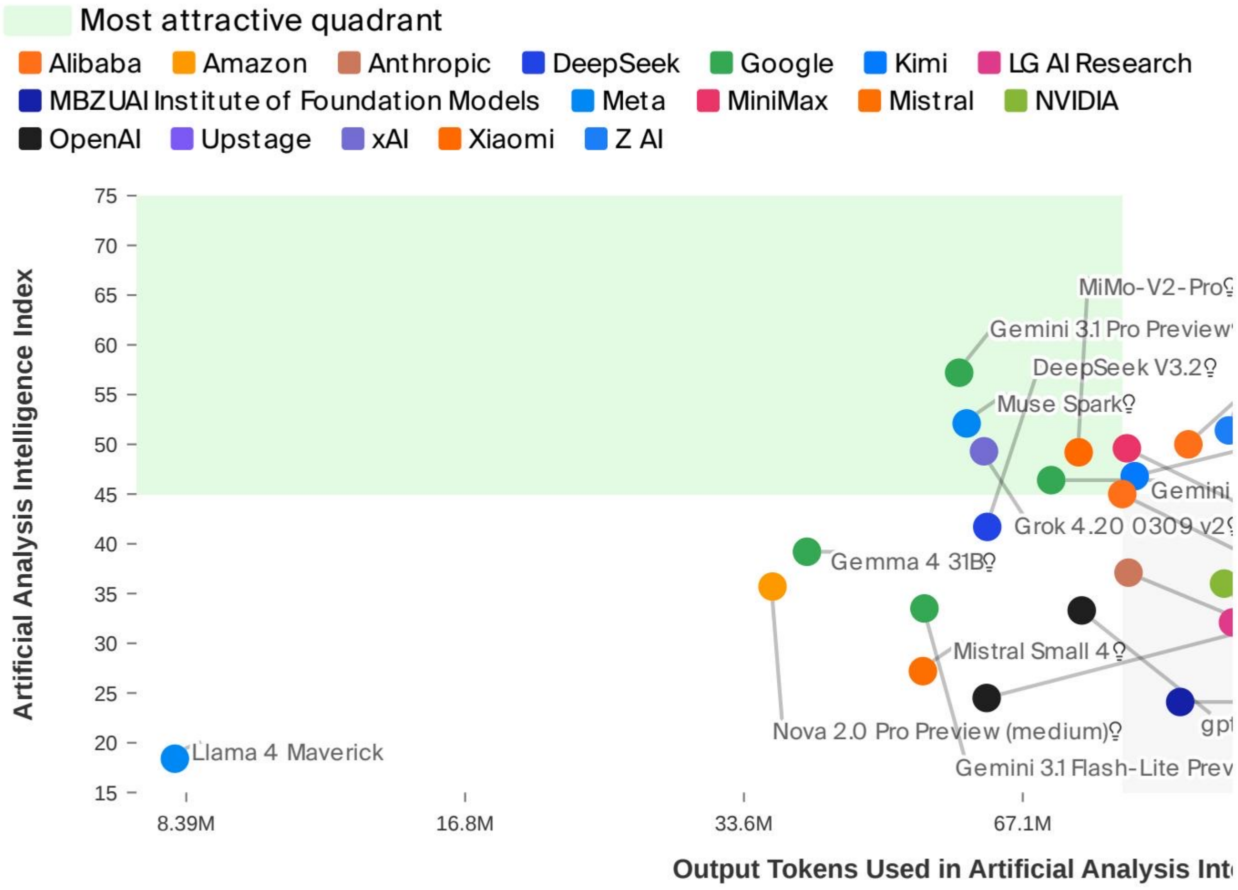

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price ($/1M tokens)	Gemini 3 Flash	$1.10	gpt-oss-120B (high)	$0.30

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at $2.10/M Tokens

- Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index.
- Speed and intelligence are inversely correlated: The fastest model (gpt-oss-120B at 218 tokens/s) scores only 33 on intelligence, while the smartest models operate at 49-118 tokens/s.
- Massive pricing spread: A 33x cost gap separates the cheapest ($0.30/1M tokens) from the most expensive ($10.00/1M tokens).
- Open-weight models are closing the gap: GLM-5.1 at 51 intelligence is within 10% of the proprietary leaders, while being significantly cheaper at $2.10/1M tokens.
- Conclusion: No single model wins on all three axes; choosing a model requires explicit trade-offs among intelligence, speed, and cost depending on the use case.