



AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost
Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33x pricing gap separates the cheapest and most expensive frontier models.

11

Models Evaluated

3


Evaluation Axes

8+

Organizations

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Two models share the top at 57; the full spread from leader to laggard is 57-to-33 — open-weight GLM-5.1 scores 51, within 10% of the co-leaders.




Artificial Analysis Intelligence Index Score

Proprietary Open-Weight

Open-weight models like GLM-5.1 (51) are within 10% of the proprietary leaders, demonstrating the narrowing gap in model intelligence.

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6

Open-weight and smaller models dominate throughput. The intelligence leaders (57 score) sit mid-speed at 83–118 tokens/s — an inverse correlation.




Output Tokens per Second

Proprietary Open-Weight

Key Insight: The fastest model (gpt-oss-120B, 218 tok/s) scores only 33 on intelligence. The smartest models operate at 49–118 tok/s. Muse Spark excluded — no published speed data.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

Budget-sensitive workloads can save 33x by choosing gpt-oss-120B or DeepSeek V3.2. Gemini 3 Flash offers a compelling mid-range value at \$1.10.




USD per 1 Million Output Tokens

\$ 33x Gap: Claude Opus 4.6 costs \$10.00/1M tokens vs. gpt-oss-120B and DeepSeek V3.2 at \$0.30 — a decisive factor for budget-sensitive workloads. Muse Spark excluded — no published pricing data.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Every model requires trade-offs among intelligence, token efficiency, and cost. Model selection must be driven by specific workload priorities.



Intelligence Index vs Output Tokens Consumed During Evaluation
Upper-left quadrant = most attractive (high intelligence, fewer tokens consumed)

Proprietary vs Open-Weight Leaders

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120B	218
Price (\$/1M tok)	Gemini 3 Flash	\$1.10	gpt-oss-120B	\$0.30

Key Observations from the Scatter Plot

- Gemini 3.1 Pro Preview and DeepSeek V3.2 cluster in the upper region — strong intelligence scores
- Grok 4.20 and Google models sit in the middle-right, balancing intelligence and token efficiency
- Llama 4 Maverick occupies the low-intelligence, low-token region — neither efficient nor highly capable
- MiMo-V2-Pro shows very high intelligence but with substantial token usage — upper-right corner

↔ The intelligence leaderboard spans 8+ organizations across the US, China, and Europe — a genuinely global competition.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

No single model wins on all three axes — choose based on your workload priorities.

Key Findings



Two-Way Tie at the Top
Gemini 3.1 Pro Preview and GPT-5.4 both score 57 on the Artificial Analysis Intelligence Index.



Speed ≠ Intelligence
Fastest model (gpt-oss-120B, 218 tok/s) scores only 33; smartest models operate at 49–118 tok/s.



33x Pricing Gap
\$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2) to \$10.00/1M tokens (Claude Opus 4.6).



Open-Weight Models Closing the Gap
GLM-5.1 reaches 51 intelligence — within 10% of leaders — at only \$2.10/1M tokens.

Actionable Model Selection Guidance



Latency-Critical

Best pick: gpt-oss-120B (high)
218 tok/s | \$0.30/1M tokens
Intelligence: 33

Real-time apps, high-volume inference, chatbot routing, simple classification



Intelligence-Critical

Best picks: Gemini 3.1 Pro / GPT-5.4
Score: 57 | 83–118 tok/s
\$4.50–\$5.60/1M tokens

Complex reasoning, code generation, research, agentic workflows



Cost-Sensitive

Best pick: DeepSeek V3.2
\$0.30/1M tokens | Intelligence: 42
46 tok/s

Batch processing, data extraction, summarization at scale

 **Bottom line:** Define your workload profile first — then pick the model that optimizes for your primary axis.