

Artificial Analysis · April 2026

AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost
Across 11 Frontier Models


For AI Engineering Leads & Product Managers

April 2026



Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57


Source: Artificial Analysis Intelligence Index, April 2026



gpt-oss-120B Tops Output Speed at 218 Tokens/s

— More Than 4× Faster Than Claude Opus 4.6

Source: Artificial Analysis, April 2026 · Output tokens per second (median) · Note: Muse Spark excluded — no published speed data




Key Trade-off:

Intelligence leaders run slower — Gemini 3.1 Pro (118 tok/s) & GPT-5.4 (83 tok/s) sit mid-table on speed, while the fastest model (gpt-oss-120B, 218 tok/s) scores only 33 on intelligence.

From \$0.30 to \$10 per Million Tokens: A 33× Pricing Gap Across Models


Source: Artificial Analysis, April 2026 · USD per 1M tokens (blended input+output) · Note: Muse Spark excluded — no published pricing data



33× cost gap between cheapest and most expensive frontier model

Gemini 3 Flash (\$1.10/M) offers strong mid-range value: 46 intelligence + 153 tok/s at 9× less than Claude Opus 4.6

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Proprietary vs. Open-Weight: Best-in-Class

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview Google	57	GLM-5.1 Z AI	51
Speed (tok/s)	Gemini 3 Flash Google	153	gpt-oss-120B (high) OpenAI	218
Price (\$/1M)	Gemini 3 Flash Google	\$1.10	gpt-oss-120B (high) OpenAI	\$0.30

Intelligence Gap Closing

Open-weight GLM-5.1 (51) is only 11% behind proprietary leaders (57)
57 pts (prop) vs 51 pts (open)

Speed & Cost Won by Open

gpt-oss-120B leads speed (218 tok/s) AND price (\$0.30)
No proprietary model matches both dimensions simultaneously

No single model wins on all three axes

Choosing a model requires explicit trade-offs: the most intelligent models are 4–5× slower and up to 33× more expensive than their open-weight counterparts optimized for speed and cost.

Match Model to Workload

Max Intelligence

Gemini 3.1 Pro Preview
or GPT-5.4 (xhigh)

Speed-Critical

gpt-oss-120B (218 tok/s)
or Grok 4.20 (188 tok/s)

Budget-Sensitive

gpt-oss-120B or
DeepSeek V3.2 (\$0.30/M)

Balanced Value

Gemini 3 Flash
\$1.10 · 153 tok/s · IQ 46

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key Findings & Actionable Recommendations · Artificial Analysis, April 2026

1 Two-Way Tie at the Top

57 points — Gemini 3.1 Pro Preview & GPT-5.4 (xhigh) share the Intelligence Index lead as of April 2026

Google and OpenAI are neck-and-neck at the frontier

2 Speed ↔ Intelligence Trade-off

4x+ speed gap — fastest model (gpt-oss-120B, 218 tok/s) scores only 33 on intelligence
Top models run 49–118 tok/s

Inverse correlation: smarter models are slower

3 Massive Pricing Spread

33x cost gap separates cheapest (\$0.30/1M — gpt-oss-120B & DeepSeek V3.2) from most expensive (\$10/1M — Claude Opus)

Budget-sensitive workloads: price is the decisive factor

4 Open-Weights Closing Fast

51 intelligence — GLM-5.1 is within 11% of proprietary leaders at just \$2.10/M tokens
gpt-oss-120B dominates on speed AND price simultaneously

8+ organizations across US, China & Europe — global competition

Actionable Model Selection Guidance

MAX INTELLIGENCE

Gemini 3.1 Pro Preview
or GPT-5.4 (xhigh)
IQ 57 · \$4.50–\$5.60/M
83–118 tok/s

Complex reasoning · Coding · Research

SPEED-CRITICAL

gpt-oss-120B (high)
or Grok 4.20 0309 v2
218–188 tok/s
IQ 33–49 · \$0.30–\$3.00/M

Streaming · Real-time · Chat

BUDGET-SENSITIVE

gpt-oss-120B (high)
or DeepSeek V3.2
\$0.30/M tokens
33–42 intelligence

High-volume · Batch · Classification

BALANCED VALUE

Gemini 3 Flash
\$1.10/M · 153 tok/s · IQ 46
Strong mid-range across all
three dimensions

Production workloads · General use