

AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost
Across 11 Frontier Models

57

2 models tie at top intelligence

33x


pricing gap across models

218

tokens/s top output speed


Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

KEY INSIGHT 24-point spread across 11 models | Open-weight models (GLM-5.1, Grok 4.20) within striking distance of leaders




gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6

KEY INSIGHT Speed and intelligence are inversely correlated — the two intelligence leaders sit mid-table on throughput




From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

KEY INSIGHT Massive pricing spread — budget-sensitive workloads have compelling sub-\$1 options; Gemini 3 Flash offers mid-range value



No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

KEY INSIGHT Every model requires trade-offs among intelligence, token efficiency, and cost — selection must be use-case driven



Proprietary vs. Open-Weight

Metric	Proprietary	Open-Weight
Intelligence	Gemini 3.1 Pro 57	GLM-5.1 51 -6 pts
Speed (tokens/s)	Gemini 3 Flash 153	gpt-oss-120B 218 +42%
Price (\$/1M tokens)	Gemini 3 Flash \$1.10	gpt-oss-120B \$0.30 3.7x cheaper

The Open-Weight Closing Gap

GLM-5.1 scores 51 intelligence — within 10% of the proprietary co-leaders — while costing 2.1x less than Gemini 3.1 Pro and 4.8x less than GPT-5.4.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens



Intelligence

Frontier competition tightens

57 / max score

- Two-way tie: Gemini 3.1 Pro and GPT-5.4 both score 57
- GLM-5.1 at 51 is within 10% of proprietary leaders
- 8+ organizations from US, China, and Europe compete

TIGHT RACE



Speed & Cost

Open-weight dominates

33x price gap

- gpt-oss-120B: 218 tokens/s at \$0.30/M tokens
- Speed and intelligence are inversely correlated
- Gemini 3 Flash: compelling mid-range value at \$1.10

OPEN ADVANTAGE



Model Selection

Match to your workload

No single model wins all 3 axes

Quality-critical

Gemini 3.1 Pro / GPT-5.4

Latency-sensitive

gpt-oss-120B / Grok 4.20

Budget-constrained

GLM-5.1 / Gemini 3 Flash

Best balance

Gemini 3 Flash (\$1.10, 153 t/s, 46 IQ)

USE-CASE FIRST