


AI Model Landscape

April 2026

Benchmarking Intelligence, Speed, and Cost
Across 11 Frontier Models


Independent Analysis



Intelligence Index · Output Speed · API Pricing

Intelligence Index: 11-Model Ranking

KEY INSIGHT: Gemini 3.1 Pro Preview and GPT-5.4 co-lead at 57 — but the spread from top to bottom is 24 points.



Output Speed: 10-Model Leaderboard


KEY INSIGHT: gpt-oss-120B leads at 218 tokens/s — 4.4x faster than Claude Opus 4.6 (49 tokens/s).

Inverse Pattern:

Intelligence leaders rank mid-table on speed.


5 of top 10 speed are open-weight.

- Open-Weight
- Google
- Anthropic



API Pricing: 33x Gap from Cheapest to Most Expensive

KEY INSIGHT: gpt-oss-120B and DeepSeek V3.2 at \$0.30/M tokens vs Claude Opus 4.6 at \$10.00/M — a 33x spread.




Best Value Picks:

- Gemini 3 Flash: \$1.10 + 153 tok/s
- GLM-5.1: \$2.10 + 51 intelligence

33x Price Range
\$0.30 to \$10.00
per million tokens

Trade-offs: No Model Wins on All Three Axes

KEY INSIGHT: Every model requires trade-offs — choose based on your workload priorities.



Proprietary vs. Open-Weight

Leaders across key dimensions

Metric	Proprietary Leader	Open-Weight Leader
Intelligence	Gemini 3.1 Pro Preview 57	GLM-5.1 51
Speed (tokens/s)	Gemini 3 Flash 153	gpt-oss-120B 218
Price (\$/1M tokens)	Gemini 3 Flash \$1.10	gpt-oss-120B \$0.30

Strategic Implication:

- Open-weight leads speed by 42% (218 vs 153)
- Open-weight leads price by 73% (\$0.30 vs \$1.10)
- Proprietary leads intelligence by 12% (57 vs 51)

Model Selection Guidance:

Intelligence-first → Gemini 3.1 Pro / GPT-5.4 | Speed-first → gpt-oss-120B / Grok 4.20 | B

GLM-5.1 reaches 51 intelligence at \$2.10/M tokens — within 10% of proprietary leaders at 1/5 the price.

1. Two-way tie at the frontier

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index — no clear leader.

2. Speed vs. intelligence: inverse relationship

The fastest model (gpt-oss-120B at 218 tok/s) scores only 33 on intelligence. The two intelligence leaders (Gemini 3.1 Pro Preview and GPT-5.4) are slower, at 188 tok/s.

3. 33x pricing spread

\$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2) vs \$10.00/1M tokens (Claude Opus 4.6). For context, Gemini 3.1 Pro Preview is at \$4.50/1M tokens.

4. Open-weight models are closing the gap

GLM-5.1 at 51 intelligence is within 10% of proprietary leaders, at \$2.10/1M tokens — significantly cheaper than Gemini 3.1 Pro Preview (\$4.50/1M tokens) and GPT-5.4 (\$5.60/1M tokens).

Model Selection Guidance

Intelligence-Critical Tasks

Gemini 3.1 Pro Preview (57, \$4.50) or GPT-5.4 (57, \$5.60)

Speed-Critical Tasks

Co-leaders on intelligence: gpt-oss-120B (218 tok/s, \$0.30) or Grok 4.20 (188 tok/s, \$3.00)

Budget-Sensitive Tasks

4x+ faster than V3.2, might tie for lowest cost: gpt-oss-120B (\$0.30/M tokens, 42-46 intelligence)