AI Model Landscape April 2026:

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Independent benchmarks across intelligence, output speed, and API pricing

Key snapshot

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index
while a 33x pricing gap separates the cheapest and most expensive frontier models.

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Co-leaders at 57; overall spread is 57 → 33 across 11 models.

Rank	Model	Provider	Index
1	Gemini 3.1 Pro Preview	Google	**57**
2	GPT-5.4 (xhigh)	OpenAI	**57**
3	Claude Opus 4.6 (max)	Anthropic	**53**
4	Muse Spark	Meta	**52**
5	Claude Sonnet 4.6 (max)	Anthropic	**52**
6	GLM-5.1	Z AI	**51**
7	Grok 4.20 0309 v2	xAI	**49**
8	Gemini 3 Flash	Google	**46**
9	DeepSeek V3.2	DeepSeek	**42**
10	NVIDIA Nemotron 3 Super	NVIDIA	**36**
11	gpt-oss-120B (high)	OpenAI	**33**

What this ranking implies

- Google + OpenAI are tied at the frontier (57).
- Claude Opus 4.6 trails by 4 points (53).
- Open-weight models sit mid-tier (GLM-5.1 at 51).
- Index range: 57 to 33 across the set.

Selection hint

If your workload is reasoning-heavy, treat the top tier as your default baseline—then optimize for speed/cost.

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Open-weight models dominate throughput; frontier leaders sit mid-table on speed.

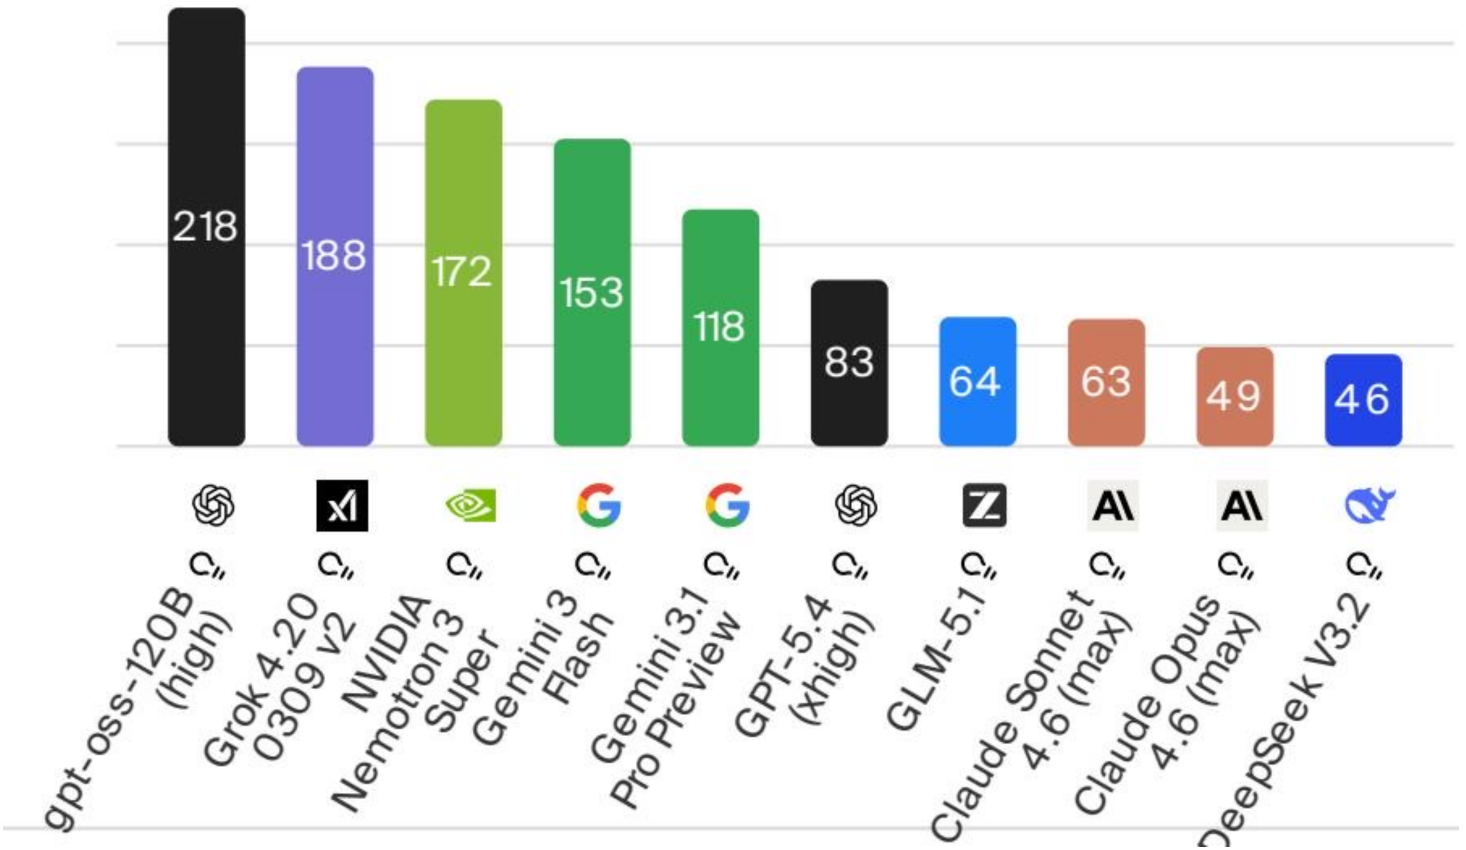

Key callouts

- #1: gpt-oss-120B (high) at 218 tokens/s.
- More than 4× Claude Opus 4.6 (49 tokens/s).
- Speed and intelligence are inversely correlated.
- Gemini 3.1 Pro: 118; GPT-5.4: 83 tokens/s.

Practical guidance

For latency-sensitive apps (chat, agents, UI loops), throughput can be the dominant constraint—even when the most intelligent models exist.

From /bin/bash.30 to 0 per Million Tokens: A 33x Pricing Gap Across Models

Pricing varies by 33× across models; cost structure can dominate total system design.

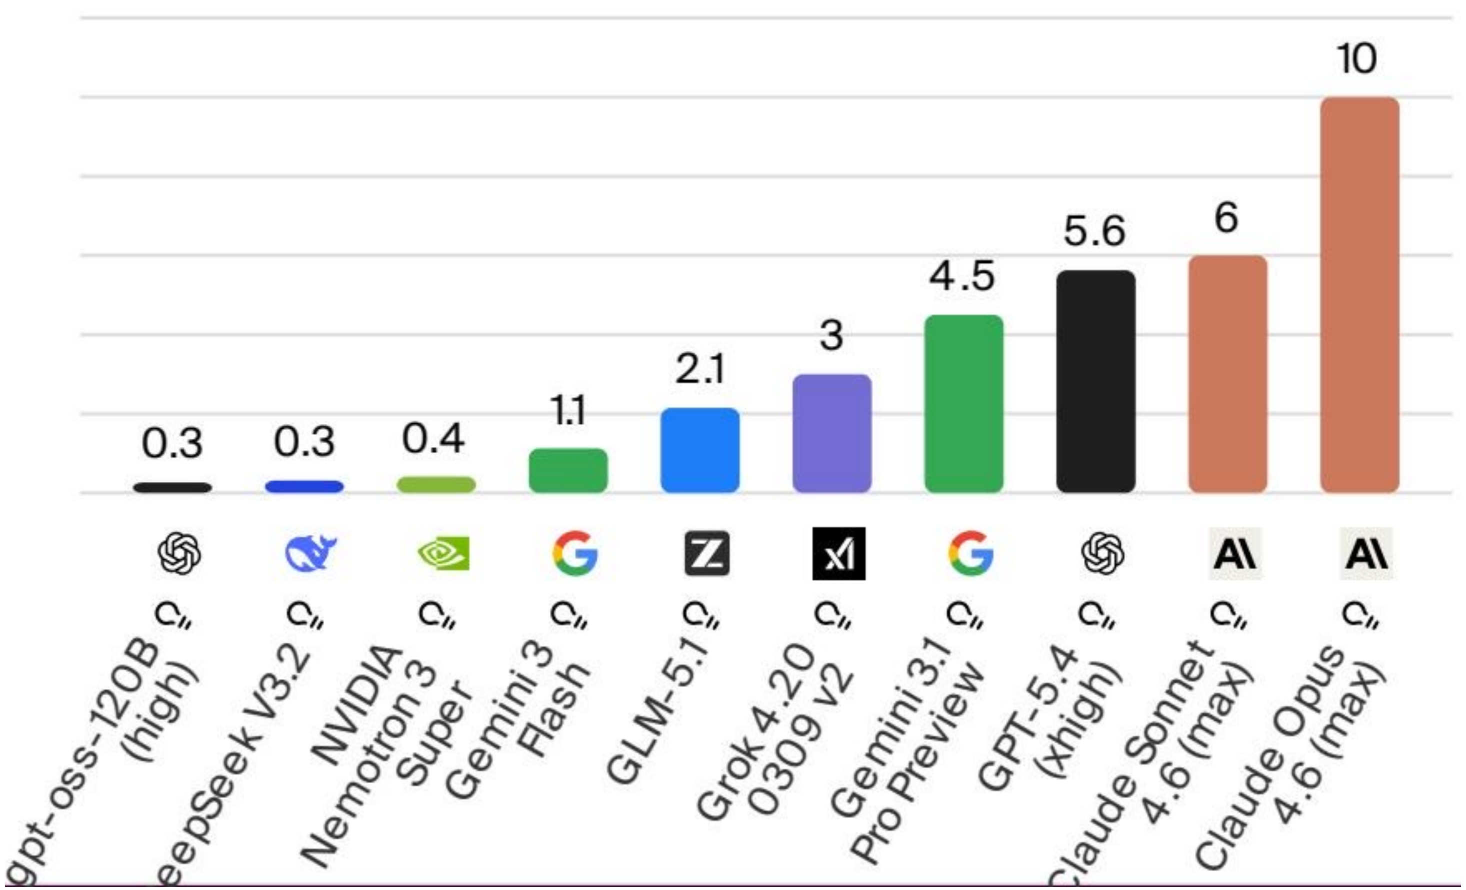

Cost anchors (USD per 1M tokens)

- Cheapest: /bin/bash.30 (gpt-oss-120B, DeepSeek V3.2)
- Most expensive: 0.00 (Claude Opus 4.6)
- Mid-range value: Gemini 3 Flash at .10

How to use this slide

Price is a first-order driver for batch + high-volume apps (retrieval, summarization, offline eval). Combine with intelligence + speed to pick the right tier per workload.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Every provider lands on a different point of the frontier—optimize per workload, not globally.

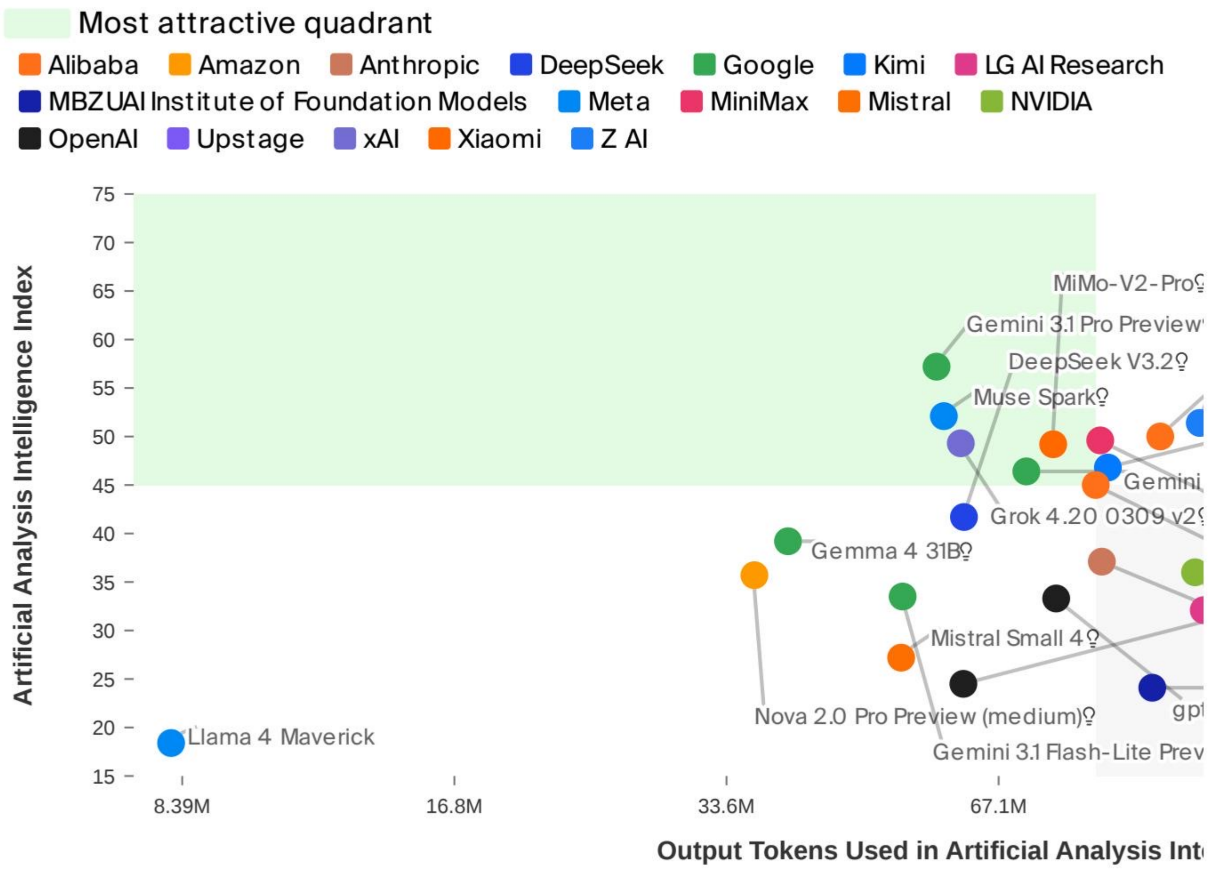

Scatter x-axis = output tokens consumed during evaluation (efficiency), not tokens/s throughput.

Proprietary vs open-weight leaders

Metric	Proprietary leader	Open-weight leader
Intelligence	Gemini 3.1 Pro Preview Score: 57	GLM-5.1 Score: 51
Speed (tokens/s)	Gemini 3 Flash Score: 153	gpt-oss-120B (high) Score: 218
Price ($/1M tokens)	Gemini 3 Flash Score: .10	gpt-oss-120B (high) Score: /bin/bash.30

Interpretation

Open weights lead on speed + cost; proprietary still edge out intelligence.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at .10/M Tokens

Actionable selection guidance based on workload priorities

What the April 2026 snapshot says

- Two-way intelligence tie at 57 (Gemini 3.1 Pro Preview, GPT-5.4).
- Speed & intelligence trade off; fastest model is 218 tokens/s.
- Pricing spans 33× (/bin/bash.30 to 0 per 1M tokens).
- Open weights narrow the gap: GLM-5.1 scores 51 at .10/1M.

No single winner → route by workload

- Intelligence-first: use top-tier models as baseline for hard reasoning.
- Latency-first: prioritize throughput (tokens/s) for interactive experiences.
- Cost-first: select low $/token for batch + high-volume pipelines.
- Hybrid routing: mix models across stages to hit SLOs + budgets.

REFERENCE ARCHITECTURE (MODEL ROUTING)

Workload classification

- Reasoning-heavy
- Latency-sensitive
- High-volume / batch

→

Model tier selection

- Frontier proprietary (max intelligence)
- Open-weight (speed + cost)
- Value mid-tier (balanced)

→

Outcome targets

- Quality thresholds
- Latency SLOs
- Cost budgets