

Frontier AI Model Landscape (April 2026)

Competitive dynamics across intelligence, speed, and pricing show persistent trade-offs.

Intelligence Leaders

Top scores cluster among latest proprietary releases, with open-weight models closing rapidly.

Speed Dispersion

Throughput and latency vary by architecture and serving stack; fast models are not always top-intelligence.

Pricing Spread

Token pricing differs by order(s) of magnitude across premium frontier and cost-optimized alternatives.

Core Insight

No single model dominates intelligence, speed, and pricing simultaneously.

Selection should be workload-specific, with open-weight options increasingly competitive.

Intelligence Benchmark Ranking (11 Models)

Full ranking reproduced from source brief.

Rank	Model	Score
1	GPT-5.1-codex	71
2	Gemini 2.5 Pro	70
3	Claude 4 Opus	64
4	o3	62
5	Grok 4	58
6	Qwen3-235B	56
7	Claude 3.7 Sonnet	56
8	Gemini 2.5 Flash	53
9	DeepSeek R1 0528	51
10	Llama 4 Maverick	45
11	Nemotron Ultra	42

Throughput & Latency: Speed Comparison

Inference speed varies significantly by provider and model class.

Model	Output Tokens/s	TTFT (s)	Speed Notes
GPT-5.1-codex	58.8	48.7	High output throughput
Gemini 2.5 Pro	31.9	20.0	Balanced profile
Claude 4 Opus	23.4	45.2	Lower throughput
o3	84.4	33.1	Fastest in group
Grok 4	72.0	44.6	Fast tokens, slower start
Qwen3-235B	61.0	39.7	Competitive speed
Gemini 2.5 Flash	78.6	6.2	Very low TTFT
DeepSeek R1 0528	33.2	13.3	Moderate throughput
Note: Speed metrics are distinct from output tokens consumed in benchmark evaluations.			


Pricing Comparison Across Leading Models

Input/output token pricing shows large dispersion across frontier options.

Model	Input \$/1M	Output \$/1M	Price Tier
GPT-5.1-codex	1.25	10.00	Premium
Gemini 2.5 Pro	1.25	10.00	Premium
Claude 4 Opus	15.00	75.00	Ultra-premium
o3	2.00	8.00	Premium
Grok 4	3.00	15.00	Premium
Qwen3-235B	0.80	4.00	Value
Gemini 2.5 Flash	0.30	2.50	Low-cost
DeepSeek R1 0528	0.55	2.19	Low-cost
Cost leadership and intelligence leadership remain misaligned in April 2026.			

Trade-off Frontier: Intelligence vs Output Tokens

Benchmark scatter measures output tokens consumed during evaluation (not throughput).



Proprietary vs Open-weight

Group	Avg IQ	Avg Tokens
Proprietary	65	20.4k
Open-weight	51	16.7k

Implications

- Open-weight models use fewer output tokens on average.
- Intelligence gap is narrowing, especially in upper open tier.
- Frontier selection requires workload-specific trade-offs.

Actionable Model Selection Framework

Choose models by dominant workload constraint rather than single leaderboard rank.

Workload Priority	Primary Axis	Recommended Model Class	Operating Pattern
High-stakes reasoning/coding	Intelligence	Top proprietary frontier	Quality-first routing
Real-time interactive apps	Speed	Low-TTFT/high-throughput models	Latency-aware fallback
Large-scale batch generation	Pricing	Cost-optimized/open-weight tier	Budget-capped orchestration
Balanced enterprise portfolio	All three	Hybrid proprietary + open-weight	Multi-model routing

Strategic Takeaway

Open-weight models are closing the gap; winning teams optimize model mix by workload, not by a single universal “best model.”