

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Frontier model selection is now a multi-objective problem

Gemini 3.1 Pro Preview and GPT-5.4 tie at 57 on intelligence, while the fastest and cheapest options come from different parts of the market.

11 frontier models

3 evaluation axes

April 2026

57

Co-leading

Intelligence Index score

218

Top speed

output tokens/s

33x

Pricing gap

\$0.30 to \$10 per 1M tokens

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Full 11-model Intelligence Index ranking

Composite score across reasoning, knowledge, and instruction-following benchmarks

1	Gemini 3.1 Pro Preview	Google		57
2	GPT-5.4 (xhigh)	OpenAI		57
3	Claude Opus 4.6 (max)	Anthropic		53
4	Muse Spark	Meta		52
5	Claude Sonnet 4.6 (max)	Anthropic		52
6	GLM-5.1	Z AI		51
7	Grok 4.20 0309 v2	xAI		49
8	Gemini 3 Flash	Google		46
9	DeepSeek V3.2	DeepSeek		42
10	NVIDIA Nemotron 3 Super	NVIDIA		36
11	gpt-oss-120B (high)	OpenAI		33

57

Two-way tie: Google and OpenAI are neck-and-neck at the frontier.

57-to-33 spread

Capability remains stratified even among frontier and open-weight models.


Open weights narrow the gap

GLM-5.1 scores 51, just 6 points behind the co-leaders.

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

Output-speed leaderboard

Inference speed measured in output tokens per second; Muse Spark excluded because speed data is not listed.



218 tokens/s

Fastest observed output speed:
gpt-oss-120B (high).

4x+

More than 4x Claude Opus 4.6
and nearly 5x DeepSeek V3.2.

Trade-off signal

The two intelligence leaders
sit mid-table on speed: 118 and
83 tokens/s.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

Published API price per 1M tokens

Lower is better; Muse Spark excluded because published pricing is not listed.

1	gpt-oss-120B (high)		\$0.30
2	DeepSeek V3.2		\$0.30
3	NVIDIA Nemotron 3 Super		\$0.40
4	Gemini 3 Flash		\$1.10
5	GLM-5.1		\$2.10
6	Grok 4.20 0309 v2		\$3.00
7	Gemini 3.1 Pro Preview		\$4.50
8	GPT-5.4 (xhigh)		\$5.60
9	Claude Sonnet 4.6 (max)		\$6.00
10	Claude Opus 4.6 (max)		\$10.00

\$0.30

Cheapest: gpt-oss-120B and DeepSeek V3.2 per 1M tokens.


\$10.00

Most expensive: Claude Opus 4.6 per 1M tokens.

Gemini 3 Flash

Mid-range value: \$1.10, 153 tokens/s, and 46 intelligence.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Scatter plot measures output tokens consumed during evaluation, not throughput in tokens/s.

Proprietary vs. open-weight leaders

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed	Gemini 3 Flash	153	gpt-oss-120B (high)	218h
Price	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Frontier insight

Raw intelligence, token efficiency, speed, and cost do not collapse into a single winner. Teams need workload-specific routing rather than one default model.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key findings

- **Two-way tie at the top**
Gemini 3.1 Pro Preview and GPT-5.4 both score 57.
- **Speed favors open-weight / smaller models**
gpt-oss-120B reaches 218 tokens/s while scoring 33 intelligence.
- **Cost spread is material**
Cheapest models are \$0.30/M tokens; Claude Opus 4.6 is \$10/M.
- **Open weights are closing the gap**
GLM-5.1 scores 51, within 10% of proprietary leaders.
- **No model wins all three axes**
Selection requires explicit intelligence, speed, and cost trade-offs.

Actionable model-selection guidance

- Prioritize intelligence** **57 index**
Start with Gemini 3.1 Pro Preview or GPT-5.4 for reasoning-heavy paths.
- Prioritize latency** **153–218 tokens/s**
Route interactive and high-volume generation toward gpt-oss-120B, Grok, Nemotron, or Gemini 3 Flash.
- Prioritize budget** **\$0.30–\$0.40/M**
Benchmark gpt-oss-120B, DeepSeek V3.2, and Nemotron first for cost-sensitive deployments.
- Balance capability and openness** **51 index · \$2.10/M**
Evaluate GLM-5.1 when open-weight control matters but capability must remain near frontier.