

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Mo

Independent Analysis of AI Models and API Providers

Source: Artificial Analysis, April 2026


Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Intelligence Index Ranking (Top 11 Models)

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Cl


Output Speed Leaderboard (Tokens/s)



Key Takeaway: Speed and intelligence are inversely correlated. The fastest model scores lowest on intelligence.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models


Price per 1M Tokens (USD) - Arranged Lowest to Highest



Insight: gpt-oss-120B and DeepSeek V3.2 represent the most cost-effective tier, while Gemini 3 Flash sits in the compelling mid-range (\$1.10) for solid speed-intelligence pairing.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence vs. Output Tokens Frontier



Proprietary vs. Open-Weight Summary

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B	\$0.30

Strategic Trade-offs to Consider:

- Best Intelligence: Models like Gemini 3.1 Pro & GPT-5.4
- Highest Output Speed: Dominated by open weights (gpt-oss-120B)
- Strongest Value: Gemini 3 Flash and DeepSeek V3.2 pair solid metrics with highly attractive costs

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key Findings & Recommendations



Two-way Tie at Top

Gemini 3.1 Pro Preview and GPT-5.4 co-lead with an Intelligence Index of 57. Proprietary models still hold the edge on raw capability.



Speed Dominance

gpt-oss-120B (218 tokens/s) leads the pack. Open-weights dominate throughput, showcasing an inverse relationship with intelligence.



Actionable Conclusion

GLM-5.1 scores 51 (near frontier) at just \$2.10/M tokens. No single model rules them all; selection must be explicitly workload-driven.

"For latency-sensitive applications, lean into open weights; for complex reasoning, stick to the proprietary frontier."