


AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33× pricing gap separates the cheapest and most expensive frontier models.

Data source: Artificial Analysis • April 2026

Audience: AI engineering leads & product managers evaluating LLM providers

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57




KEY INSIGHT

Google & OpenAI are neck-and-neck at the frontier. Open-weight models like GLM-5.1 (51) are within 10% of the leaders, signaling the gap is closing fast.

SPREAD

57 → 33
24-point spread across 11 models from 8+ organizations across US, China & Europe.

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4× Faster Than Claude Opus 4.6



KEY INSIGHT

Speed and intelligence are inversely correlated. The fastest model (gpt-oss-120B, 218 tok/s) scores only 33 on intelligence. The smartest models run at 49–118 tok/s.

OPEN-WEIGHT EDGE

Open-weight and smaller models dominate the top of the throughput leaderboard.

From \$0.30 to \$10 per Million Tokens: A 33× Pricing Gap Across Models

Rank	Model	USD / 1M Tokens
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00


BUDGET-FRIENDLY

gpt-oss-120B & DeepSeek V3.2 at \$0.30/M tokens are 33× cheaper than Claude Opus 4.6 (\$10.00).

MID-RANGE VALUE

Gemini 3 Flash (\$1.10) pairs solid speed (153 tok/s) with reasonable intelligence (46) — a compelling value option.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Metric	Proprietary	Score	Open-Weight	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tok/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

SCATTER PLOT HIGHLIGHTS

- Gemini 3.1 Pro & DeepSeek V3.2 → upper region (high intelligence)
- Grok 4.20 & Google models → mid-right (balanced)
- Llama 4 Maverick → low intelligence, low tokens
- MiMo-V2-Pro → high intelligence but heavy token usage

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Two-Way Tie at the Top

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Intelligence Index as of April 2026.

Speed ≠ Intelligence

The fastest model (gpt-oss-120B, 218 tok/s) scores only 33 on intelligence. The smartest models run at 49–118 tok/s.

33× Pricing Spread

\$0.30/M (gpt-oss-120B, DeepSeek V3.2) to \$10.00/M (Claude Opus 4.6). Budget-sensitive workloads have radically cheaper options.

Open-Weight Models Closing In

GLM-5.1 at 51 intelligence is within 10% of the proprietary leaders at \$2.10/M tokens.

No Single Winner

Choosing a model requires explicit trade-offs among intelligence, speed, and cost depending on the use case.

ACTIONABLE GUIDANCE

- Max Intelligence needed?
 - Gemini 3.1 Pro Preview or GPT-5.4
- Latency-sensitive workloads?
 - gpt-oss-120B (218 tok/s) or Grok 4.20
- Budget-constrained at scale?
 - gpt-oss-120B or DeepSeek V3.2 (\$0.30/M)
- Balanced value play?
 - Gemini 3 Flash (\$1.10, 153 tok/s, 46 IQ)
- Open-weight flexibility?
 - GLM-5.1 (51 IQ, \$2.10/M) — closing the gap fast