

AI Model Landscape

April 2026

Benchmarking Intelligence, Speed, and Cost
Across 11 Frontier Models

Data source: Artificial Analysis — Independent LLM Benchmarks, April 2026

For: AI Engineering Leads & Product Managers

11 Models Evaluated


33× Pricing Gap

57 Top Intelligence Score

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index — composite score across reasoning, knowledge & instruction-following benchmarks

Intelligence Index



Key Observations

🏆 Two-way tie at the top

Gemini 3.1 Pro Preview & GPT-5.4 both score 57 — first co-lead ever recorded.

📊 57 → 33 spread

24-point gap between leader and tail model (gpt-oss-120B).

🌐 Global competition

8 distinct organisations across US, China & Europe in top 11.

⚖️ Open weights closing gap

GLM-5.1 scores 51 — within 11% of co-leaders.


🔵 Anthropic

Claude Opus 4.6 (53) and Sonnet 4.6 (52) round out the top 5.

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4× Faster Than Claude Opus 4.6

Output speed measured in tokens per second during inference — critical for latency-sensitive applications

Tokens/s



Key Observations

⚡ 218 tokens/s leader

gpt-oss-120B leads — 4× faster than Claude Opus 4.6 (49 t/s).

🔄 Speed ↔ Intelligence trade-off

The two top-intelligence models (Gemini 3.1 Pro, GPT-5.4) rank 5th and 6th on speed.

🚀 Open-weight advantage

3 of top 4 fastest models are open-weight (gpt-oss, Grok, Nemotron).

💡 Flash = best mid-range

Gemini 3 Flash: 153 t/s with 46 intelligence — compelling for latency needs.

🐌 Frontier tax

Claude Opus 4.6 & DeepSeek V3.2 slowest at 49 & 46 t/s respectively.

From \$0.30 to \$10 per Million Tokens: A 33× Pricing Gap Across Models

API pricing per 1M tokens — blended input/output rate as published by providers, April 2026

Rank	Model	Provider	USD / 1M Tokens
1	gpt-oss-120B (high)	OpenAI	\$0.30
2	DeepSeek V3.2	DeepSeek	\$0.30
3	NVIDIA Nemotron 3 Super	NVIDIA	\$0.40
4	Gemini 3 Flash	Google	\$1.10
5	GLM-5.1	Z AI	\$2.10
6	Grok 4.20 0309 v2	xAI	\$3.00
7	Gemini 3.1 Pro Preview	Google	\$4.50
8	GPT-5.4 (xhigh)	OpenAI	\$5.60
9	Claude Sonnet 4.6 (max)	Anthropic	\$6.00
10	Claude Opus 4.6 (max)	Anthropic	\$10.00

Key Observations

💰 33× cost gap

Cheapest (gpt-oss-120B / DeepSeek, \$0.30) vs. most expensive (Claude Opus 4.6, \$10.00).

🏆 Best value: Flash

Gemini 3 Flash at \$1.10 delivers 153 t/s + intelligence score 46.

🌐 Open-weight dominates low cost

Top 3 cheapest are all open-weight / open-access models.


🦖 Frontier premium

GPT-5.4, Claude Sonnet, and Claude Opus carry a \$5.60–\$10.00 price tag.

📌 Budget guidance

For cost-sensitive workloads, gpt-oss-120B or DeepSeek V3.2 offer strong ROI at \$0.30.

No Model Wins on All Three Axes: Navigating the Intelligence–Efficiency–Cost Frontier



Proprietary vs. Open-Weight Summary

Metric	Proprietary Leader	Score	Open-Weight Leader
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1
Speed (t/s)	Gemini 3 Flash	153	gpt-oss-120B
Price (\$/1M)	Gemini 3 Flash	\$1.10	gpt-oss-120B

Core Trade-offs

🧠 Intelligence ↔ Speed: Smartest models (57 index) run at 83–118 t/s; fastest (218 t/s) scores only 33.

💰 Intelligence ↔ Cost: Top models cost \$4.50–\$10/1M tokens; cheapest leaders score 33–42 on intelligence.

⚡ Speed ↔ Cost: High-speed open-weight models (\$0.30–\$0.40) vastly undercut proprietary fast options.

No single model dominates all three dimensions simultaneously — workload type must drive selection.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key Findings

Two-way tie at the frontier

Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 — marking the first recorded co-lead on the Artificial Analysis Intelligence Index.

Speed and intelligence are inversely correlated

The fastest model (gpt-oss-120B, 218 t/s) scores 33 on intelligence; the smartest operate at 49–118 t/s.

Massive 33× pricing spread

A \$0.30 floor (gpt-oss-120B, DeepSeek V3.2) vs. a \$10.00 ceiling (Claude Opus 4.6) creates extreme cost dispersion.

Open-weight models are closing the gap

GLM-5.1 scores 51 (within 11% of co-leaders) at \$2.10/1M tokens — a compelling value proposition.

Actionable Recommendations

🎯 Prioritise intelligence?

Use Gemini 3.1 Pro Preview or GPT-5.4 (xhigh). Expect \$4.50–\$5.60/1M and moderate throughput.

⚡ Need high throughput?

gpt-oss-120B (218 t/s, \$0.30/1M) or Grok 4.20 (188 t/s, \$3.00/1M) for latency-sensitive apps.

💰 Optimising cost?

gpt-oss-120B or DeepSeek V3.2 at \$0.30/1M tokens. Consider GLM-5.1 (\$2.10) for higher quality.

⚖️ Balanced workloads?

Gemini 3 Flash (\$1.10, 153 t/s, score 46) provides the best all-round value proposition.

🔍 Evaluate open-weight models

Open-weight gap is ≤11% on intelligence. Audit security & compliance before production deployment.