AI Model Landscape April 2026

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33× pricing gap separates the cheapest and most expensive frontier models.

Data source: Artificial Analysis | Captured April 10, 2026

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index — composite score across reasoning, knowledge, and instruction-following benchmarks

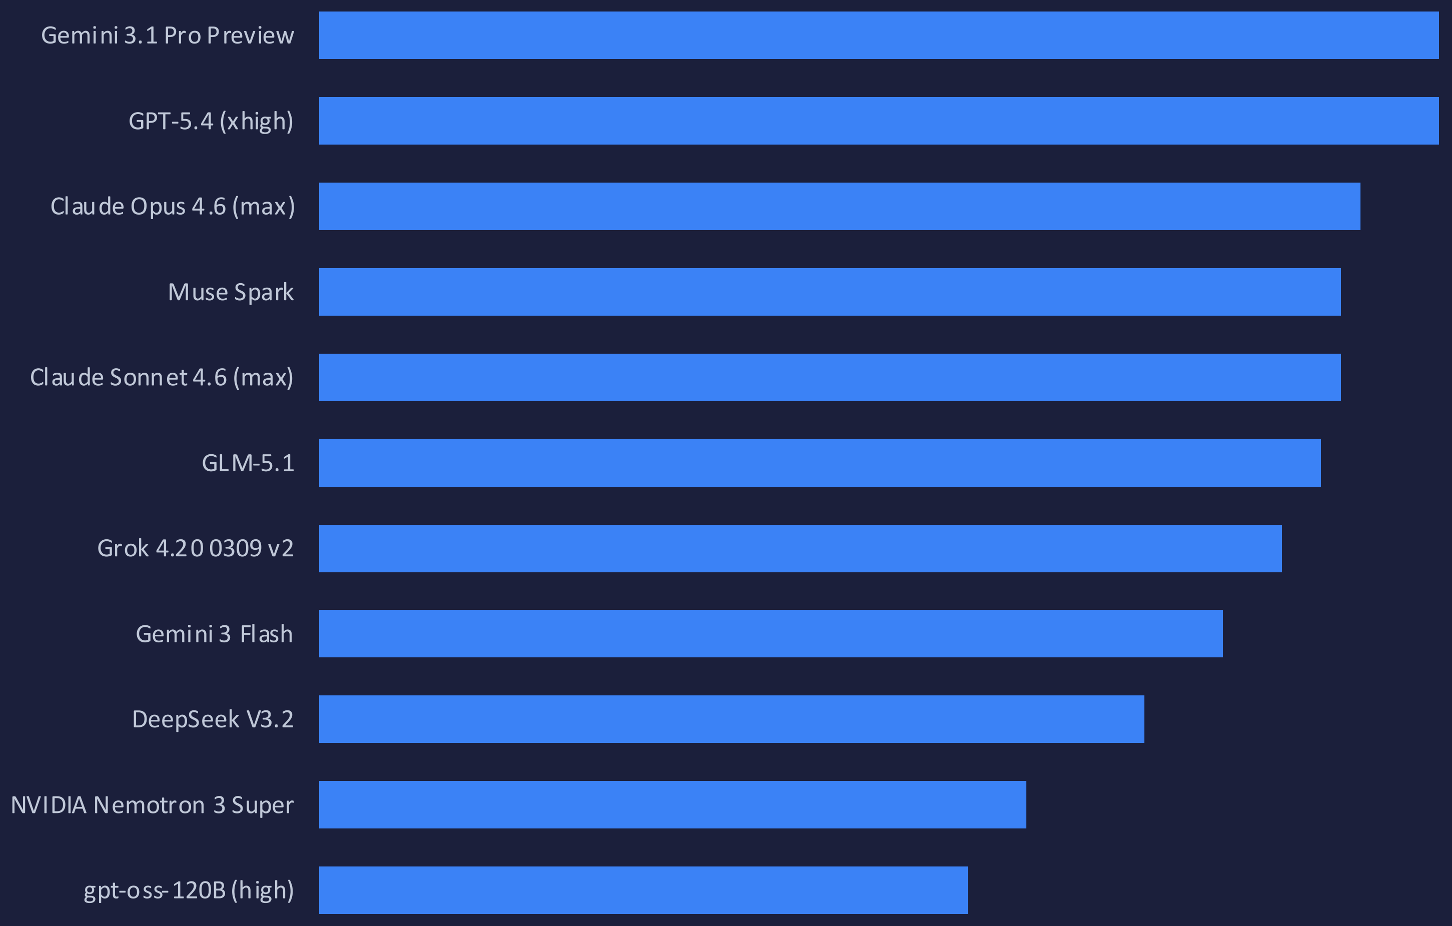

Two-way tie at the top

Gemini 3.1 Pro Preview and GPT-5.4 both score 57, making them co-leaders as of April 2026.

24-point spread

The range from 57 (co-leaders) down to 33 (gpt-oss-120B) reveals significant capability stratification.

Open-weight in the mix

GLM-5.1 scores 51 — just 6 points behind the co-leaders — showing open-weight models are competitive on intelligence.

Anthropic depth

Both Claude Opus 4.6 (53) and Sonnet 4.6 (52) rank in the top 5.

gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4× Faster Than Claude Opus 4.6

Output tokens per second — critical for latency-sensitive applications | Muse Spark excluded (no speed data)

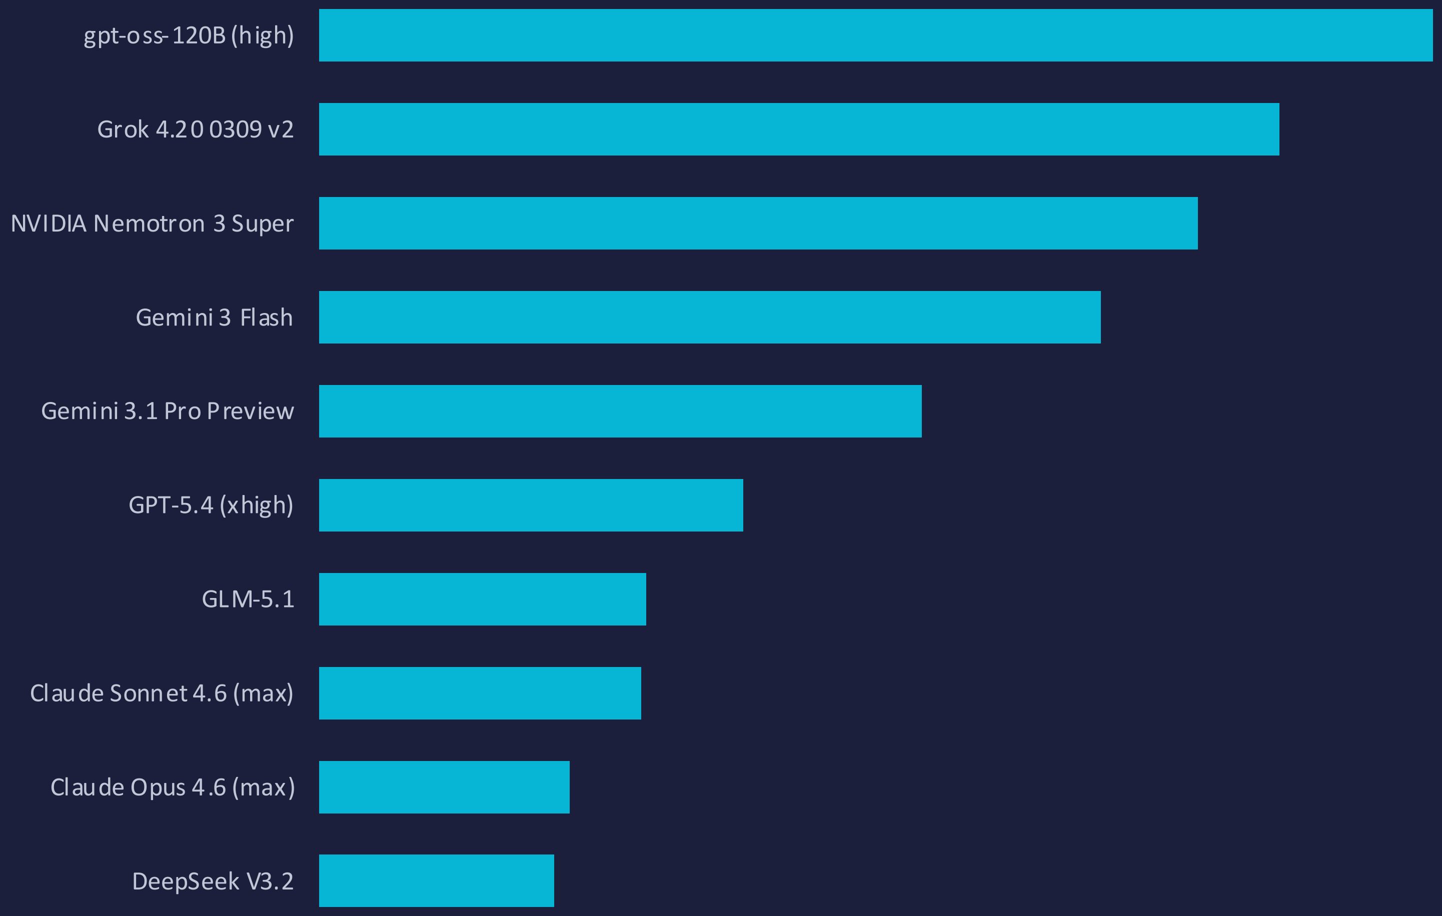

4.4× speed gap

The fastest model (gpt-oss-120B at 218 tokens/s) is more than 4× faster than Claude Opus 4.6 (49 tokens/s) and nearly 5× faster than DeepSeek V3.2 (46 tokens/s).

Speed vs. intelligence trade-off

The two intelligence leaders sit mid-table on speed: Gemini 3.1 Pro at 118 tokens/s and GPT-5.4 at 83 tokens/s.

Open-weight speed dominance

Three of the top four fastest models are open-weight, demonstrating the throughput advantage of smaller or more efficient architectures.

Budget + speed sweet spot

Gemini 3 Flash pairs solid speed (153 tokens/s) with reasonable intelligence (46), making it a strong choice for latency-sensitive workloads.

From $0.30 to $10 per Million Tokens: A 33× Pricing Gap Across Models

API pricing per million tokens | Muse Spark excluded (no pricing data)

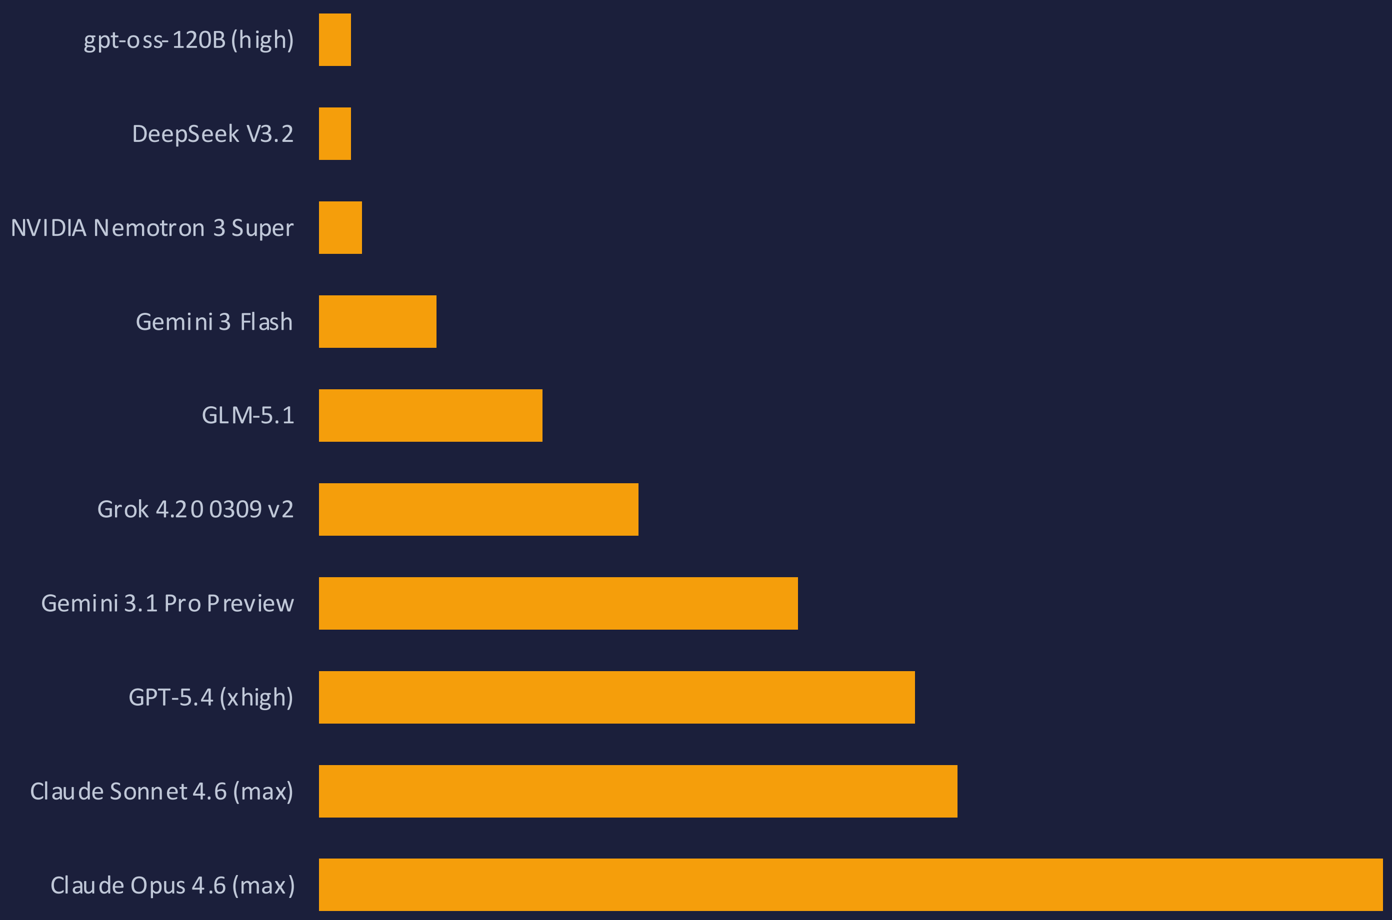

33× cost gap

The cheapest options ($0.30/1M tokens for gpt-oss-120B and DeepSeek V3.2) are over 33× less expensive than Claude Opus 4.6 ($10.00/1M tokens).

Budget-friendly cluster

Three models cost ≤ $0.40/1M tokens — all open-weight — making them ideal for high-volume, cost-sensitive workloads.

Mid-range value

Gemini 3 Flash at $1.10/1M tokens offers a compelling balance of speed (153 tokens/s) and intelligence (46).

Premium intelligence

The top-3 intelligence models (Gemini 3.1 Pro $4.50, GPT-5.4 $5.60, Claude Opus $10.00) carry premium pricing that must be justified by use-case requirements.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence vs. output tokens consumed during evaluation — higher intelligence with fewer tokens = more efficient architecture

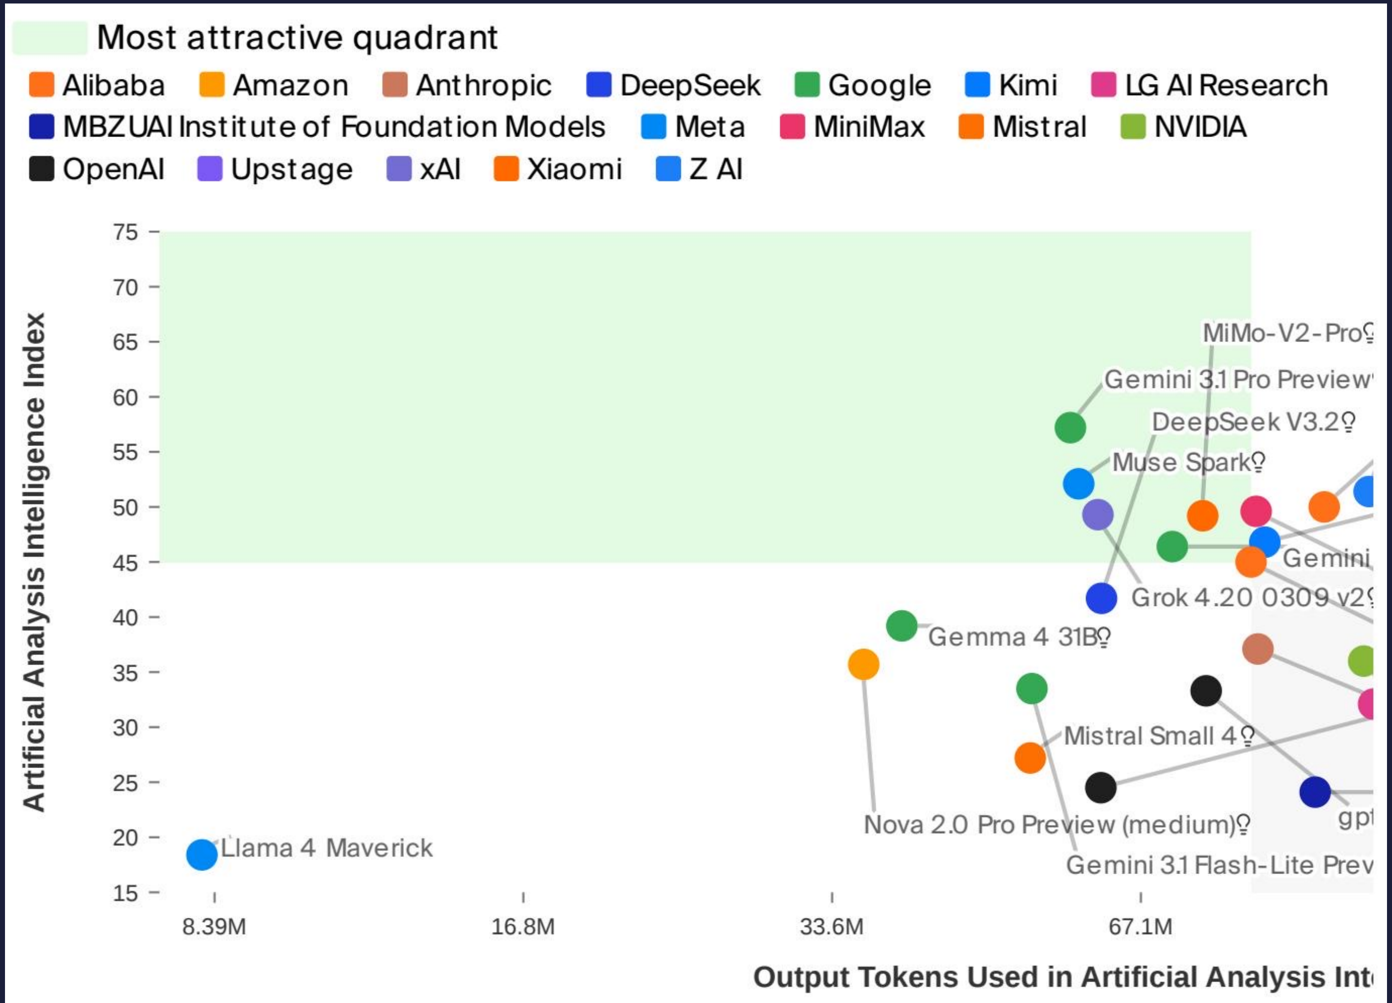

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	**57**	GLM-5.1	**51**
Speed (tokens/s)	Gemini 3 Flash	**153**	gpt-oss-120B (high)	**218**
Price ($/1M tokens)	Gemini 3 Flash	**$1.10**	gpt-oss-120B (high)	**$0.30**

- Speed and intelligence are inversely correlated — the fastest model (gpt-oss-120B, 218 tokens/s) scores only 33 on intelligence.
- Open-weight models lead on speed and price, but proprietary models still edge on raw intelligence.
- GLM-5.1 at 51 intelligence is within 10% of the proprietary leaders, while costing $2.10/M tokens — well below the frontier average.
- Every model requires explicit trade-offs — choose based on your workload's priority axis.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at $2.10/M Tokens

KEY FINDINGS

Two-way tie at the top

Gemini 3.1 Pro Preview and GPT-5.4 both score 57 on the Intelligence Index — co-leaders as of April 2026.

Speed–intelligence trade-off

The fastest model (gpt-oss-120B at 218 tokens/s) scores only 33 on intelligence; the smartest models operate at 49–118 tokens/s.

33× pricing spread

$0.30/1M tokens (gpt-oss-120B, DeepSeek V3.2) to $10.00 (Claude Opus 4.6) — cost must be weighed against capability needs.

Open-weight gap closing

GLM-5.1 at 51 intelligence is within 10% of proprietary leaders, at $2.10/M tokens — well below the frontier average.

No single model dominates

No model wins on intelligence, speed, and cost simultaneously — selection must be use-case driven.

ACTIONABLE RECOMMENDATIONS

Max-intelligence workloads

Gemini 3.1 Pro Preview or GPT-5.4 for complex reasoning, research synthesis, and high-stakes decision support where output quality is paramount.

Latency-sensitive pipelines

gpt-oss-120B or Grok 4.20 for real-time chat, streaming assistants, and interactive tools where sub-second token delivery matters.

Cost-optimized volume

DeepSeek V3.2 or gpt-oss-120B at $0.30/1M tokens for bulk classification, summarization, and data enrichment tasks at scale.

Balanced mid-range

Gemini 3 Flash ($1.10/1M tokens, 153 tokens/s, intelligence 46) for workloads that need reasonable quality at speed and moderate cost.

Future-proofing

Monitor open-weight releases — the gap is narrowing fast. Build evaluation harnesses that let you swap models as the landscape shifts.